

Advanced Topic in Research Data-centric Deep Learning

Lec 13: Generative Adaversal Networks : GAN



hsyi@kisti.re.kr

Hongsuk Yi (이홍석)



Reviewing the last class: Attention

- Machine Learning is to find a function f
- Un-Supervised
 - ✓ Data – X
 - ✓ Goal – Learn Hidden structure of data
- Supervised
 - ✓ Data – X, y
 - ✓ Goal – Learn mapping from $X \rightarrow Y$
- Output
 - ✓ Regression
 - ✓ Classification
 - ✓ Prediction/Structure Learning

$$f : X \rightarrow Y$$

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

Maximum Likelihood Models

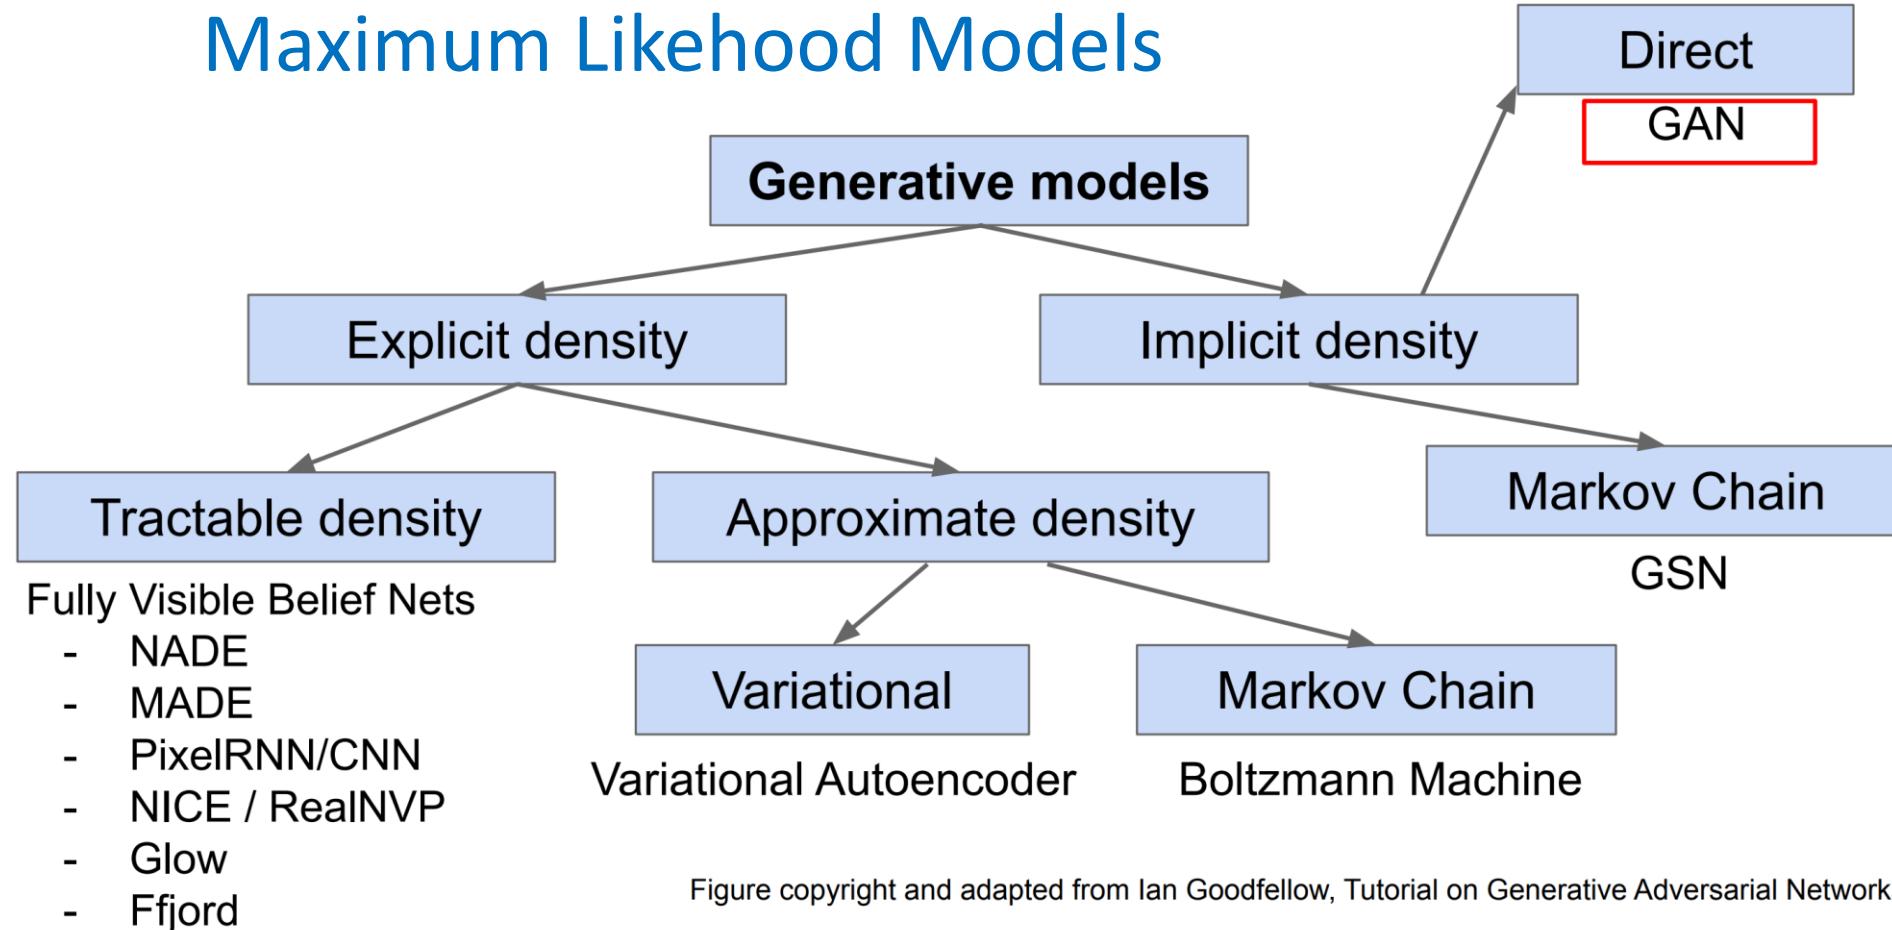
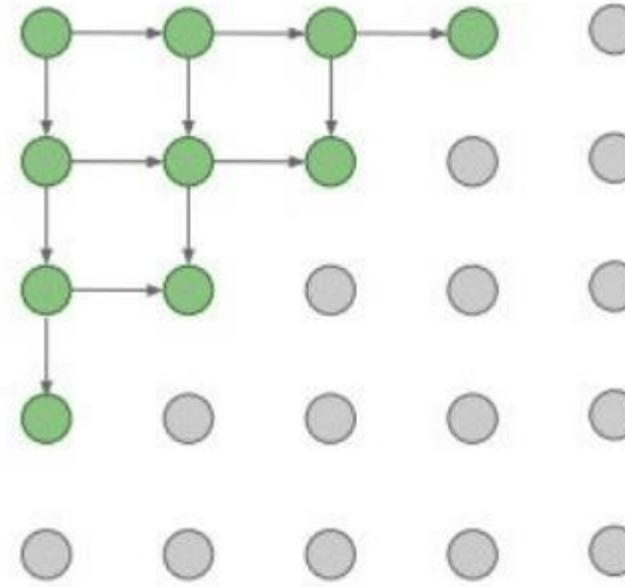


Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017.

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

- Generate image pixels from corner
- Training Faster
- Generation Slow / Sequential
- Cannot generate samples based on some latent code



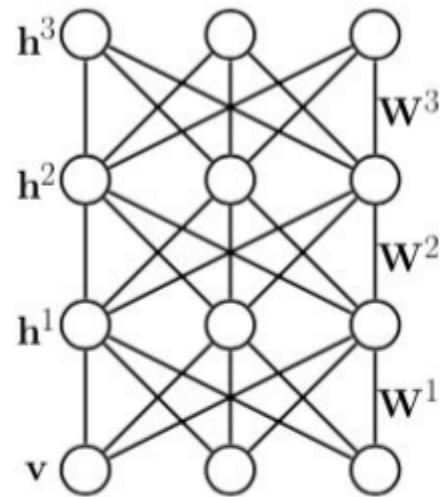
$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1})$$

Chain Rule

Maximum Likelihood based Training

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

Boltzmann Machine



- Energy Function based models
- Markov chains don't work for long sequences
- Hard to scale on large dataset

$$p(x, h) = \exp(-E(x, h)) \mid Z$$

$$Z = \sum_{x,h} \exp(-E(x, h))$$

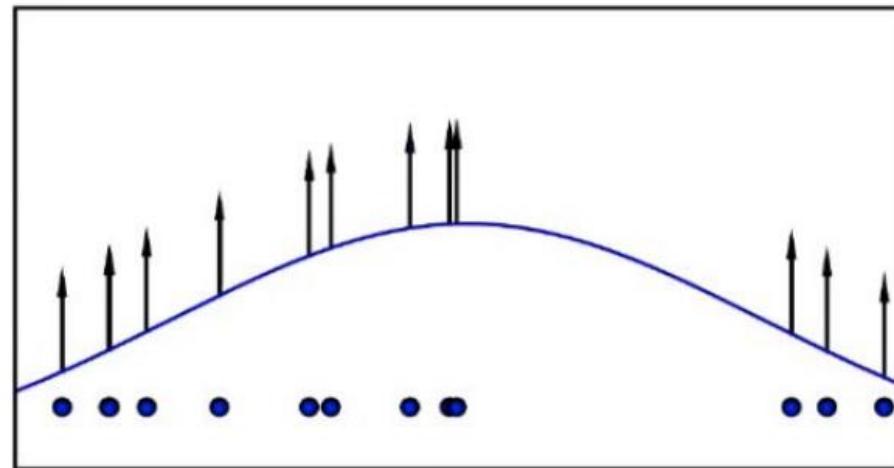
<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

Maximum Likelihood based Models

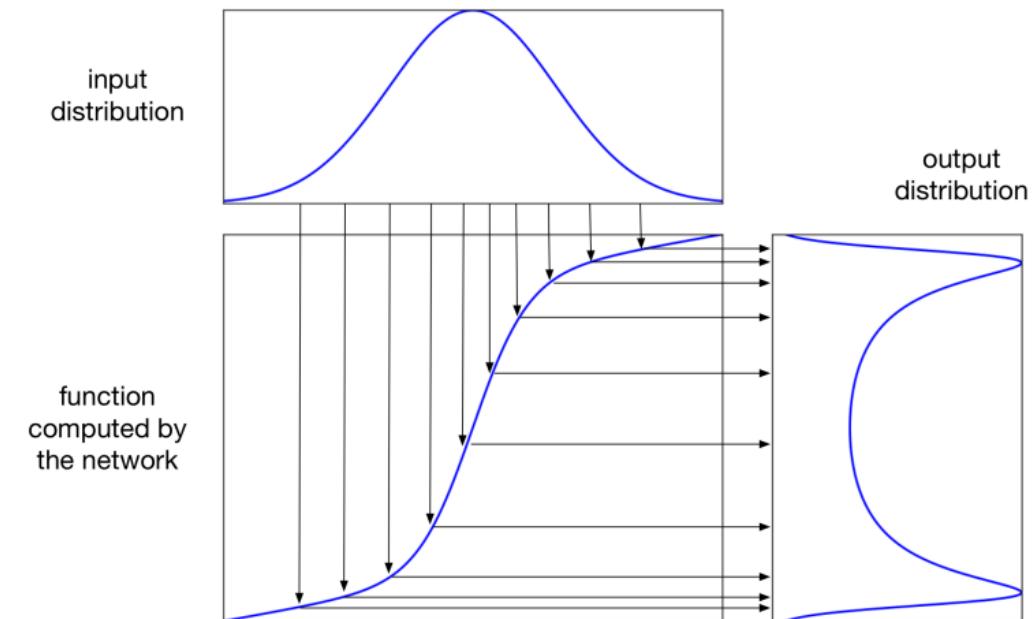
- Maximum likelihood tries increase the likelihood of data given the parameters

$P(x)$

$$\theta^* = \arg \max_{\theta} E_{x \sim P_{data}} \log P(x/\theta)$$

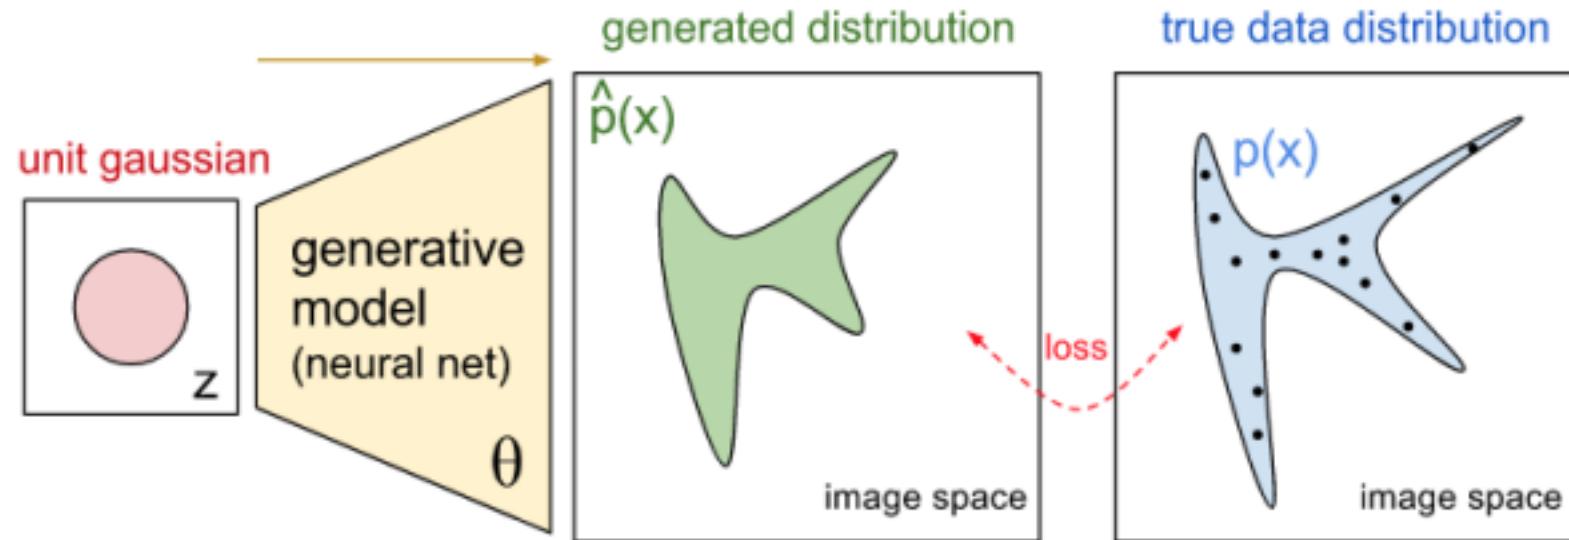


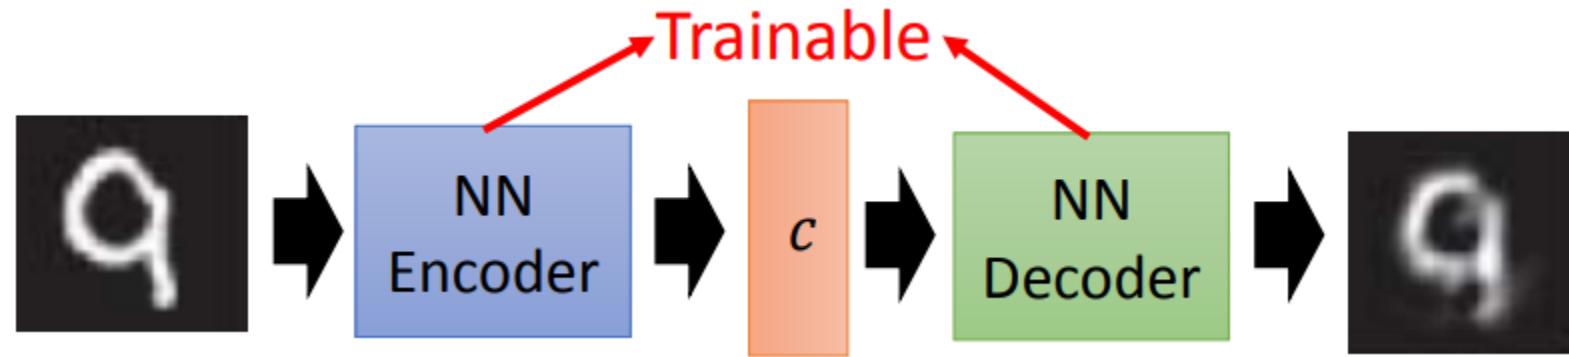
A 1-dimensional example



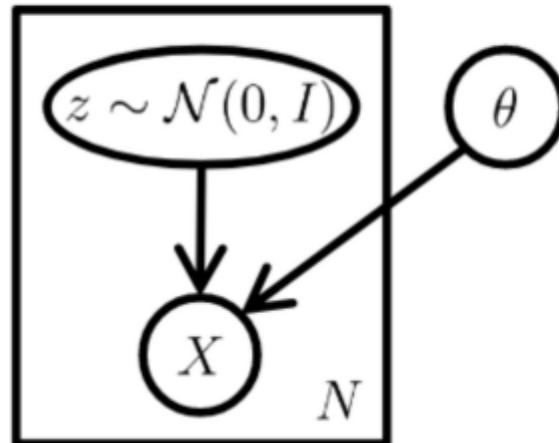
<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

- Generative models define a probability distribution
 - ✓ Start by sampling the **code vector z** from a fixed, simple distribution
 - ✓ The **generator** network computes a differentiable function **G** mapping **z** to an **x** in data space





Variational Auto-encoder

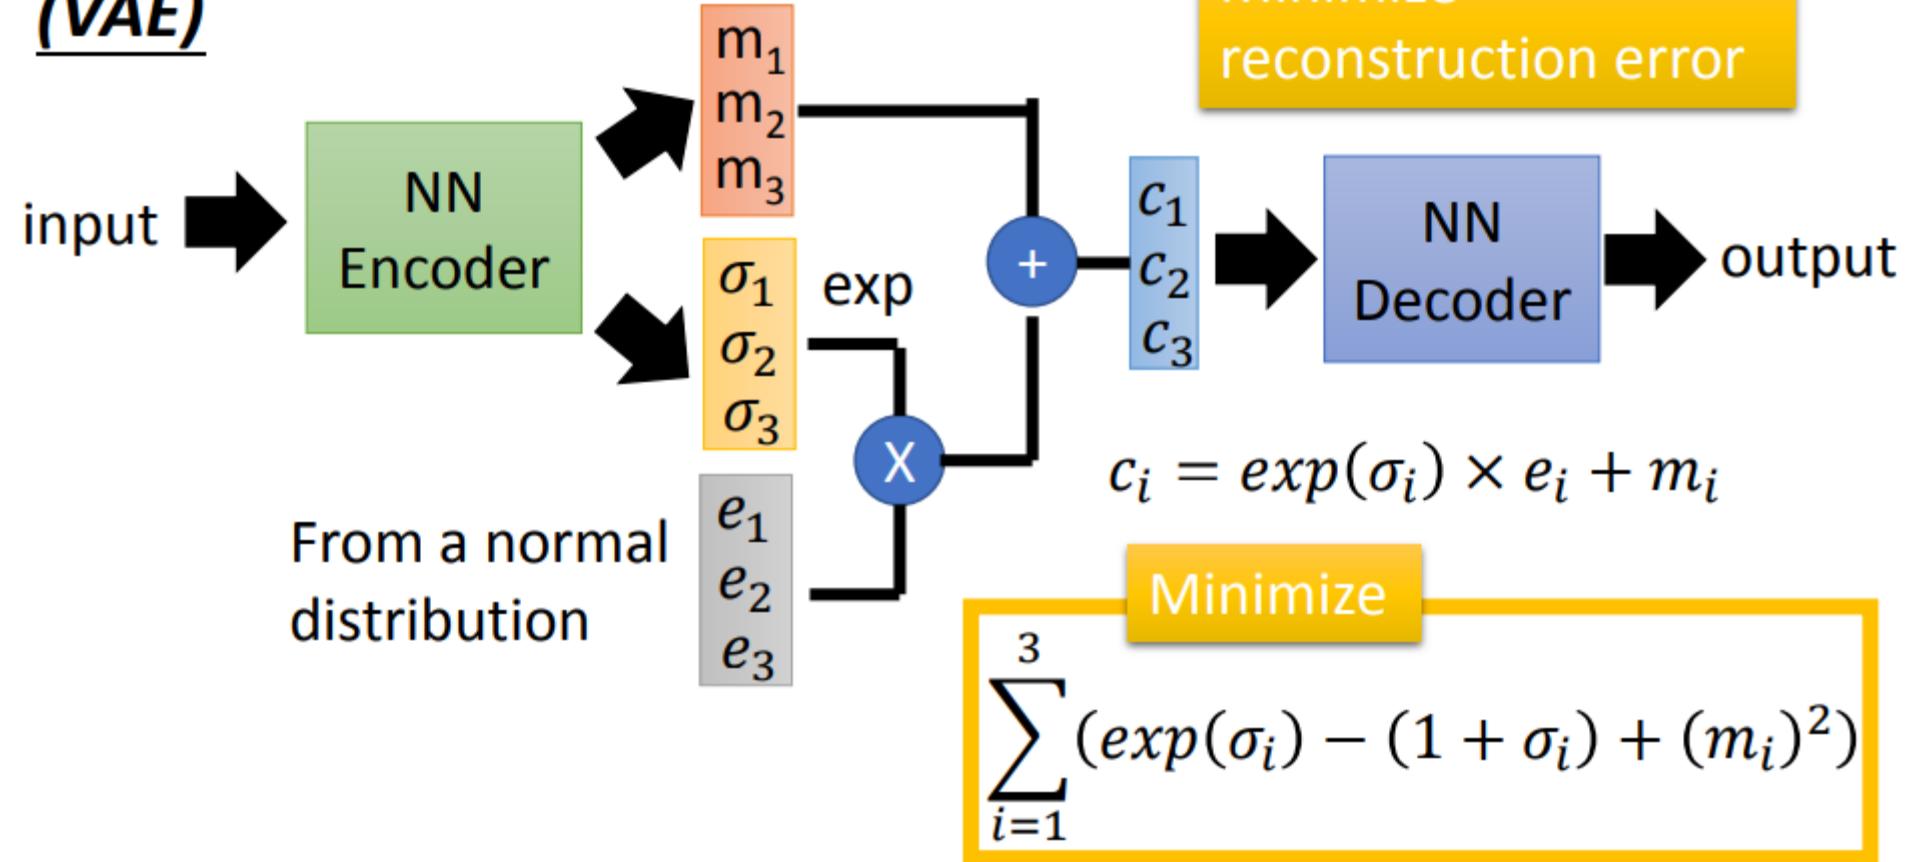


- Model is able to achieve high likelihood
- Model is not asymptotically consistent unless q is perfect
- Samples tend to have lower quality

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

Variational Auto-encoder : VAE

Variational Auto-encoder (VAE)

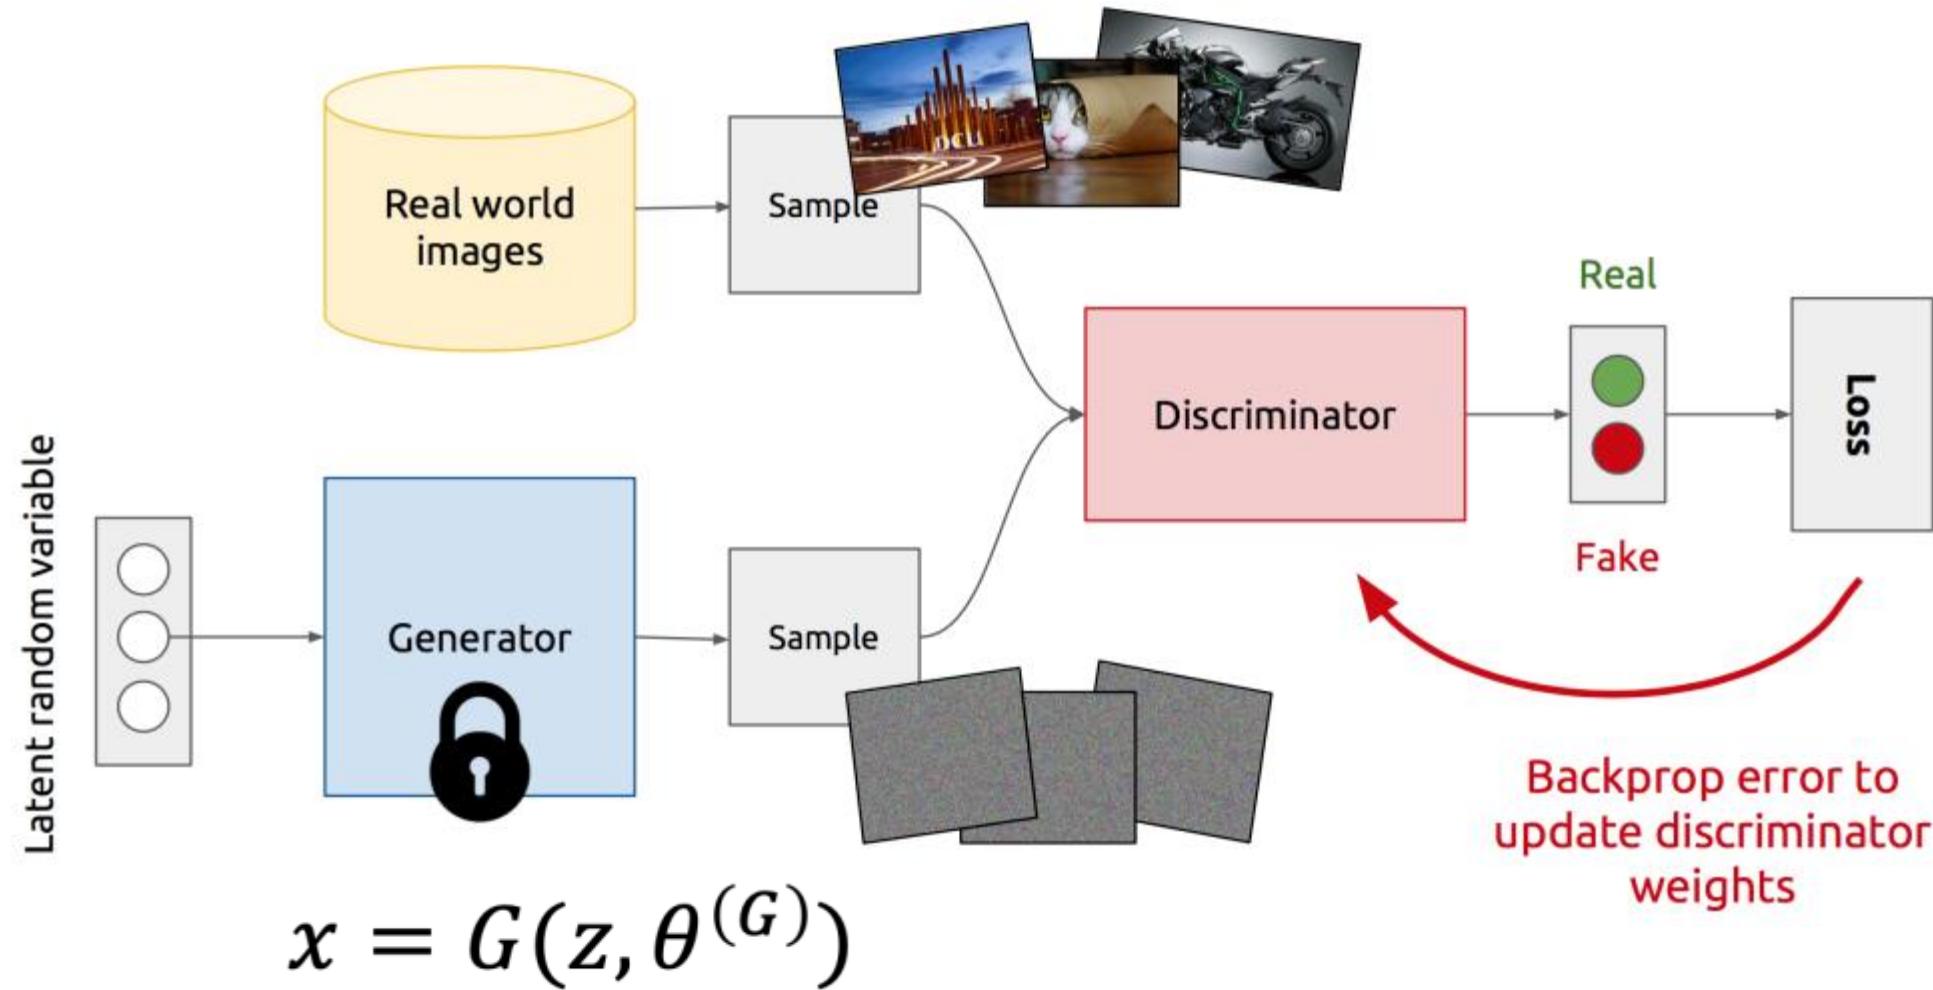


Generative Adversarial Networks (GANs)

- Plenty of existing work on Deep Generative Models
 - ✓ Boltzmann Machine
 - ✓ Deep Belief Nets
 - ✓ Variational AutoEncoders (VAE)

- Why GANs?
 - ✓ Sampling (or generation) is straightforward
 - ✓ Training doesn't involve Maximum Likelihood estimation
 - ✓ Robust to Overfitting since Generator never sees the training data
 - ✓ Empirically, GANs are good at capturing the modes of the distribution.

➤ Classic GAN Framework



- GANs extend that idea to generative models:
 - ✓ generate adversarial samples to fool a discriminative model
 - ✓ use those adversarial samples to make models robust
 - ✓ Repeat this and we get better discriminative model
- Generator:
 - ✓ generate fake samples, tries to fool the Discriminator
- Discriminator:
 - ✓ tries to distinguish between real and fake samples
 - ✓ Train them against each other
 - ✓ Repeat this and we get better Generator and Discriminator

- D tries to identify real data from fakes created by the generator
- G tries to create imitations of data to trick the discriminator
- Objective function:

Train GAN jointly via **minimax** game:

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log \left(1 - D_{\theta_d} (G_{\theta_g}(z)) \right) \right]$$

Discriminator wants to maximize objective s.t. $D(x)$ close to 1, $D(G(z))$ close to 0.

Generator wants to minimize objective s.t. $D(G(z))$ close to 1.

Discriminator

- Discriminator is a function D (network, can deep)

$$D : X \rightarrow R$$

- Input x : an object x (e.g. an image)
- Output $D(x)$: scalar which represents how “good” an object x is



Can we use the discriminator to generate objects?

Yes.

Discriminator

- Suppose we already have a good discriminator $D(x)$...

Inference

- Generate object \tilde{x} that

$$\tilde{x} = \arg \max_{x \in X} D(x)$$

Enumerate all possible x !!!

It is feasible ???

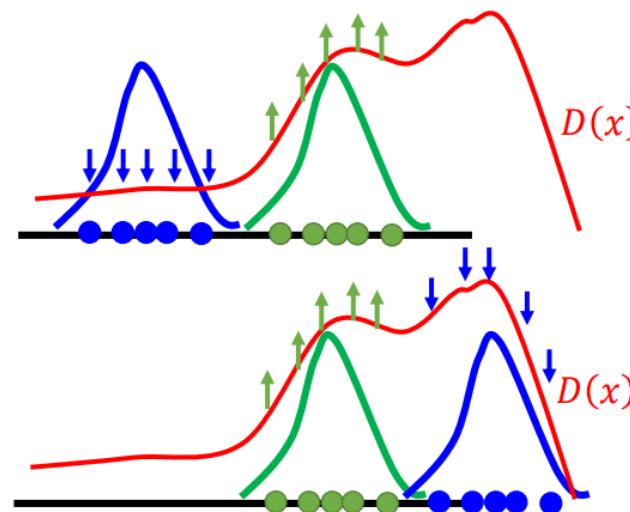
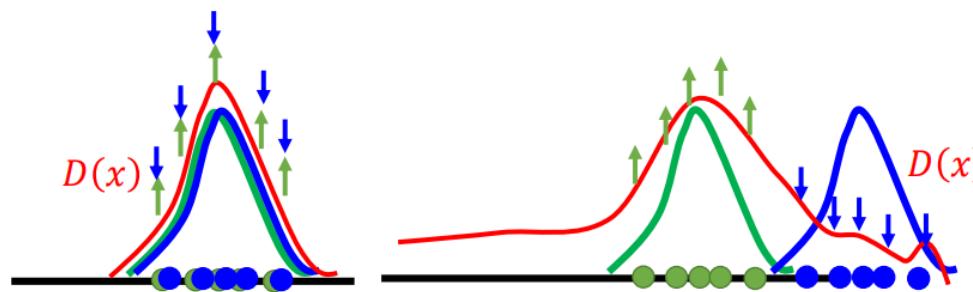
How to learn the discriminator?

Discriminator - Training

Discriminator - Training



In the end



• General Algorithm

- Given a set of **positive examples**, randomly generate a set of **negative examples**.



- In each iteration

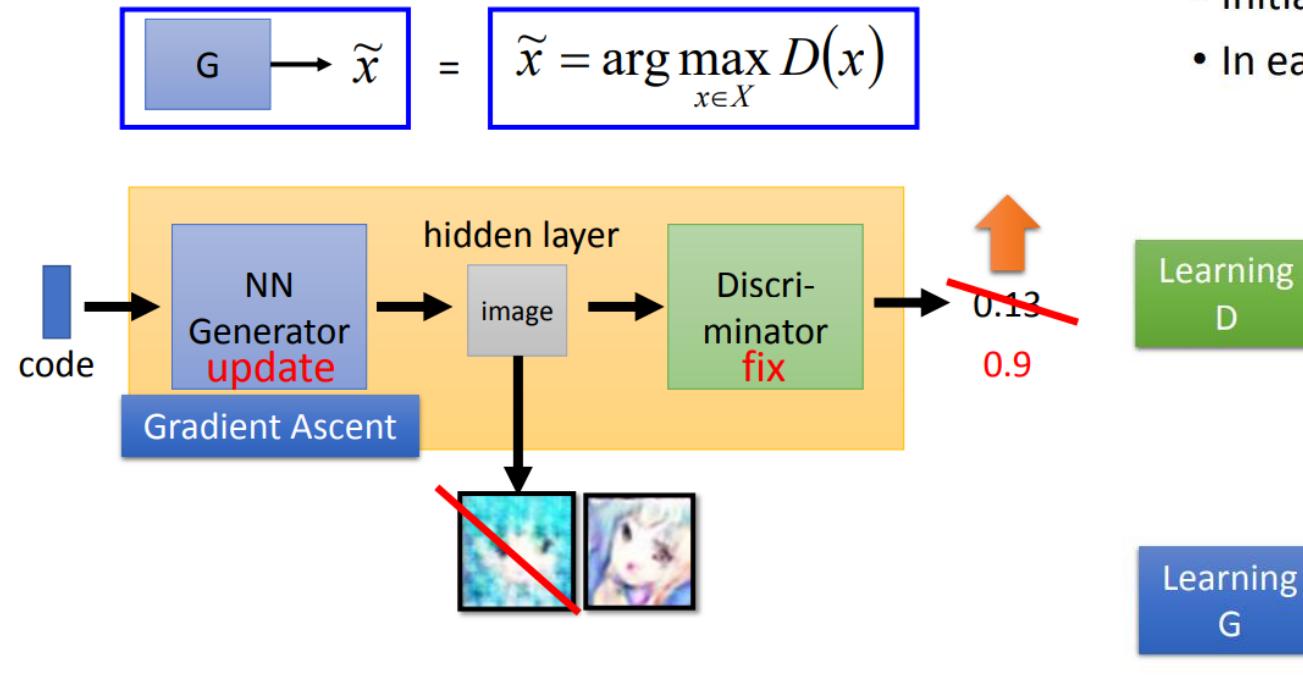
- Learn a discriminator D that can discriminate positive and negative examples.



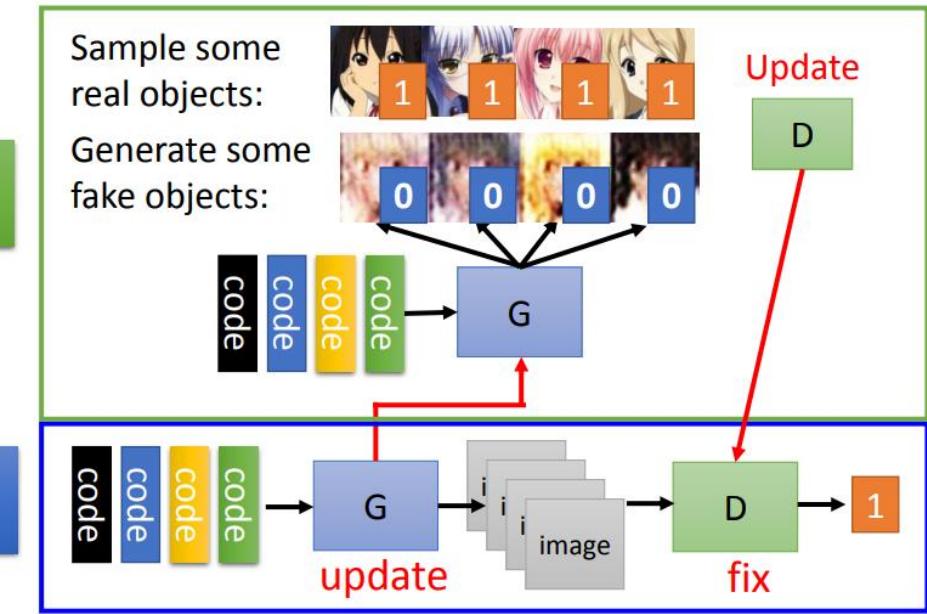
- Generate negative examples by discriminator D

$$G \rightarrow \tilde{x} = \arg \max_{x \in X} D(x)$$

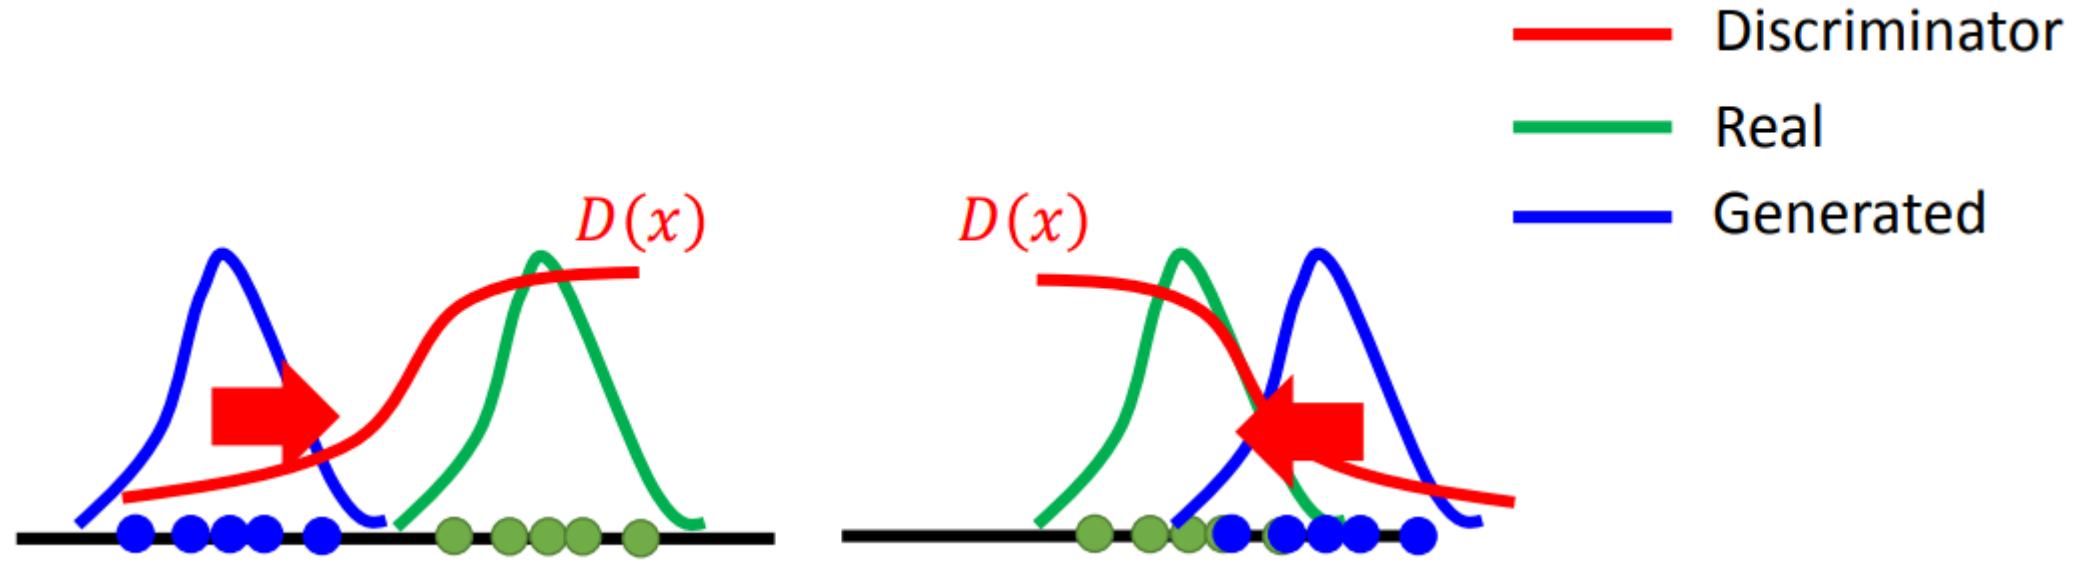
Generating Negative Examples



- Initialize generator and discriminator  
- In each training iteration:



- Discriminator leads the generator



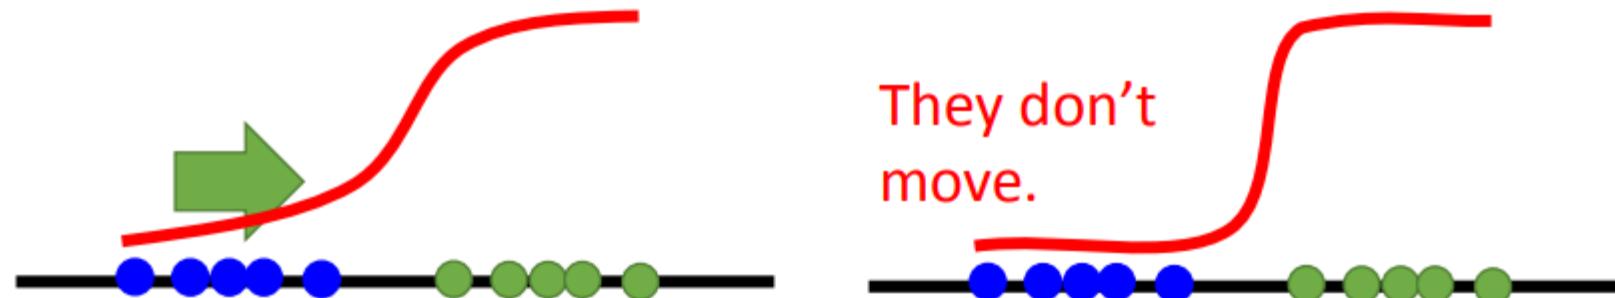
Binary Classifier as Discriminator



Typical binary classifier uses sigmoid function
at the output layer

1 is the largest, 0 is the smallest

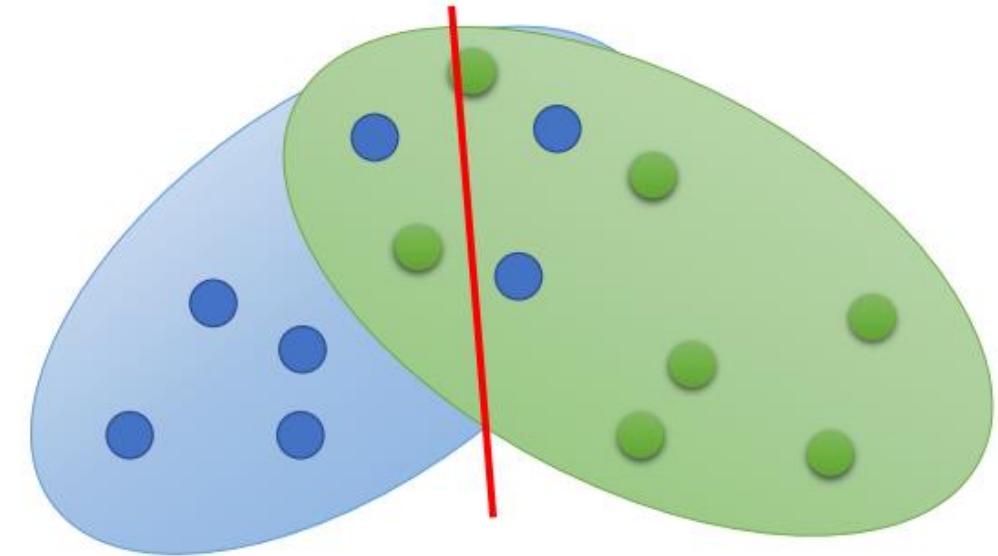
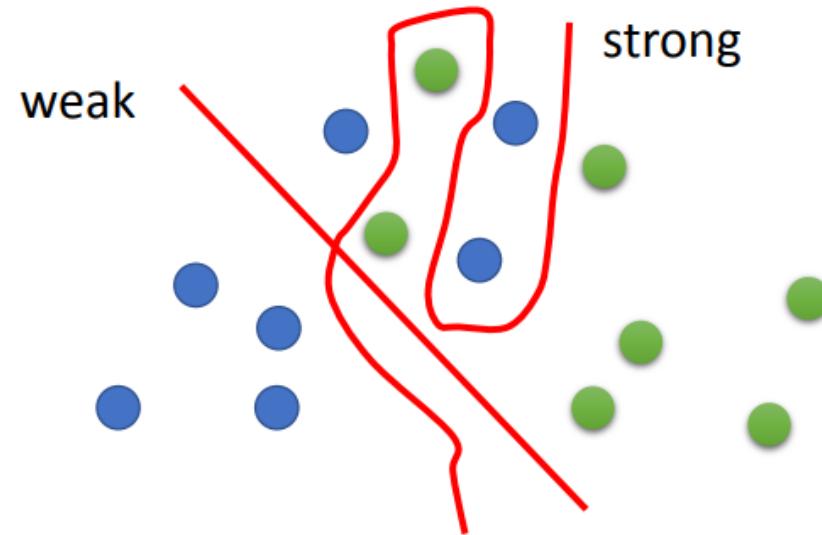
- real
- generated



You cannot train your classifier too good.....

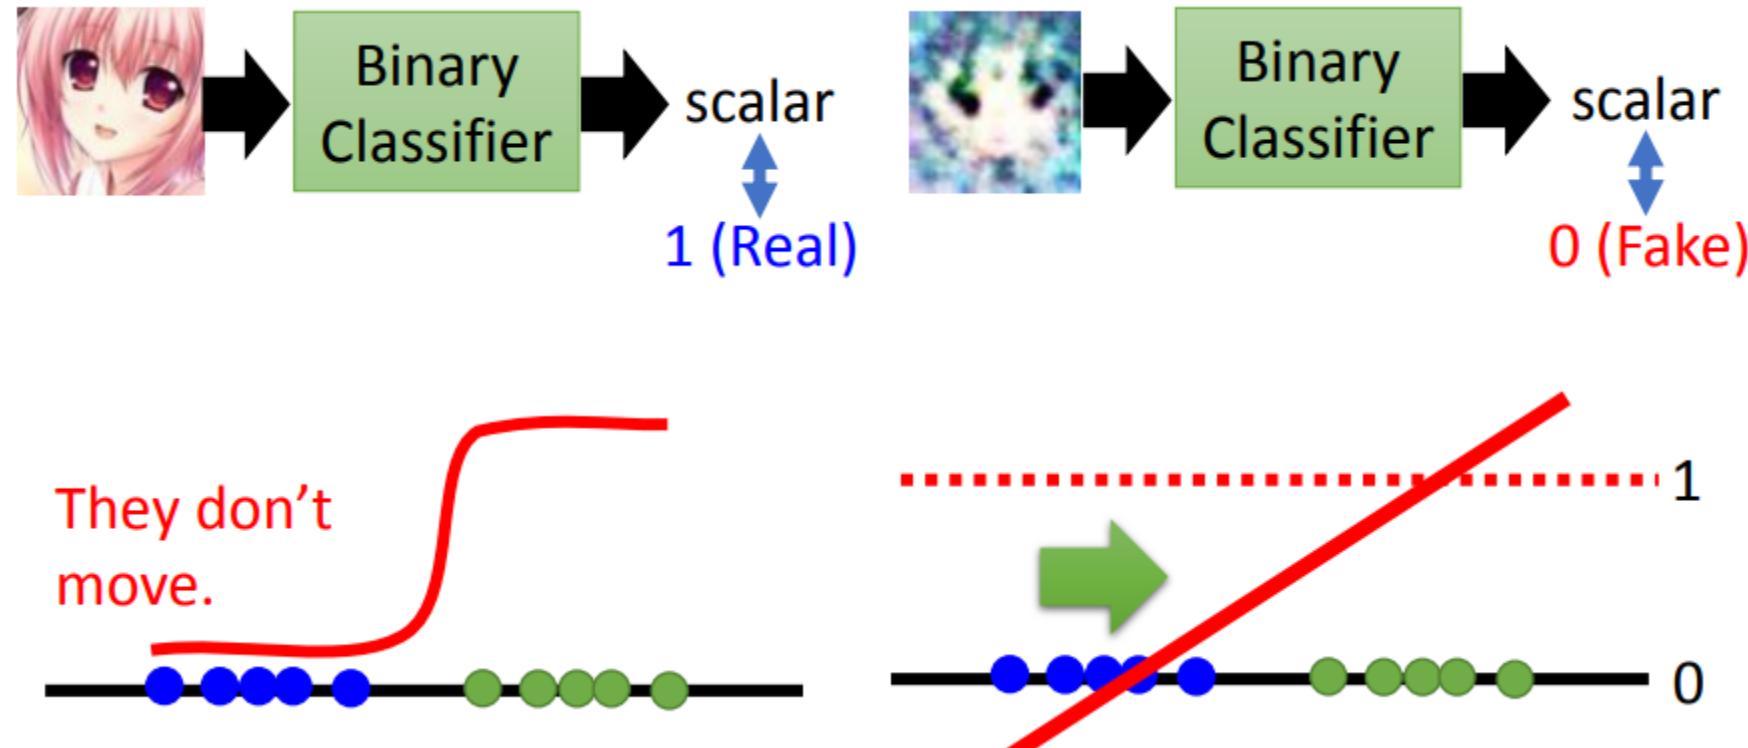
Binary Classifier as Discriminator

- Don't let the discriminator perfectly separate real and generated data
 - ✓ Add noise to input or label?



Least Square GAN (LSGAN)

- Replace sigmoid with linear (replace classification with regression)



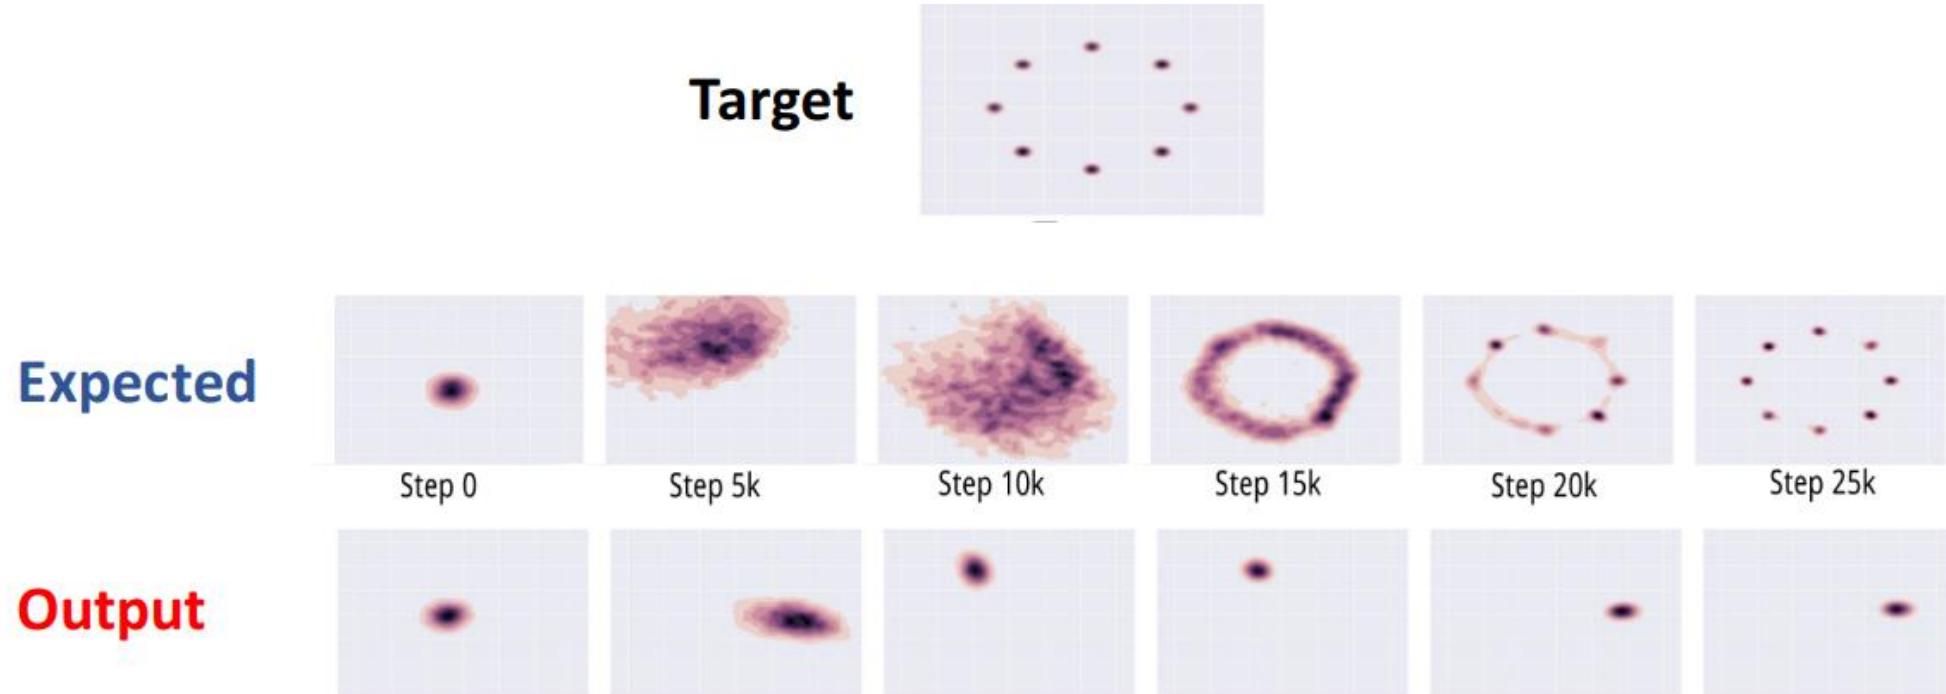
The Difficulties of Training GANs

- zero-sum game:
 - ✓ During training, the generator and the discriminator constantly try to outsmart each other
- Nash equilibrium:
 - ✓ no player would be better off changing their own strategy, assuming the other players do not change theirs
- The biggest difficulty : mode collapse
 - ✓ the generator's outputs gradually become less diverse
- GANs are very sensitive to the hyperparameters:
 - ✓ you may have to spend a lot of effort fine-tuning them.

- Experience replay to avoid the mode collapse issue:
 - ✓ Google 2018 paper.
 - ✓ storing the images produced by the generator at each iteration in a replay buffer
- Mini-batch discrimination:
 - ✓ how similar images are across the batch and provides this statistic to the discriminator
 - ✓ it can easily reject a whole batch of fake images that lack diversity

Mode-Collapse and Solutions

- Generator fails to output diverse samples



- Basic Solutions:
 - ✓ Mini-Batch GANs

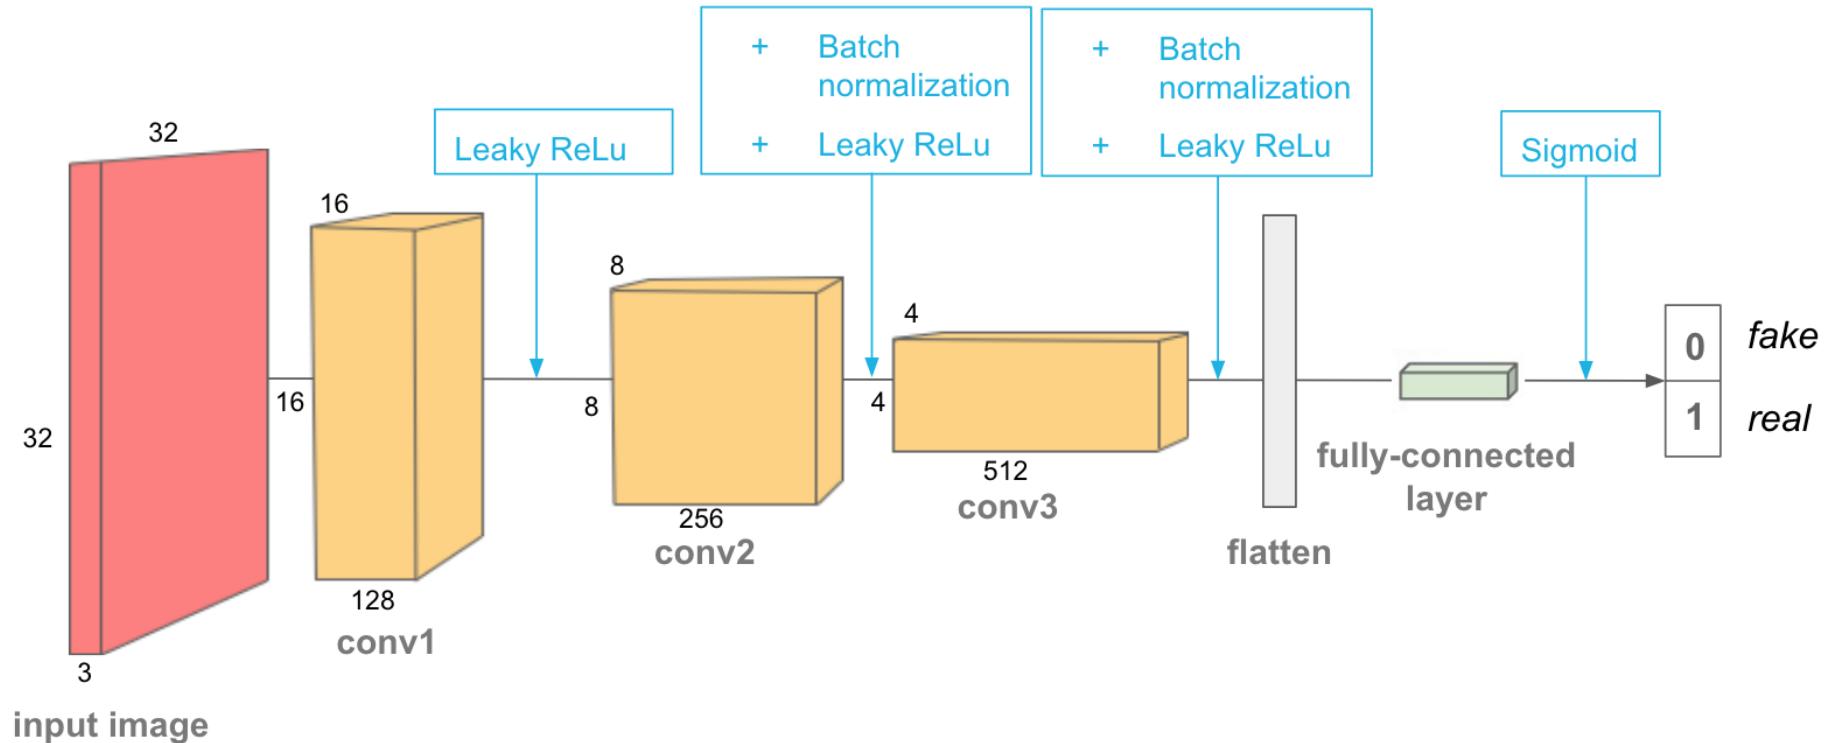
More Recent GANs

- The original GAN paper in 2014 experimented with convolutional layers:
 - ✓ But, only tried to generate small images
- DCGAN : Deep Convolutional GANs: 2015, Alec Radford
 - ✓ Replace any pooling layers with strided convolutions (in the discriminator) and transposed convolutions (in the generator).
 - ✓ Use Batch Normalization in both the generator and the discriminator, except in the generator's output layer and the discriminator's input layer.
 - ✓ Remove fully connected hidden layers for deeper architectures.
 - ✓ Use ReLU activation in the generator for all layers except the output layer, which should use tanh.
 - ✓ Use leaky ReLU activation in the discriminator for all layers.

Deep Convolutional GAN or DCGAN

➤ discriminator:

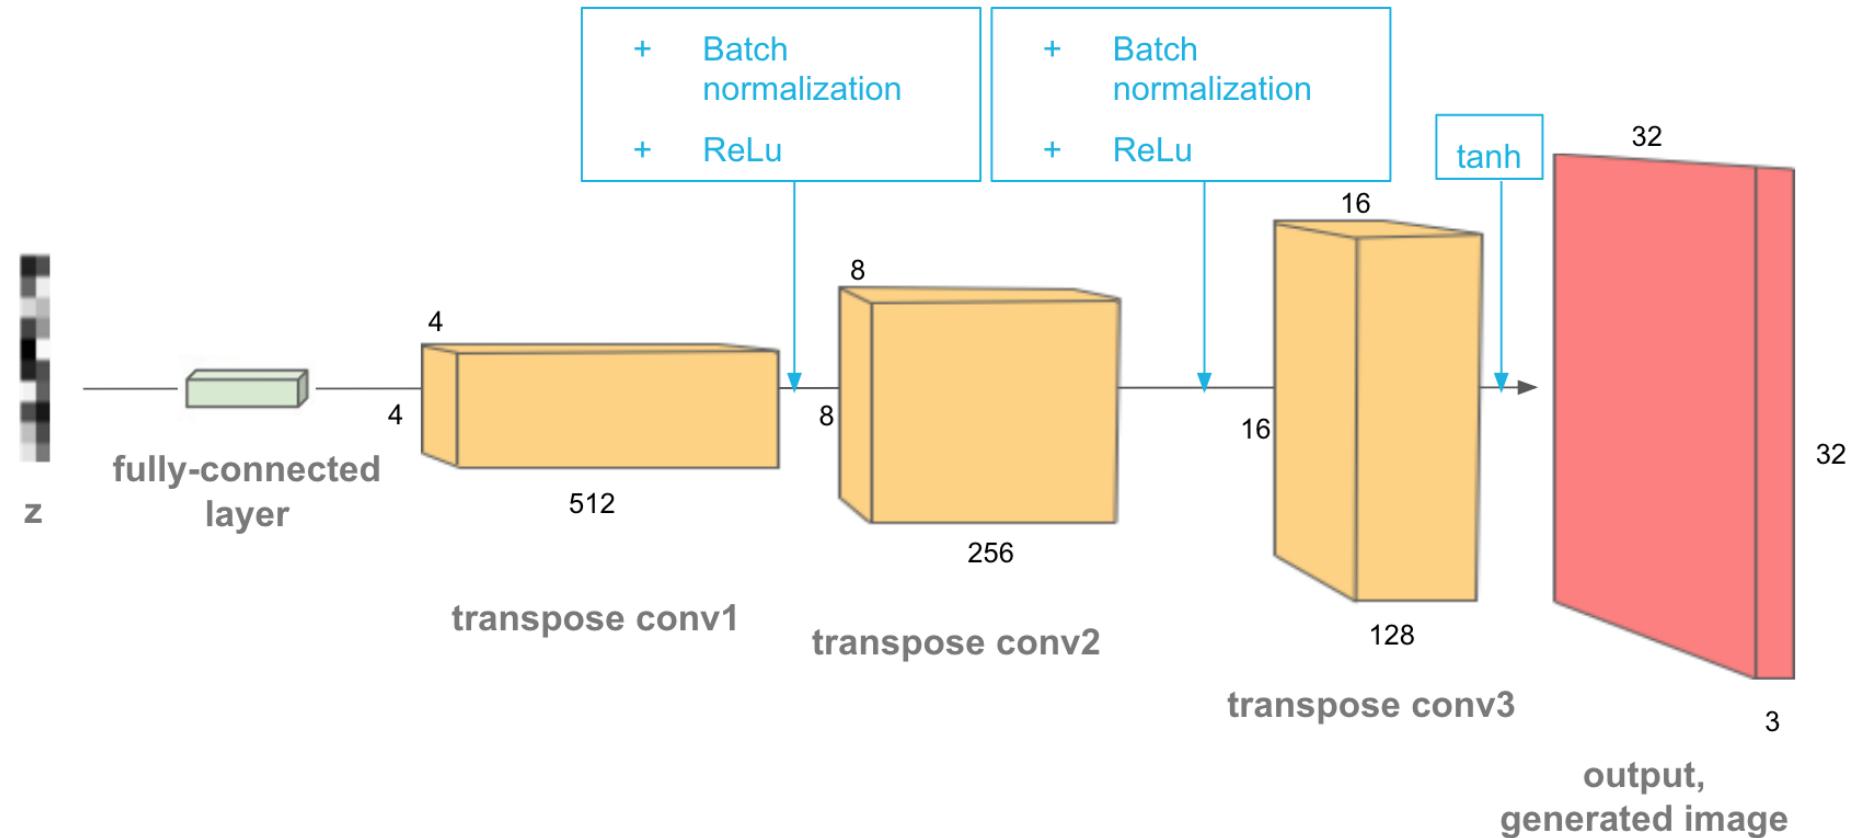
✓ convolution > batch norm > leaky ReLU.



DCGAN (Deep Convolutional GAN)

➤ The generator:

- ✓ transpose convolution > batch norm > ReLU.



Deep Convolutional GANs

- Images generated by the DCGAN after 50 epochs of training
 - ✓ Fashion MNIST



Figure 17-17. Images generated by the DCGAN after 50 epochs of training

DCGAN : Vector arithmetic for visual concepts

- DCGAN can learn quite meaningful latent representations,

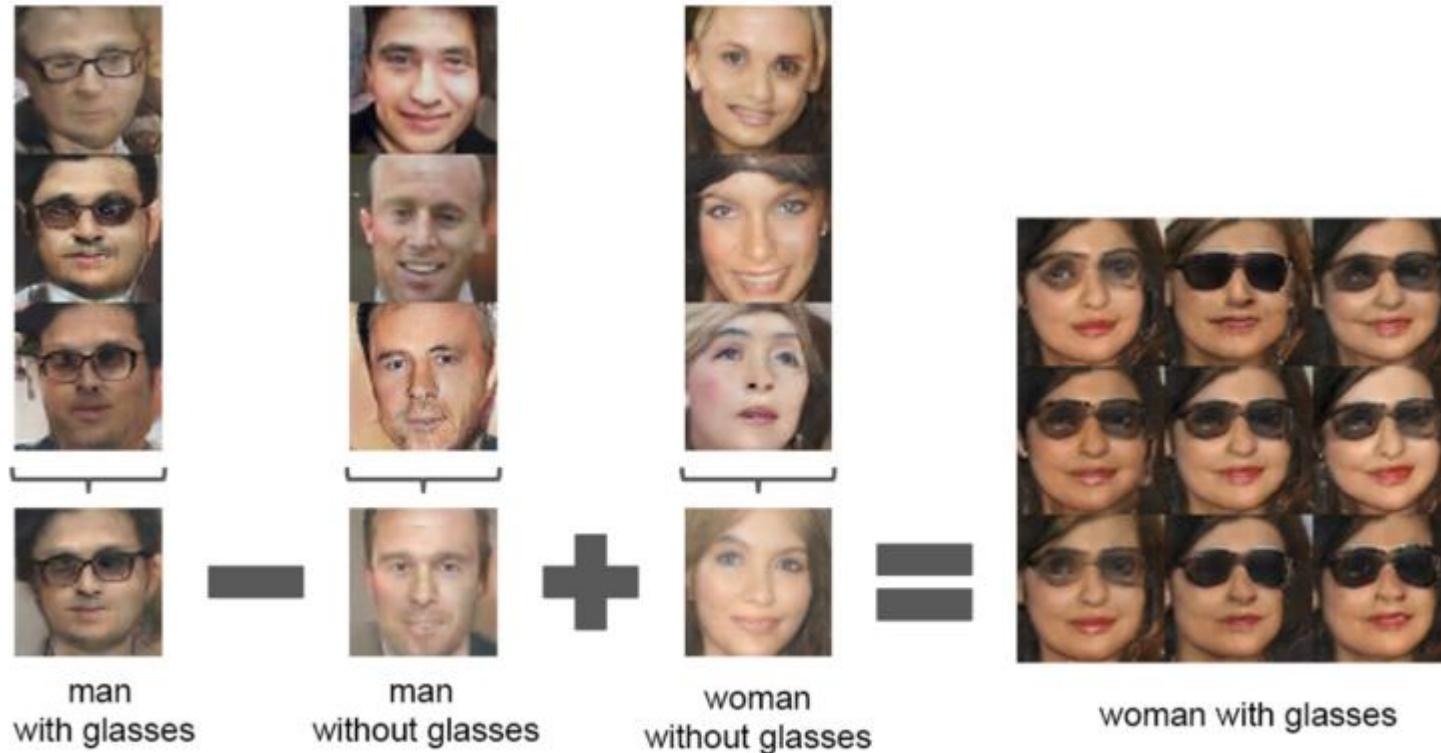
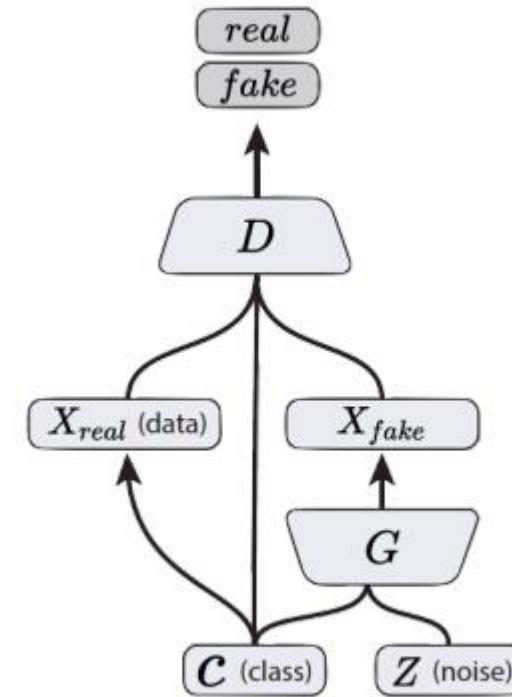


Figure 17-18. Vector arithmetic for visual concepts (part of figure 7 from the DCGAN paper)¹⁴

Conditional GANs

- Simple modification to the original GAN framework that conditions the model on additional information for better multi-modal learning
- Lends to many practical applications of GANs when we have explicit supervision available.



Conditional GAN
(Mirza & Osindero, 2014)

Conditional GANs

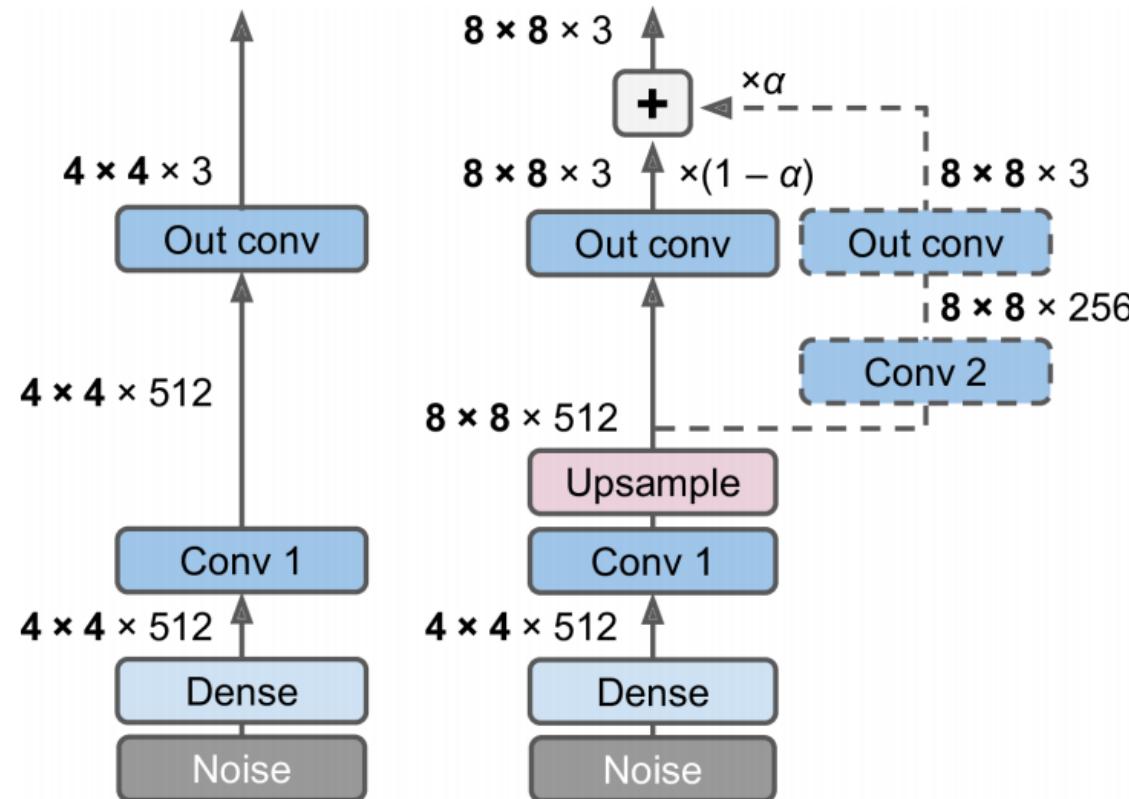
MNIST digits generated conditioned on their class label.



Figure 2 in the original paper.

Progressive Growing of GANs

- Progressive Growing of GANs: Nivida Tero Karras, in 2018 paper
 - ✓ a GAN generator outputs 4×4 color images (left)
 - ✓ extend it to output 8×8 images (right)



- StyleGAN: Nvidia team
 - ✓ the state of the art in high-resolution image generation in 2018
 - ✓ style transfer techniques in the generator
- Adaptive Instance Normalization (AdaIN) :
 - ✓ each noise layer is followed by an AdaIN

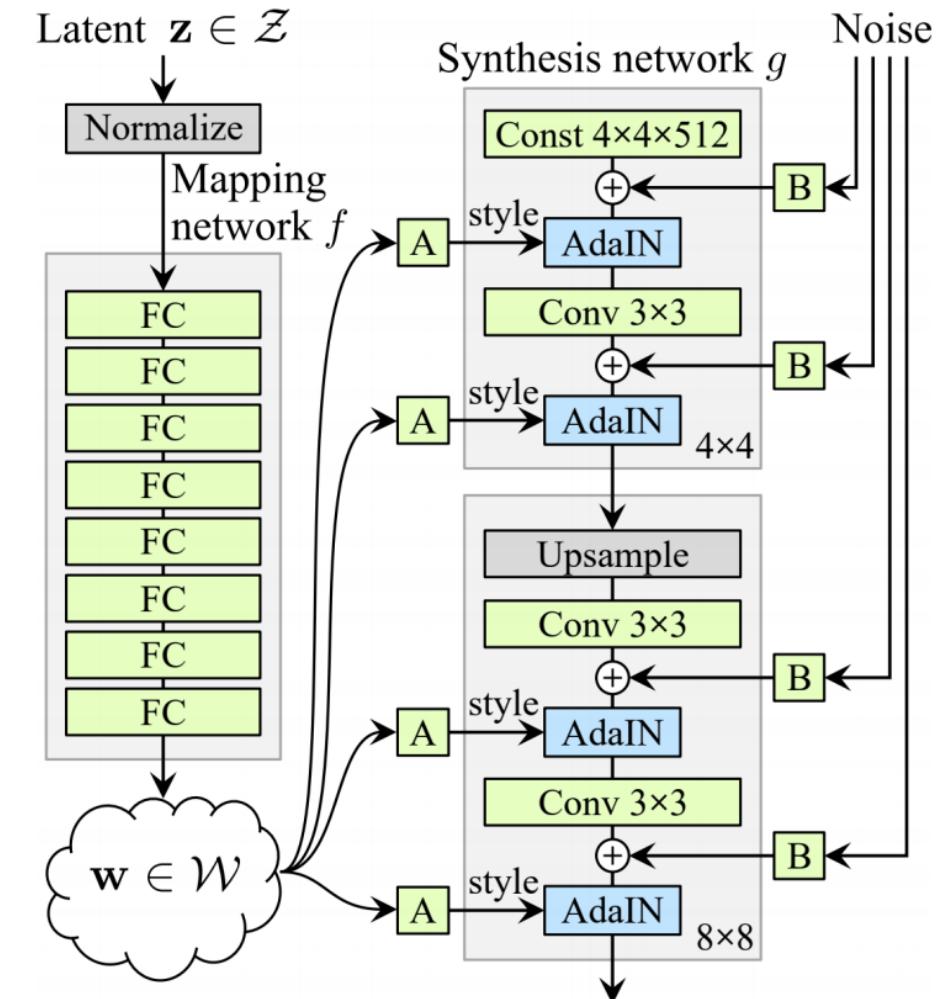


Figure 17-20. StyleGAN's generator architecture (part of figure 1 from the StyleGAN paper)¹⁹

Pix2pix

Labels to Street Scene

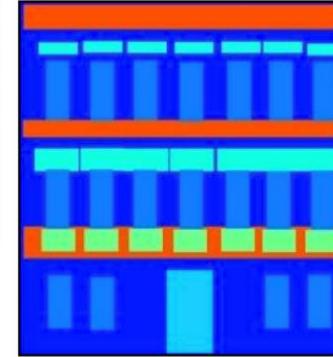


input



output

Labels to Facade



input



output

BW to Color

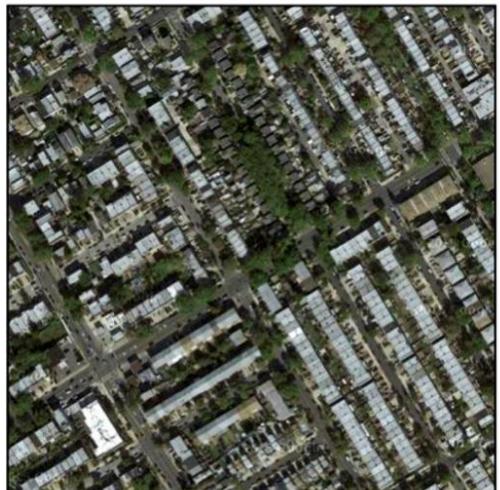


input

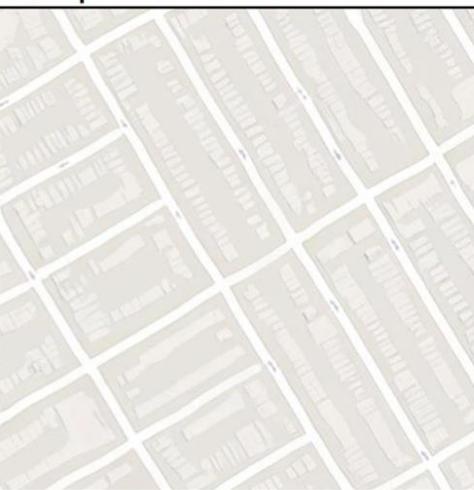


output

Aerial to Map



input



output

Day to Night



input



output

Edges to Photo



input



output

Image-to-Image Translation

- Architecture: *DCGAN-based* architecture
- Training is conditioned on the images from the source domain.
- Conditional GANs provide an effective way to handle many complex domains without worrying about designing *structured loss* functions explicitly.

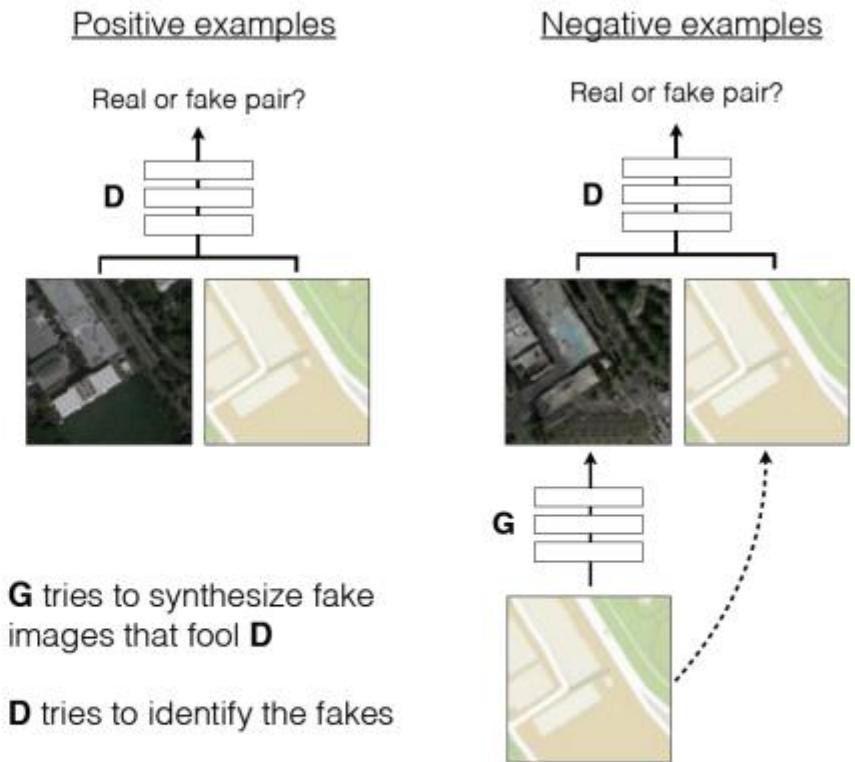
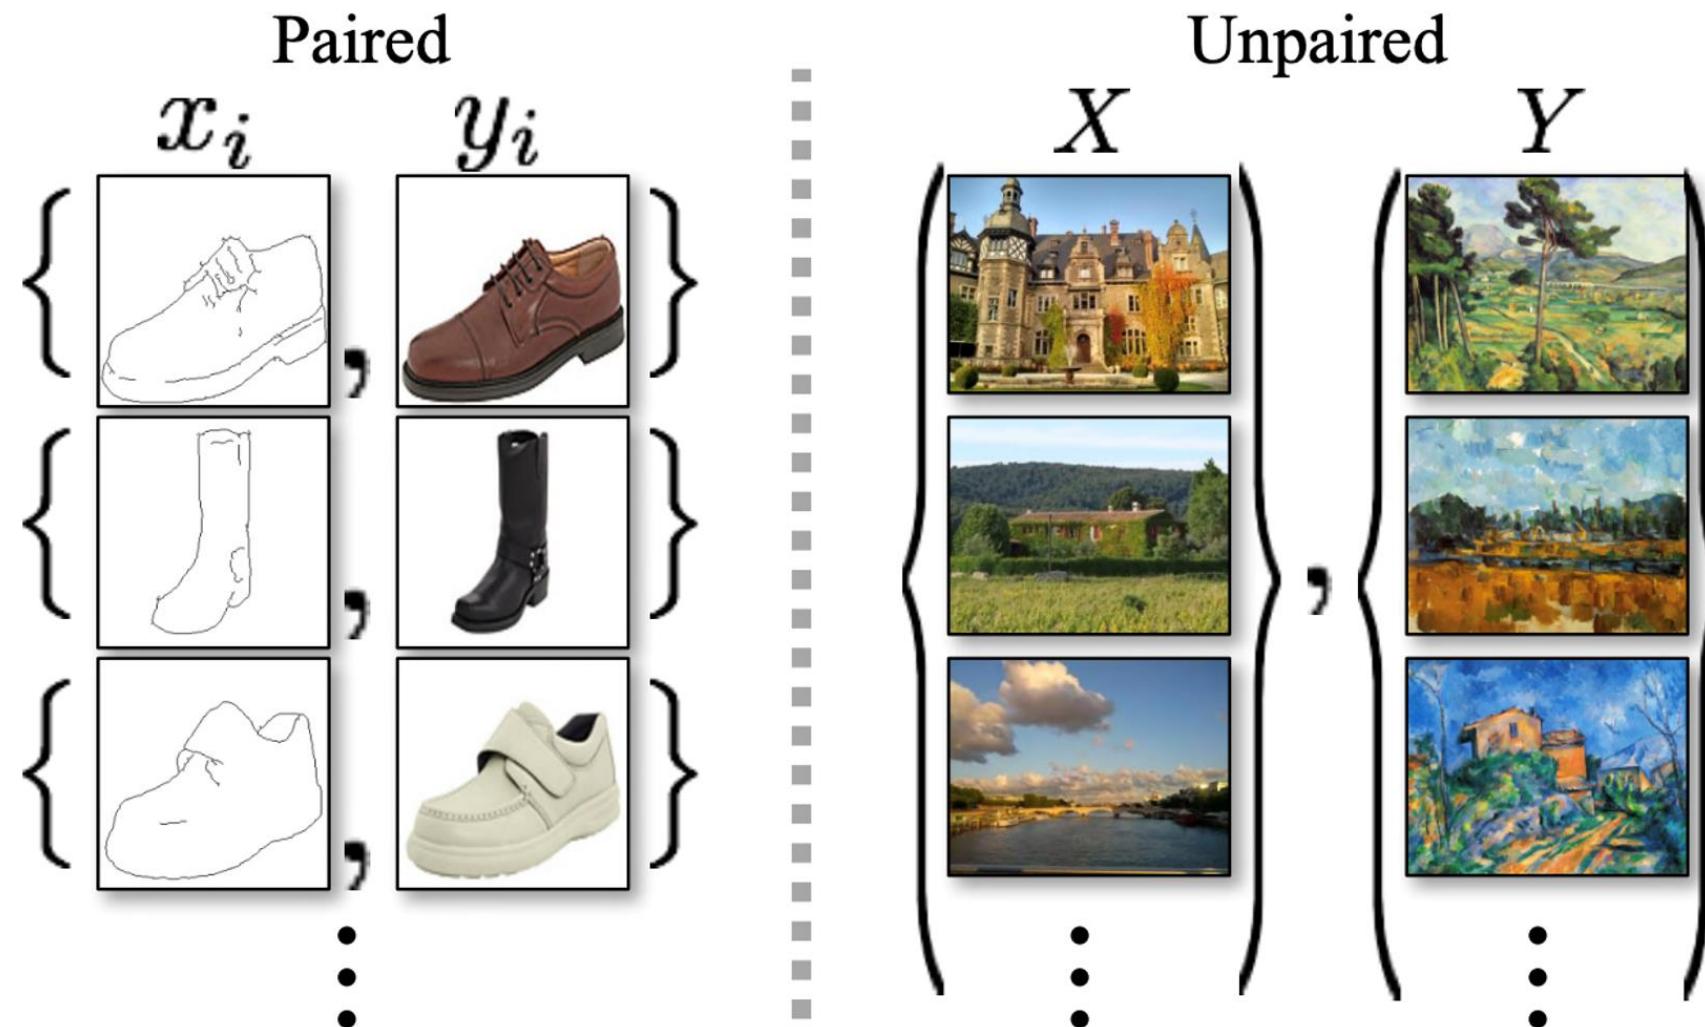


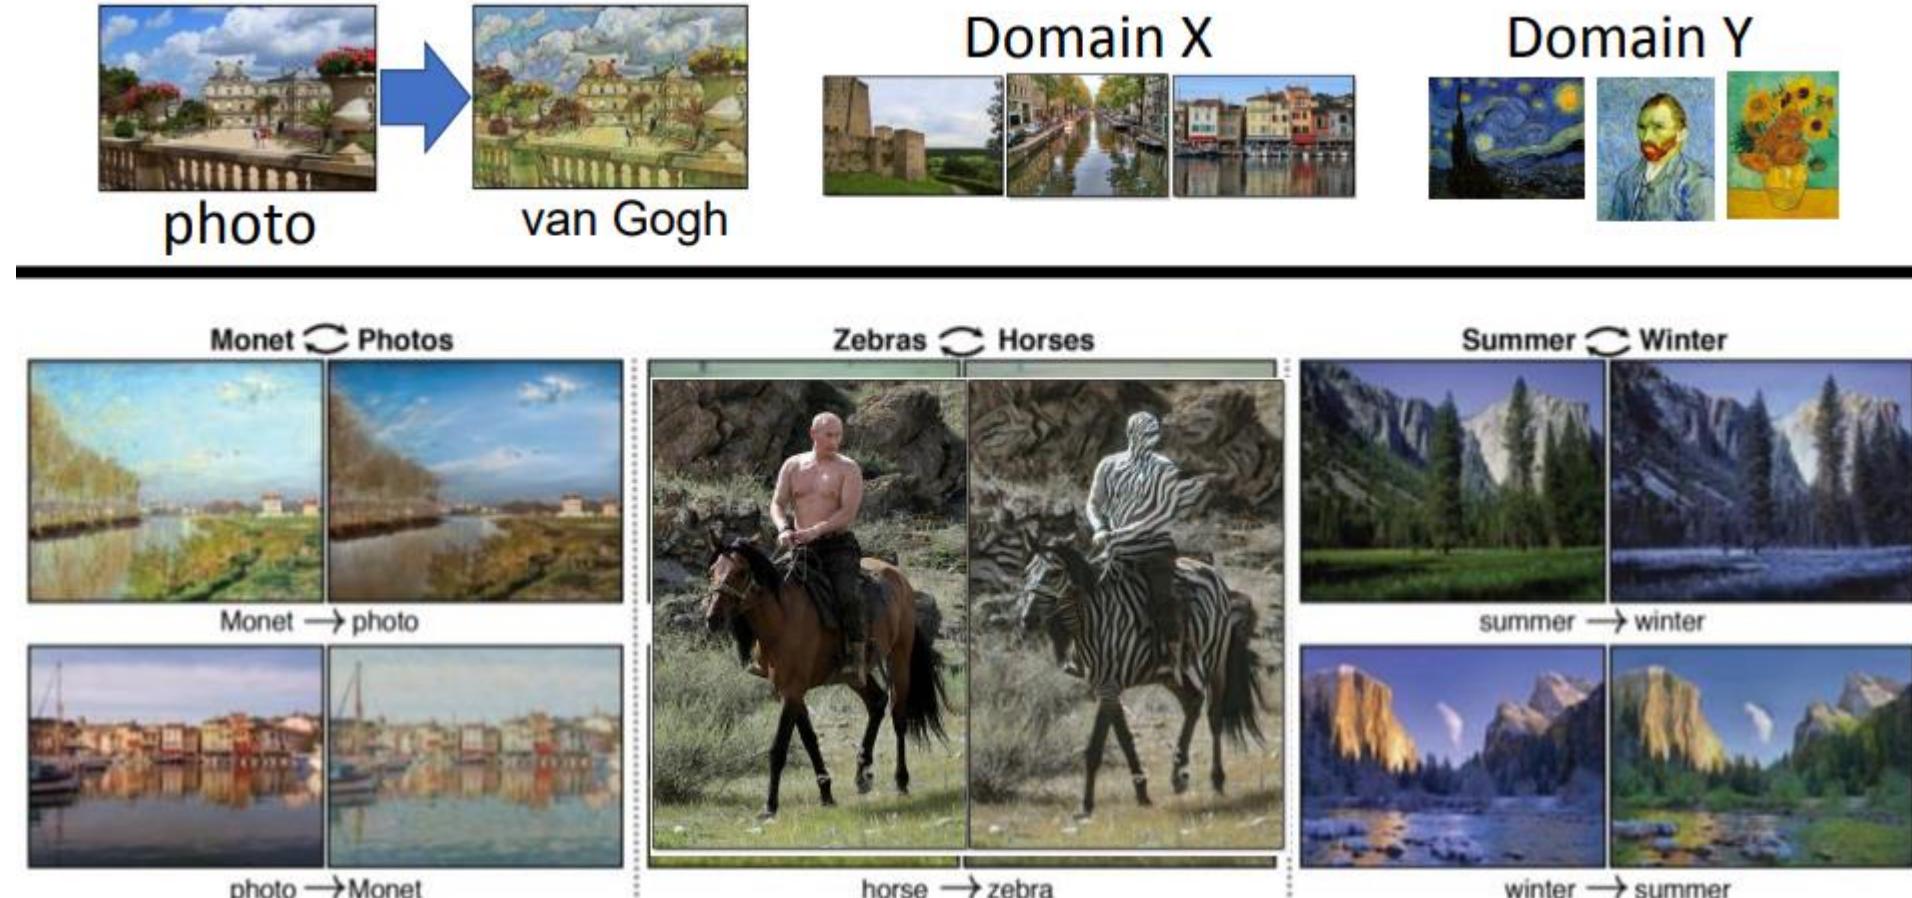
Figure 2 in the original paper.

Isola, P., Zhu, J. Y., Zhou, T., & Efros, A. A. "Image-to-image translation with conditional adversarial networks". arXiv preprint arXiv:1611.07004. (2016).



CycleGAN: domain transformation

- Transform an object from one domain to another without paired data



- "Photo-realistic single image super-resolution using a generative adversarial network."
 - ✓ Ledig, Christian, et al.
 - ✓ arXiv preprint arXiv:1609.04802 (2016).



2022

Korea Institute of Science
and Technology Information

TRUST
KISTI

