

2022

# Advanced Topic in Research Data-centric Deep Learning

## Lec 13: Generative Adversal Networks : GAN



hsyi@kisti.re.kr

Hongsuk Yi (이홍석)



# Reviewing the last class: AutoEncoder

- Machine Learning is to find a function  $f$
- Supervised
  - ✓ Data –  $X, y$
  - ✓ Goal – Learn mapping from  $X \rightarrow Y$
- Un-Supervised
  - ✓ Data –  $X$
  - ✓ Goal – Learn Hidden structure of data
- Output
  - ✓ Regression
  - ✓ Classification
  - ✓ Prediction/Structure Learning

$$f : X \rightarrow Y$$

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

## Maximum Likelihood Models

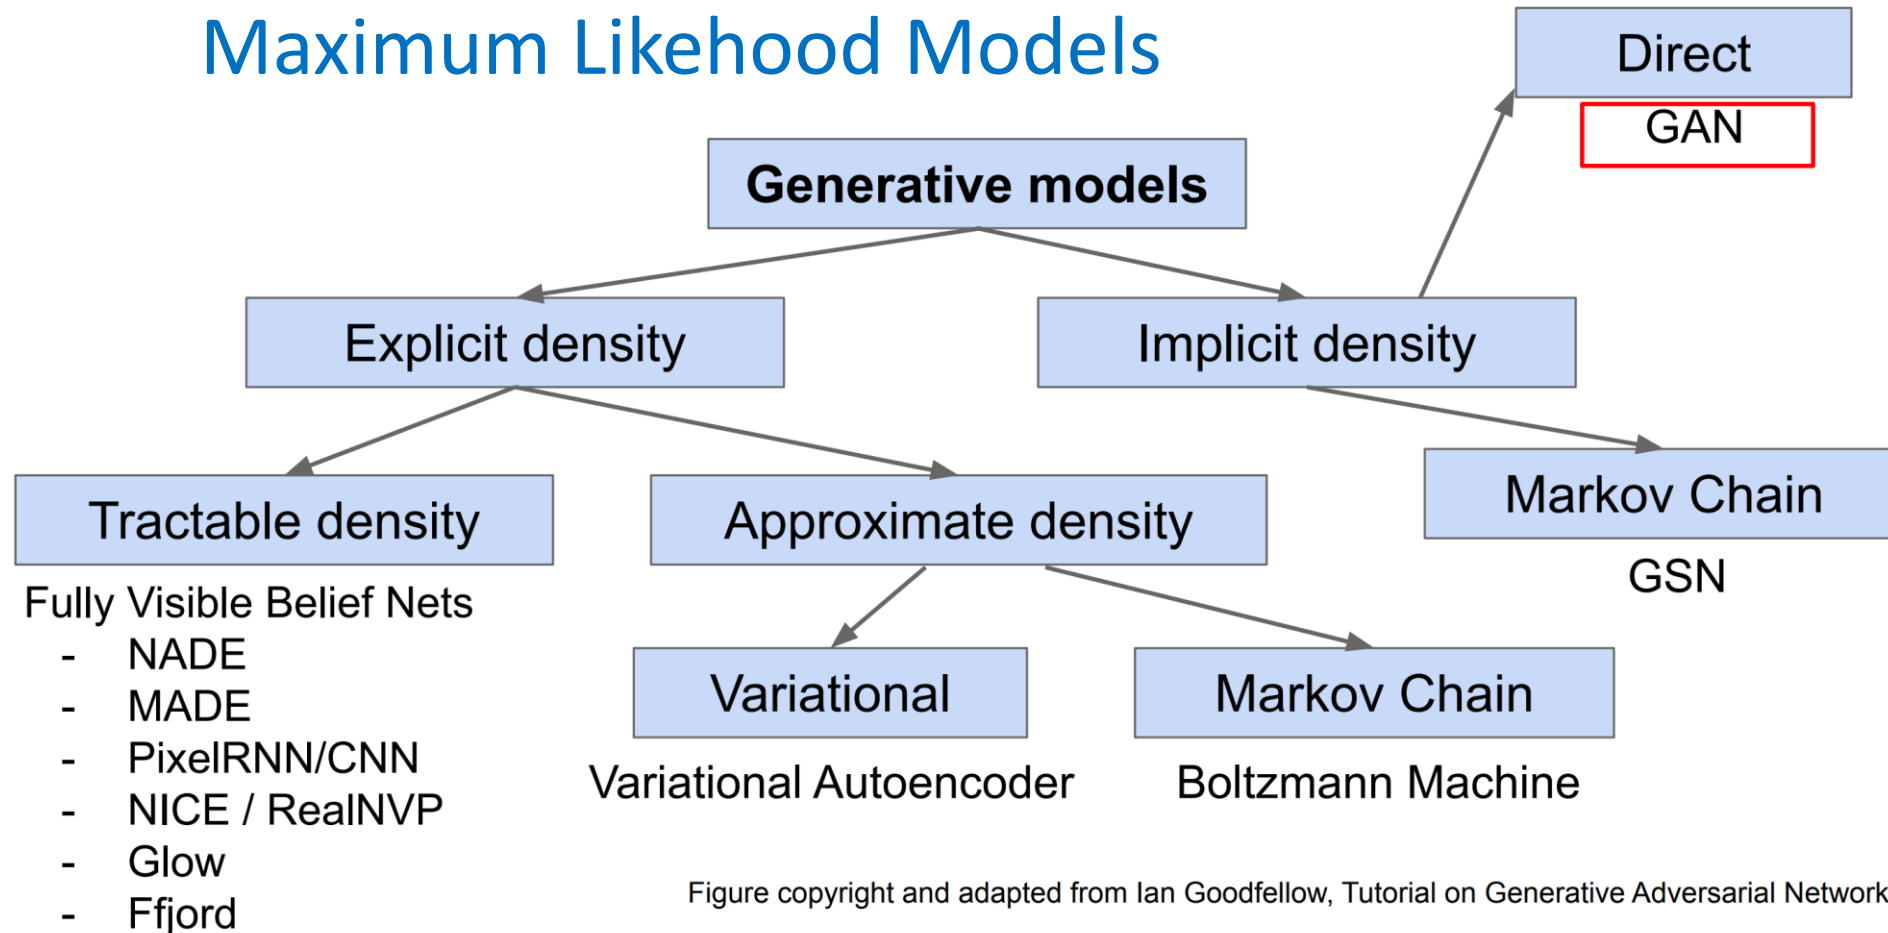


Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017.

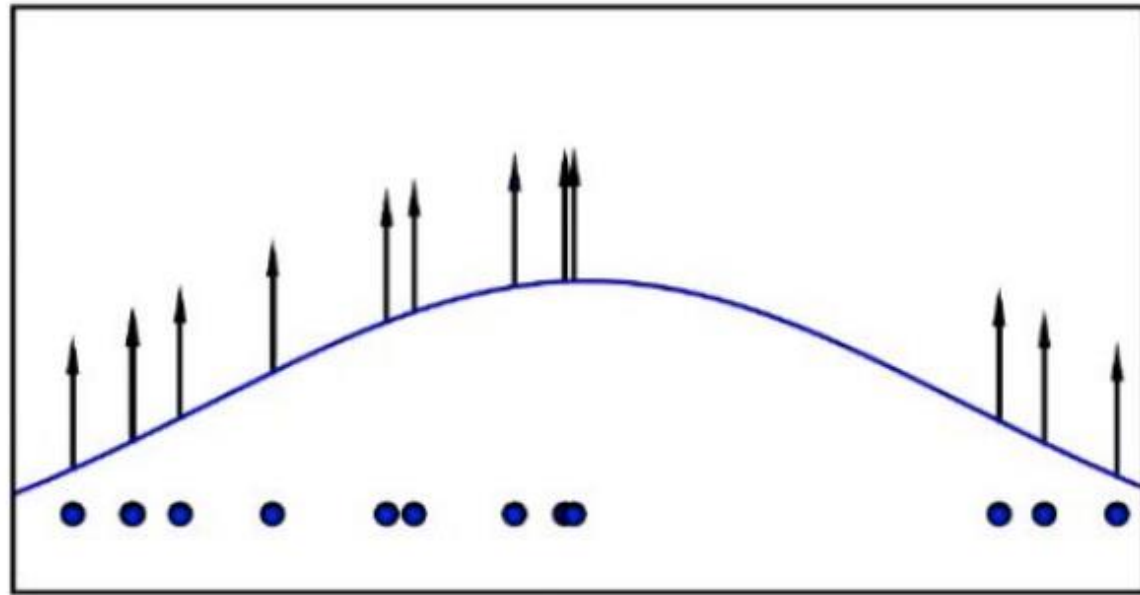
<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

- Maximum likelihood tries increase the likelihood of data given the parameters

$$P(x)$$

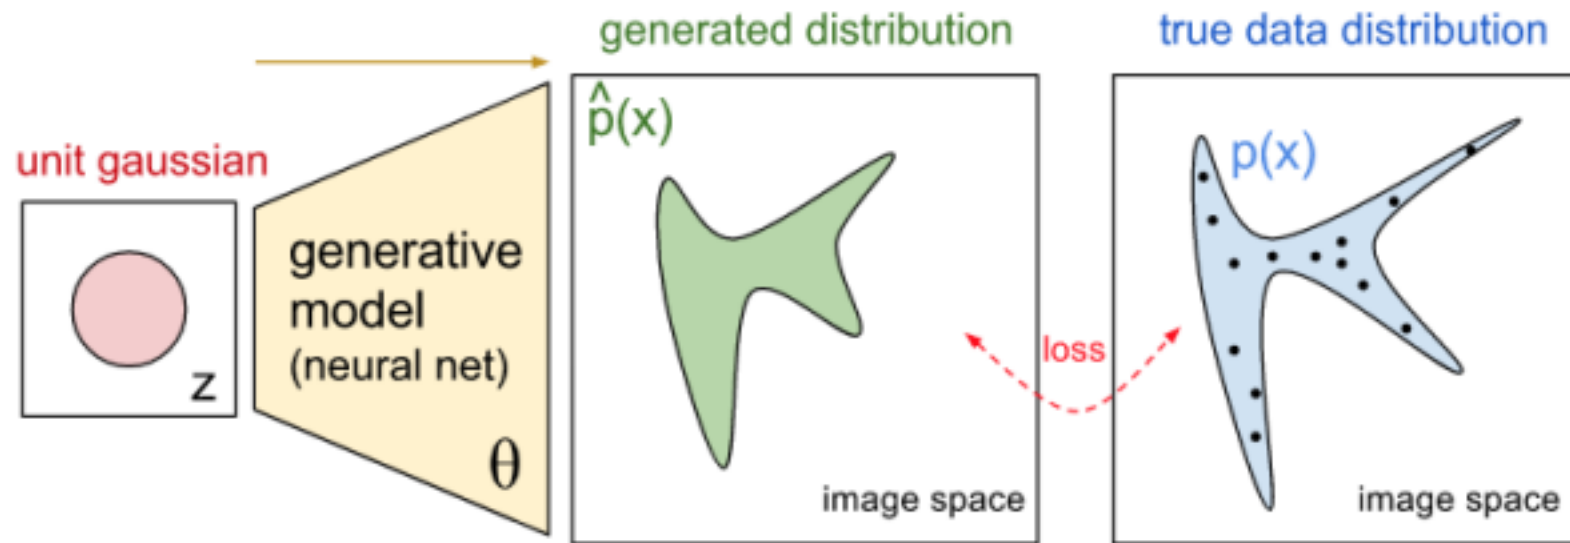
$$\theta^* = \arg \max_{\theta}$$

$$E_{x \sim P_{data}} \log P(x/\theta)$$



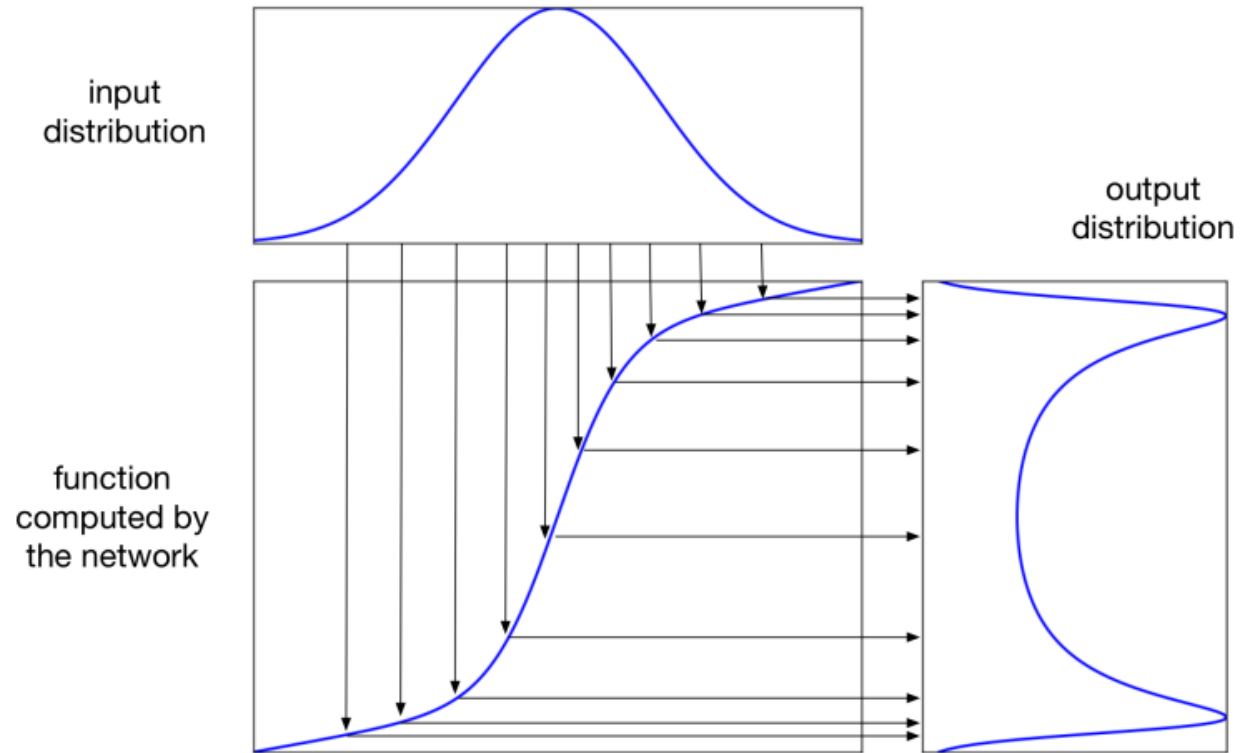
<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

- **Generative models** define a probability distribution
  - ✓ Start by sampling the **code vector  $z$**  from a fixed, simple distribution
  - ✓ The **generator** network computes a differentiable function  **$G$**  mapping  **$z$**  to an  **$x$**  in data space



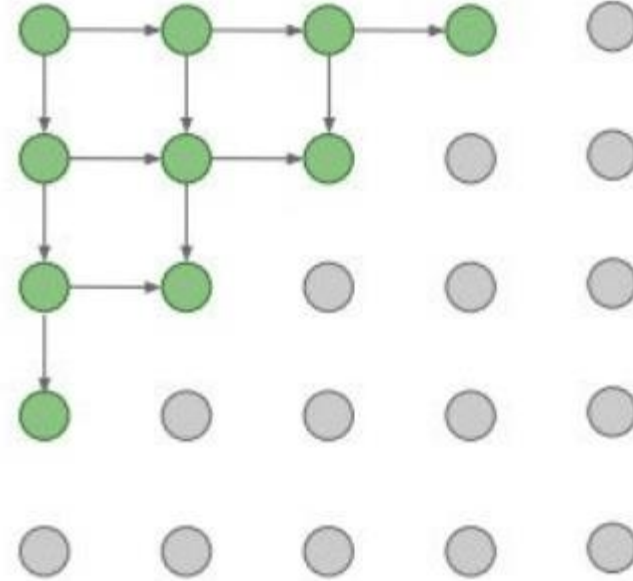


## ➤ 1Dimensional example



<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>

- Generate image pixels from corner
- Training Faster
- Generation Slow / Sequential
- Cannot generate samples based on some latent code



$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, \dots, x_{i-1})$$

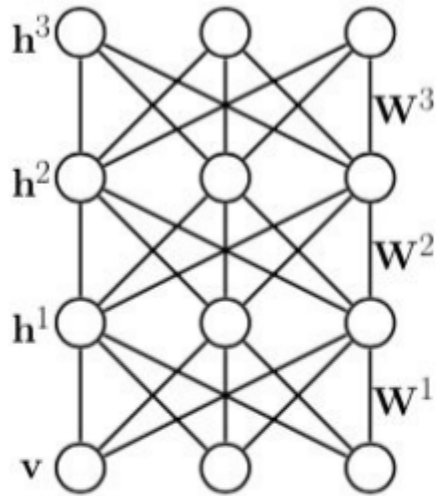
Chain Rule

Maximum Likelihood based Training

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>



## Boltzmann Machine

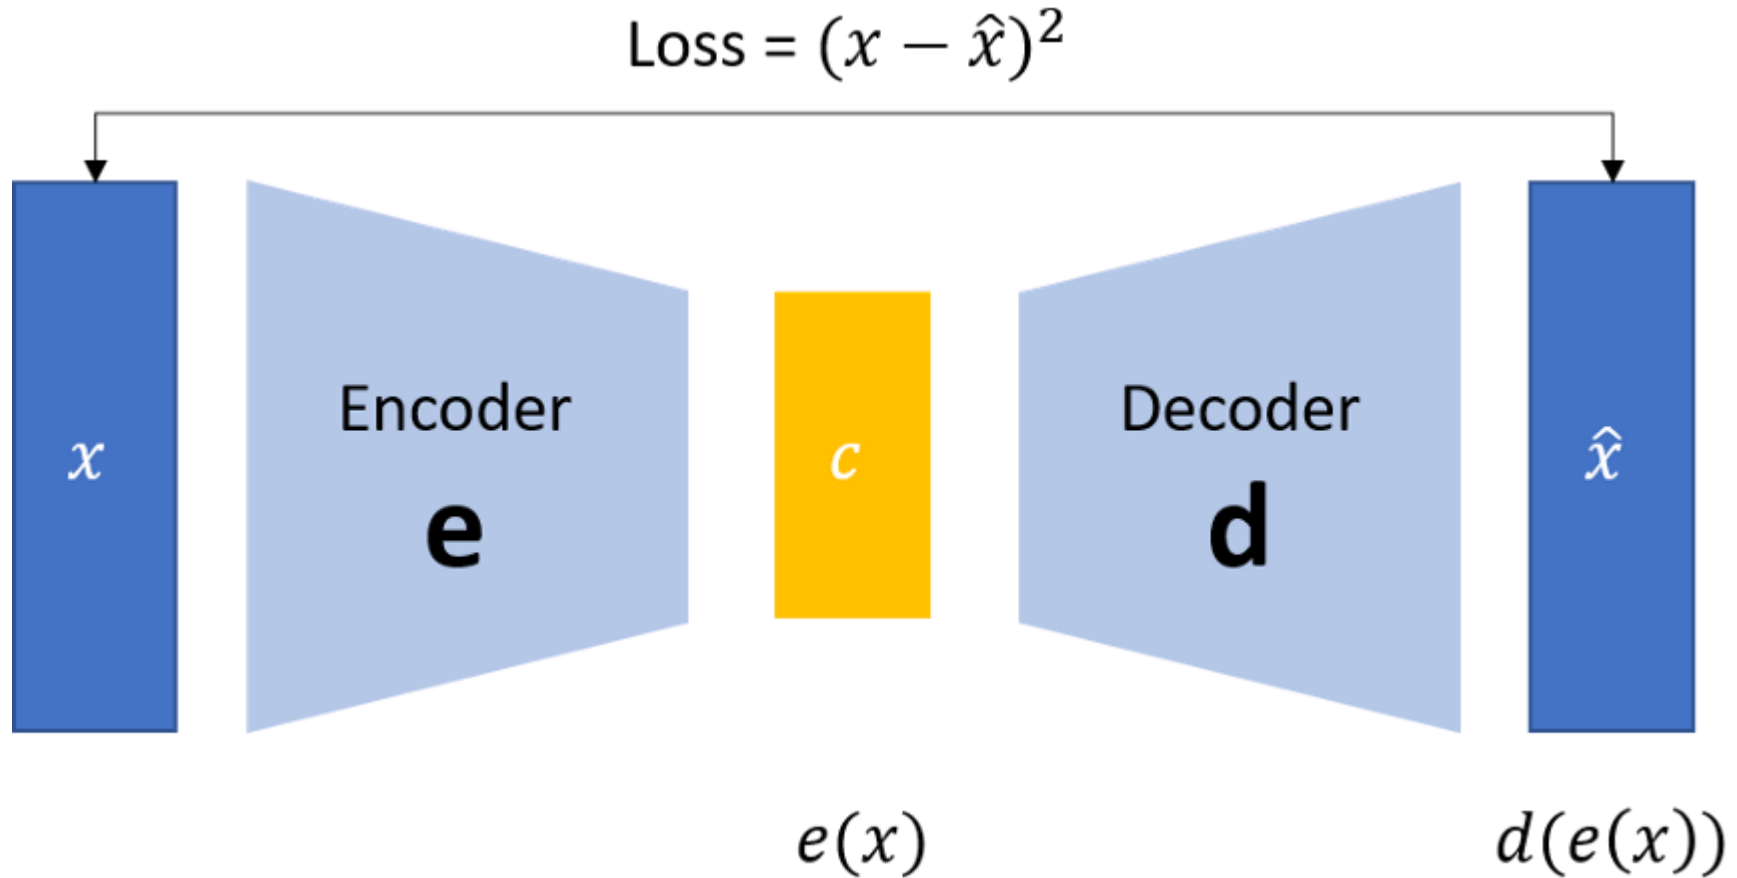


- Energy Function based models
- Markov chains don't work for long sequences
- Hard to scale on large dataset

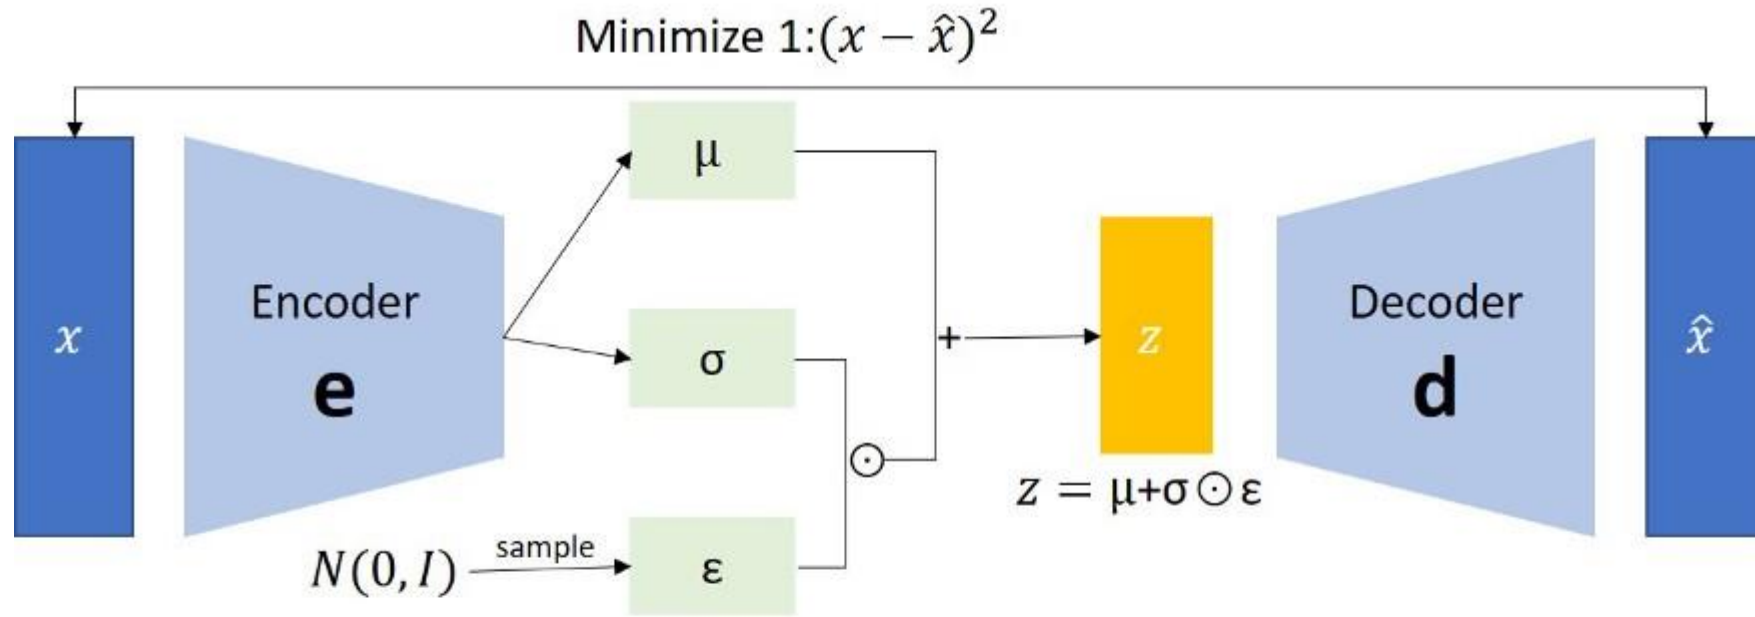
$$p(x, h) = \exp(-E(x, h)) / Z$$

$$Z = \sum_{x, h} \exp(-E(x, h))$$

<source> <https://people.cs.pitt.edu/~milos/courses/cs3750/lectures/class23.pdf>



<https://medium.com/geekculture/variational-autoencoder-vae-9b8ce5475f68>



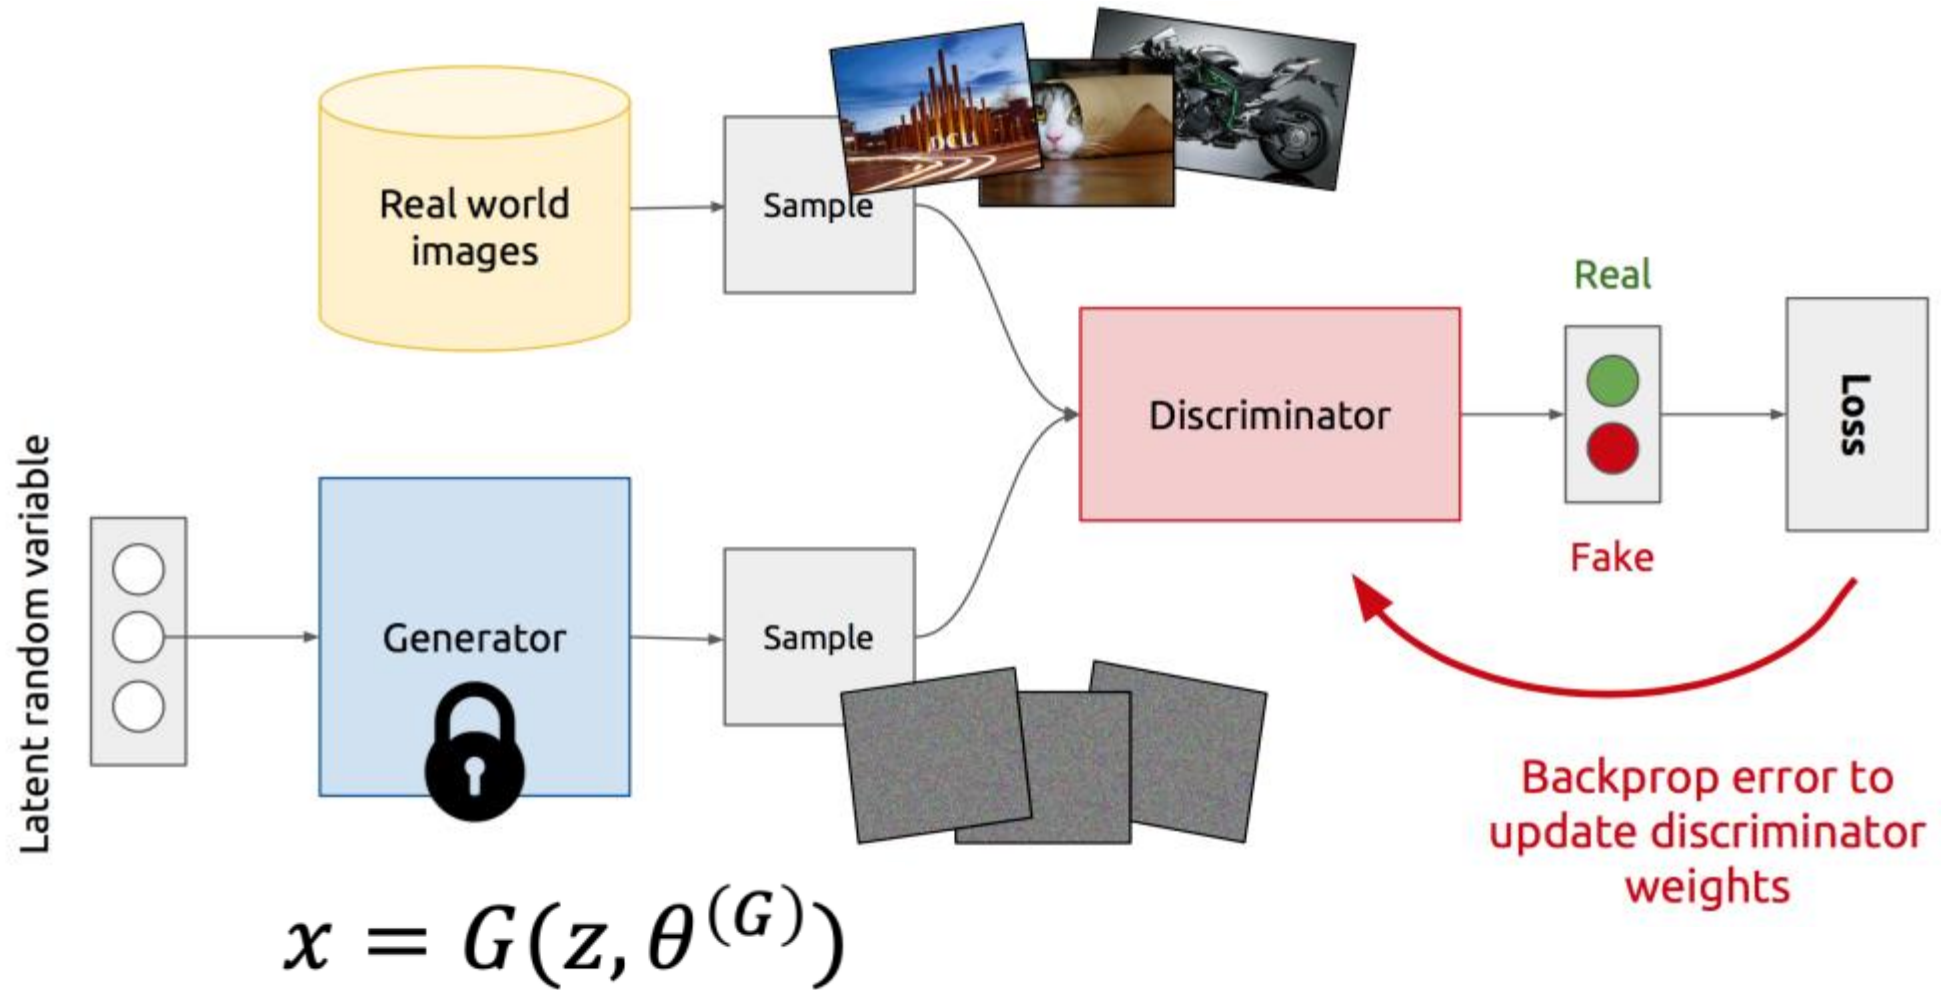
Minimize 2:  $\frac{1}{2} \sum_{i=1}^N (\exp(\sigma_i) - (1 + \sigma_i) + \mu_i^2)$  reconstruction error

**Reparameterization trick**  $z = \mu + \sigma \odot \epsilon$

# Generative Adversarial Networks (GANs)

- Plenty of existing work on Deep Generative Models
  - ✓ Boltzmann Machine
  - ✓ Deep Belief Nets
  - ✓ Variational AutoEncoders (VAE)
  
- Why GANs?
  - ✓ Sampling (or generation) is straightforward
  - ✓ Training doesn't involve Maximum Likelihood estimation
  - ✓ Robust to Overfitting since Generator never sees the training data
  - ✓ Empirically, GANs are good at capturing the modes of the distribution.

## ➤ Classic GAN Framework



- GANs extend that idea to generative models:
  - ✓ generate adversarial samples to fool a discriminative model
  - ✓ use those adversarial samples to make models robust
  - ✓ Repeat this and we get better discriminative model
- Generator:
  - ✓ generate fake samples, tries to fool the Discriminator
- Discriminator:
  - ✓ tries to distinguish between real and fake samples
  - ✓ Train them against each other
  - ✓ Repeat this and we get better Generator and Discriminator



- D tries to identify real data from fakes created by the generator
- G tries to create imitations of data to trick the discriminator
- Objective function:

Train GAN jointly via **minimax** game:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log \left( 1 - D_{\theta_d} \left( G_{\theta_g}(z) \right) \right) \right]$$

**Discriminator** wants to maximize objective s.t.  $D(x)$  close to 1,  $D(G(z))$  close to 0.

**Generator** wants to minimize objective s.t.  $D(G(z))$  close to 1.

- Discriminator is a function  $D$  (network, can deep)

$$D: X \rightarrow \mathbb{R}$$

- Input  $x$ : an object  $x$  (e.g. an image)
- Output  $D(x)$ : scalar which represents how “good” an object  $x$  is



Can we use the discriminator to generate objects?

Yes.

- Suppose we already have a good discriminator  $D(x)$  ...

## Inference

- Generate object  $\tilde{x}$  that

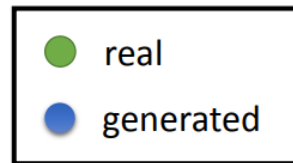
$$\tilde{x} = \arg \max_{x \in X} D(x)$$

Enumerate all possible  $x$  !!!

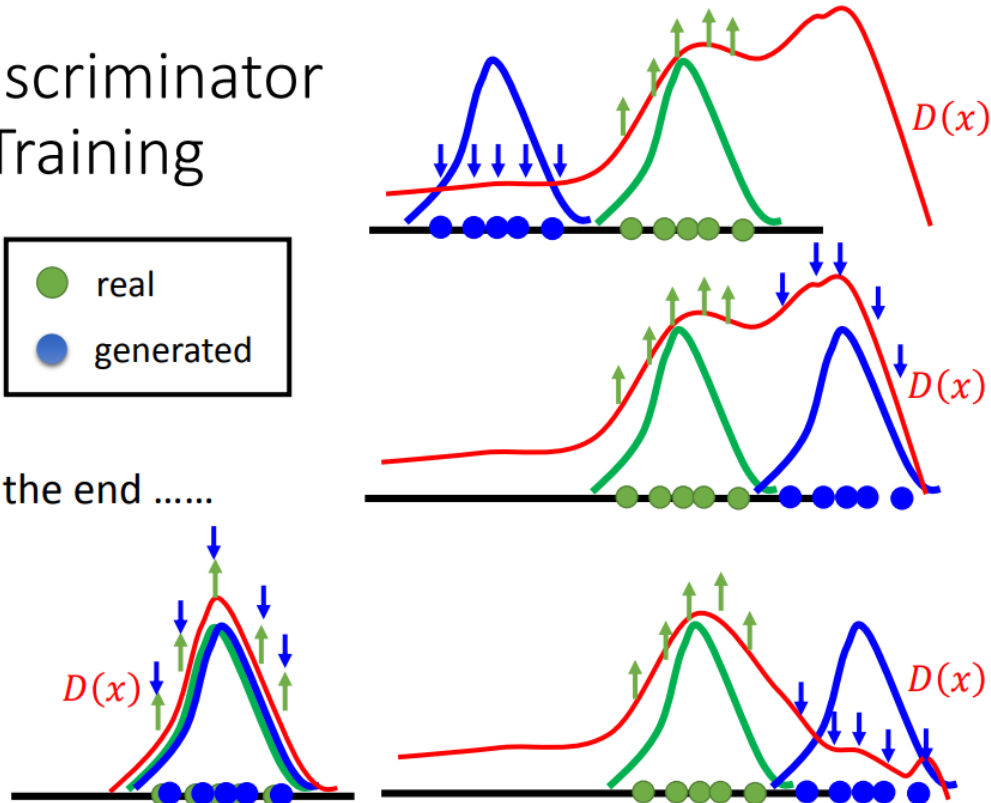
It is feasible ???

How to learn the discriminator?

## Discriminator - Training



In the end .....



### General Algorithm

- Given a set of **positive examples**, randomly generate a set of **negative examples**.



- In each iteration



- Learn a discriminator D that can discriminate positive and negative examples.



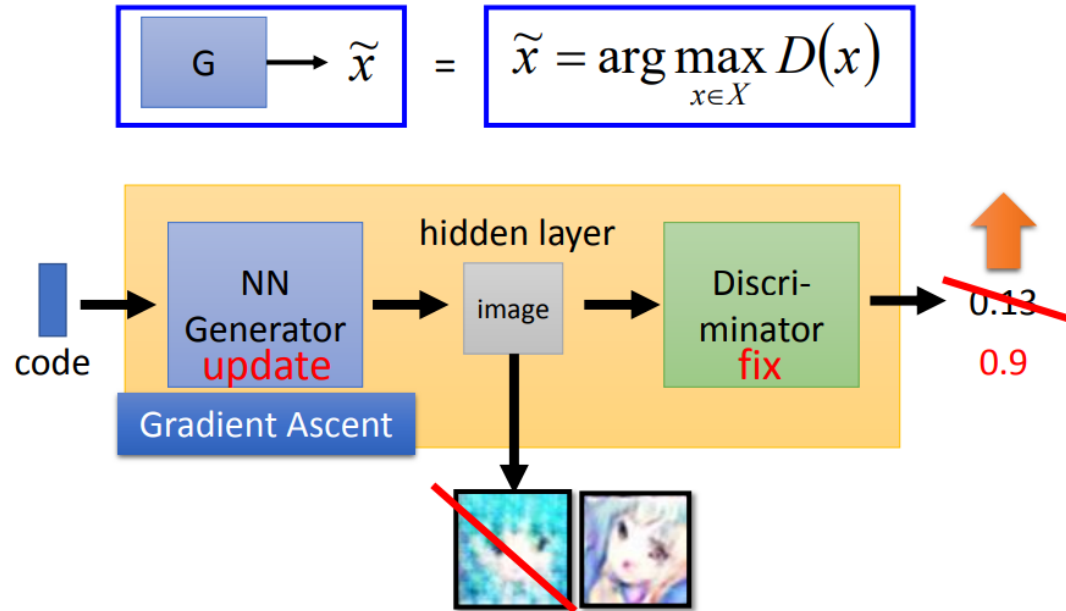
v.s.



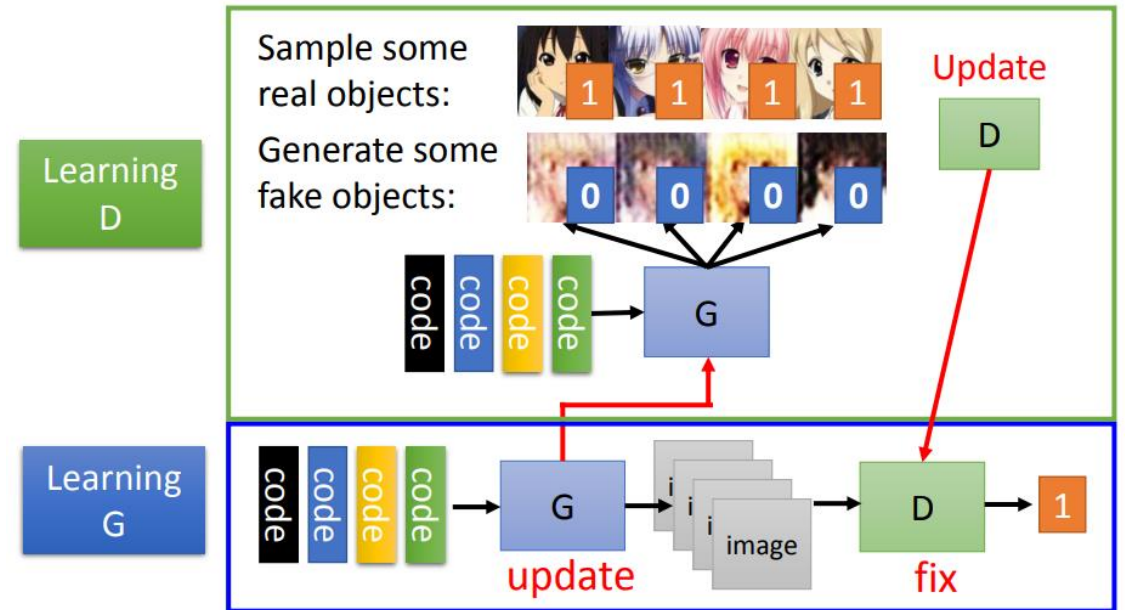
D

- Generate negative examples by discriminator D

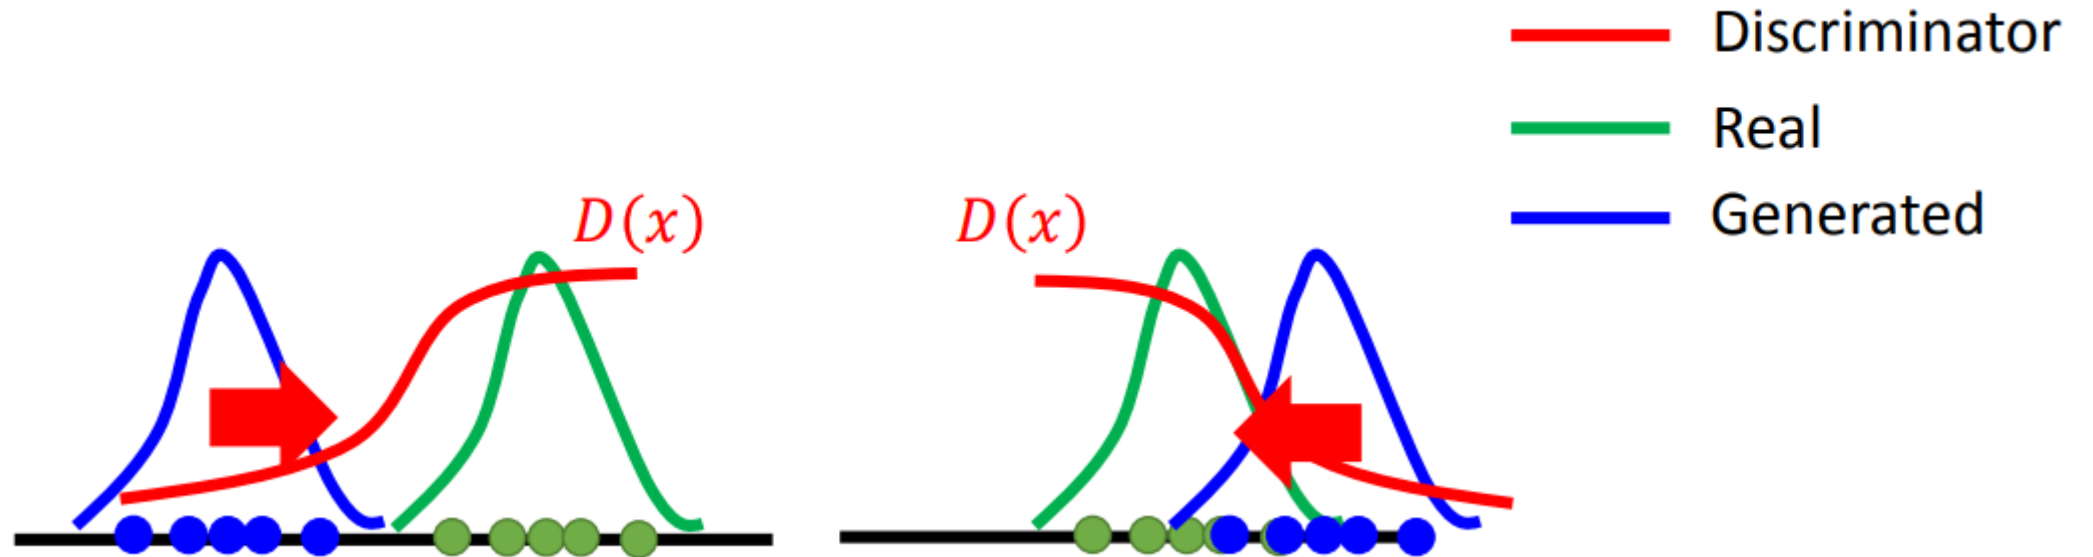
$$\boxed{G \rightarrow \tilde{x}} = \boxed{\tilde{x} = \arg \max_{x \in X} D(x)}$$



- Initialize generator and discriminator **G** **D**
- In each training iteration:



- Discriminator leads the generator

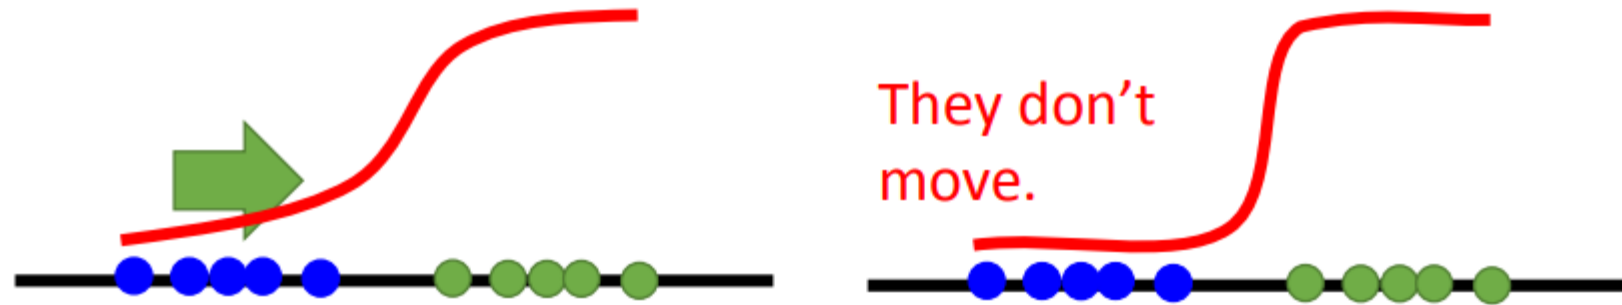




Typical binary classifier uses sigmoid function at the output layer

1 is the largest, 0 is the smallest

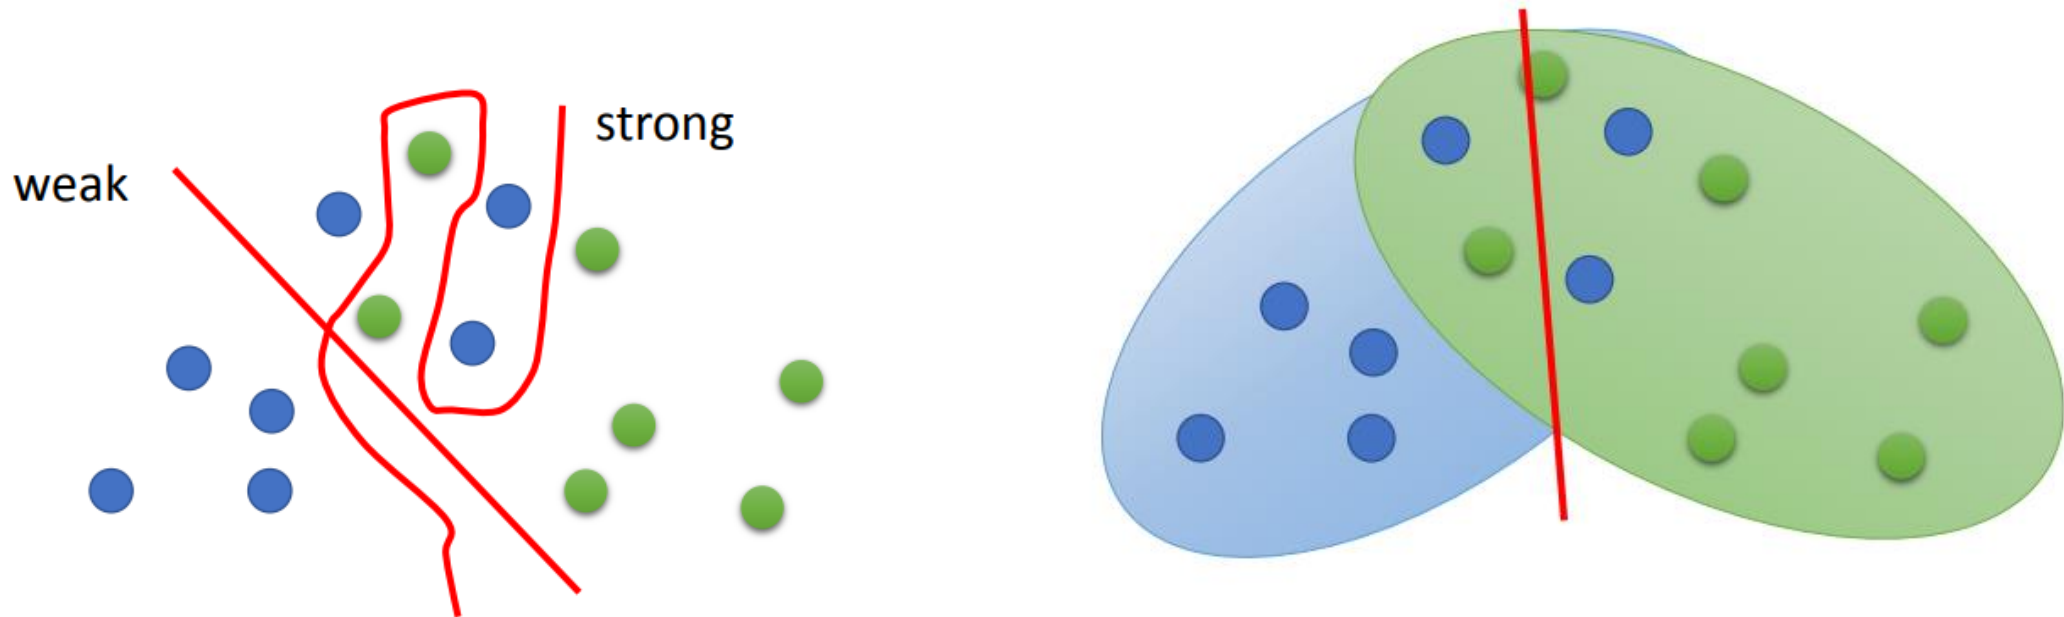
● real  
● generated



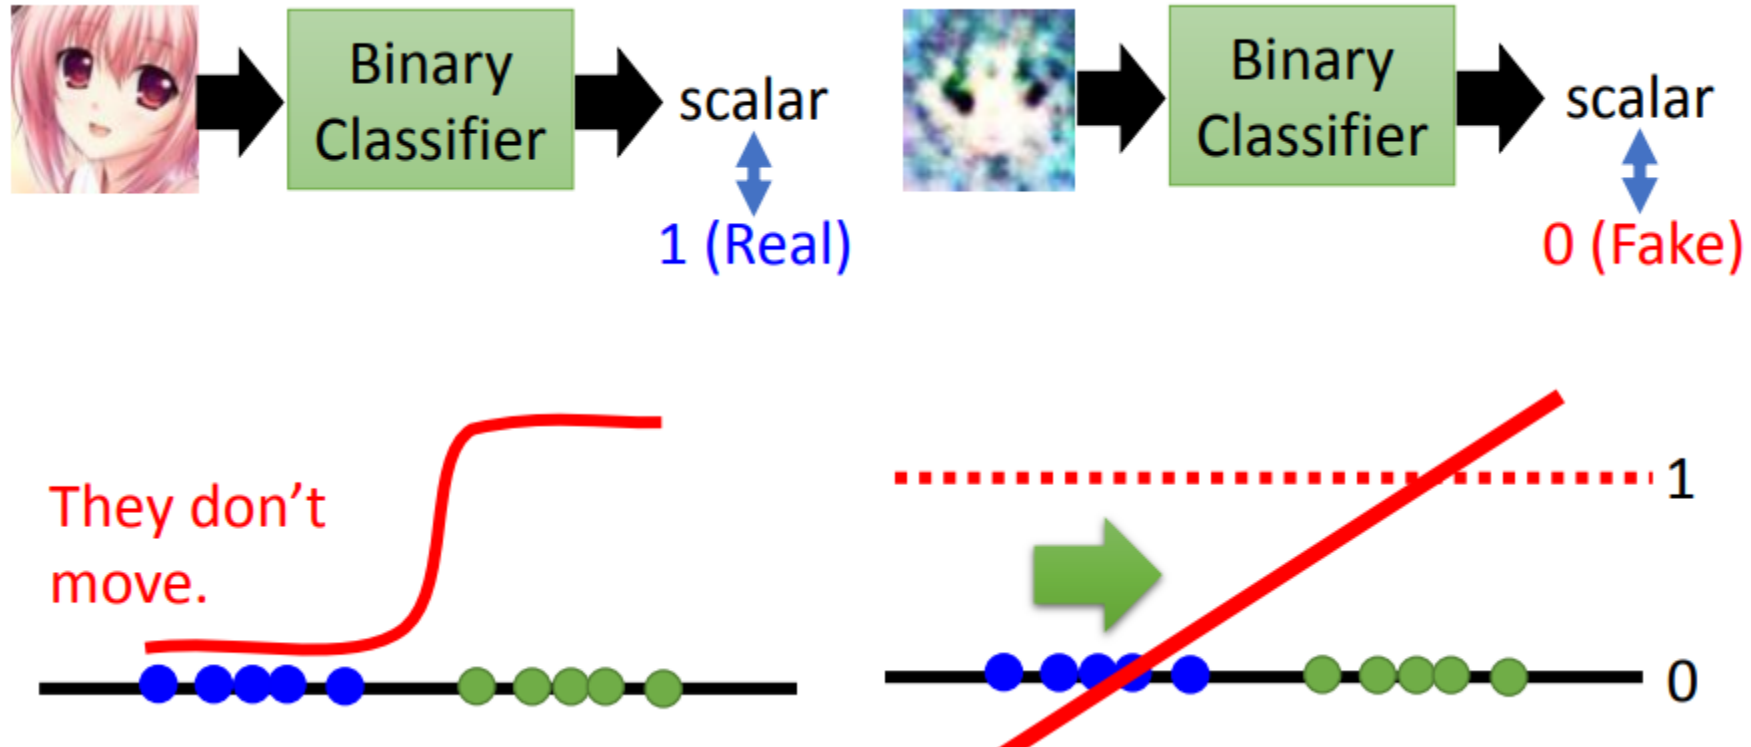
You cannot train your classifier *too good*.....



- Don't let the discriminator perfectly separate real and generated data
  - ✓ Add noise to input or label?



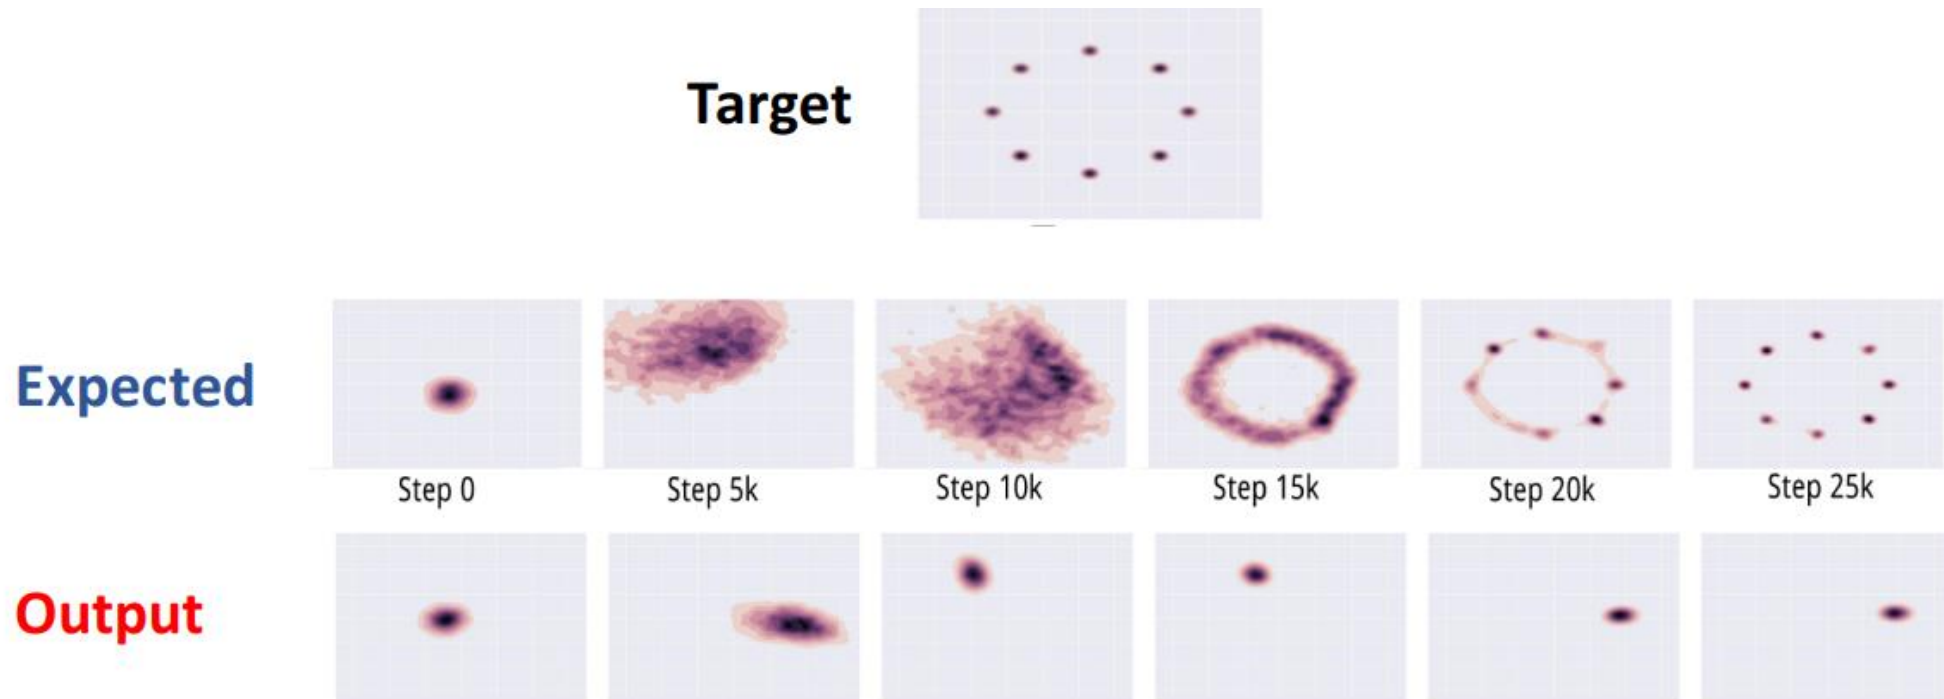
- Replace sigmoid with linear (replace classification with regression)



- zero-sum game:
  - ✓ During training, the generator and the discriminator constantly try to outsmart each other
- Nash equilibrium:
  - ✓ no player would be better off changing their own strategy, assuming the other players do not change theirs
- The biggest difficulty : mode collapse
  - ✓ the generator's outputs gradually become less diverse
- GANs are very sensitive to the hyperparameters:
  - ✓ you may have to spend a lot of effort fine-tuning them.

- Experience replay to avoid the mode collapse issue:
  - ✓ Google 2018 paper.
  - ✓ storing the images produced by the generator at each iteration in a replay buffer
- Mini-batch discrimination:
  - ✓ how similar images are across the batch and provides this statistic to the discriminator
  - ✓ it can easily reject a whole batch of fake images that lack diversity

- Generator fails to output diverse samples



- Basic Solutions:
  - ✓ Mini-Batch GANs

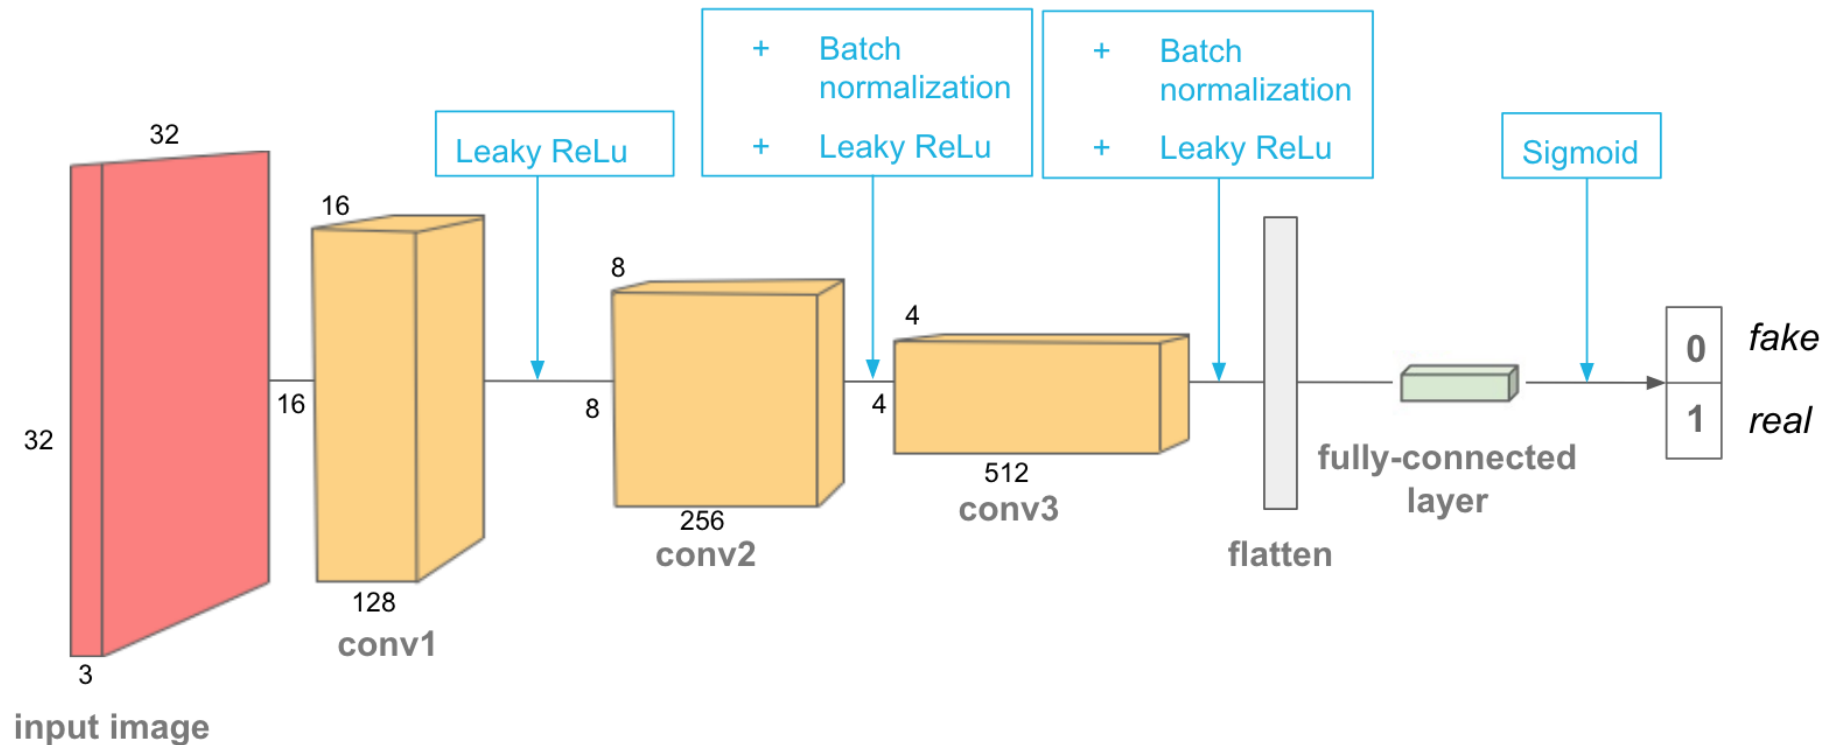
# More Recent GANs

- The original GAN paper in 2014 experimented with convolutional layers:
  - ✓ But, only tried to generate small images
- DCGAN : Deep Convolutional GANs: 2015, Alec Radford
  - ✓ Replace any pooling layers with strided convolutions (in the discriminator) and transposed convolutions (in the generator).
  - ✓ Use Batch Normalization in both the generator and the discriminator, except in the generator's output layer and the discriminator's input layer.
  - ✓ Remove fully connected hidden layers for deeper architectures.
  - ✓ Use ReLU activation in the generator for all layers except the output layer, which should use tanh.
  - ✓ Use leaky ReLU activation in the discriminator for all layers.



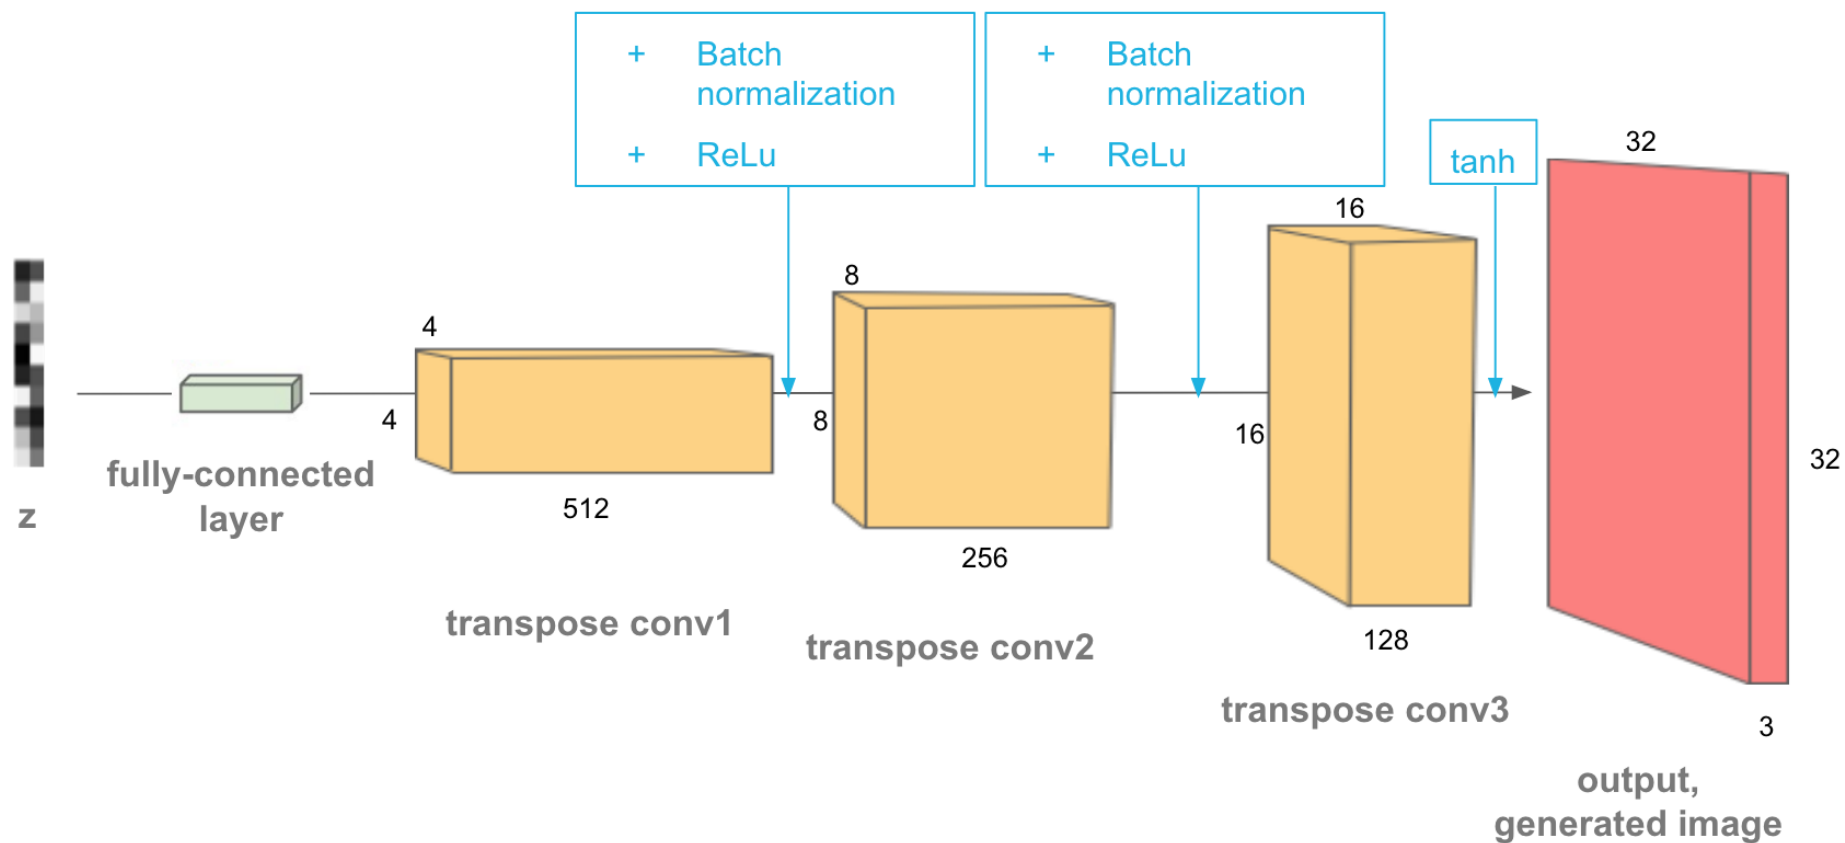
## ➤ discriminator:

✓ convolution > batch norm > leaky ReLU.



## ➤ The generator:

✓ transpose convolution > batch norm > ReLU.

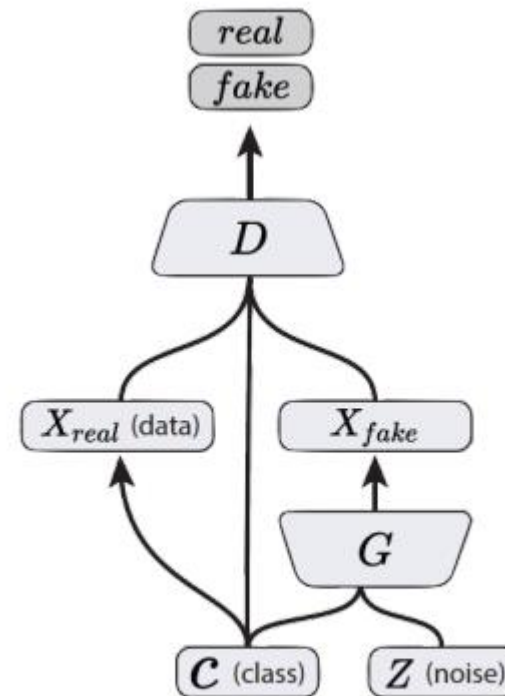


- Images generated by the DCGAN after 50 epochs of training
  - ✓ Fashion MNIST



Figure 17-17. Images generated by the DCGAN after 50 epochs of training

- Simple modification to the original GAN framework that conditions the model on additional information for better multi-modal learning
- Lends to many practical applications of GANs when we have explicit supervision available.



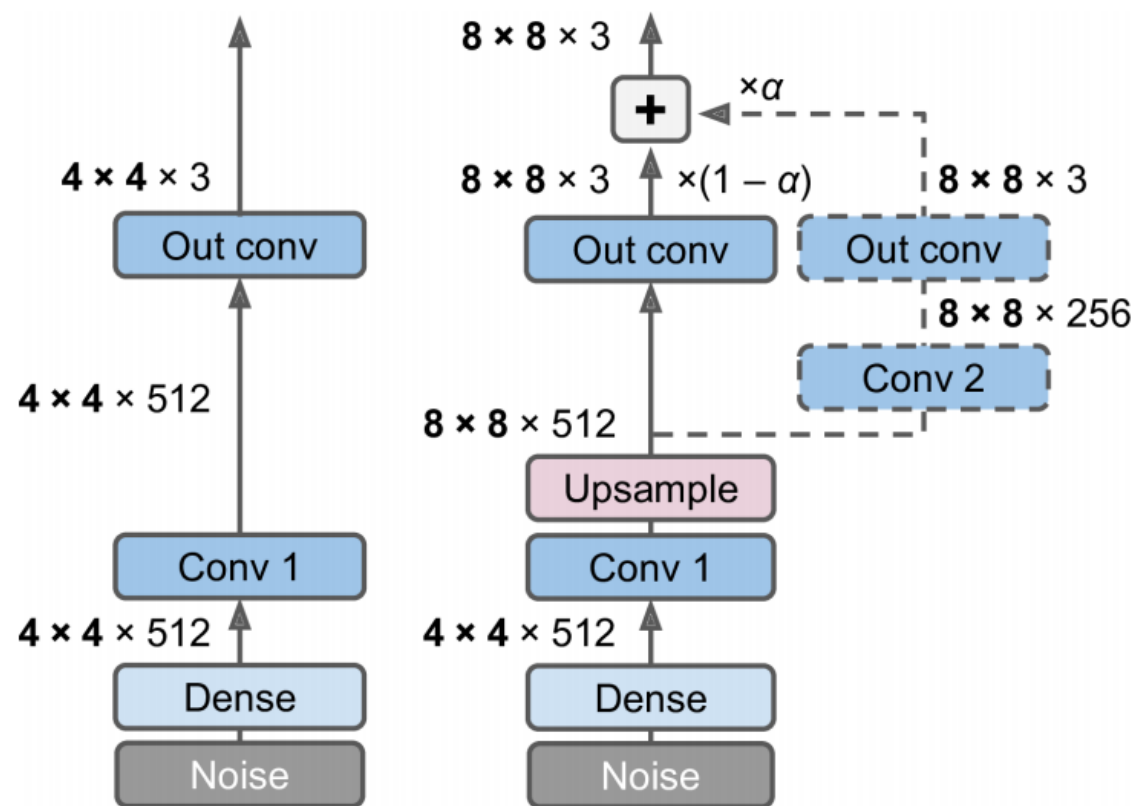
Conditional GAN  
(Mirza & Osindero, 2014)

MNIST digits generated conditioned on their class label.



Figure 2 in the original paper.

- Progressive Growing of GANs: Nivida Tero Karras, in 2018 paper
  - ✓ a GAN generator outputs  $4 \times 4$  color images (left)
  - ✓ extend it to output  $8 \times 8$  images (right)



- StyleGAN: Nvidia team
  - ✓ the state of the art in high-resolution image generation in 2018
  - ✓ style transfer techniques in the generator
- Adaptive Instance Normalization (AdaIN) :
  - ✓ each noise layer is followed by an AdaIN

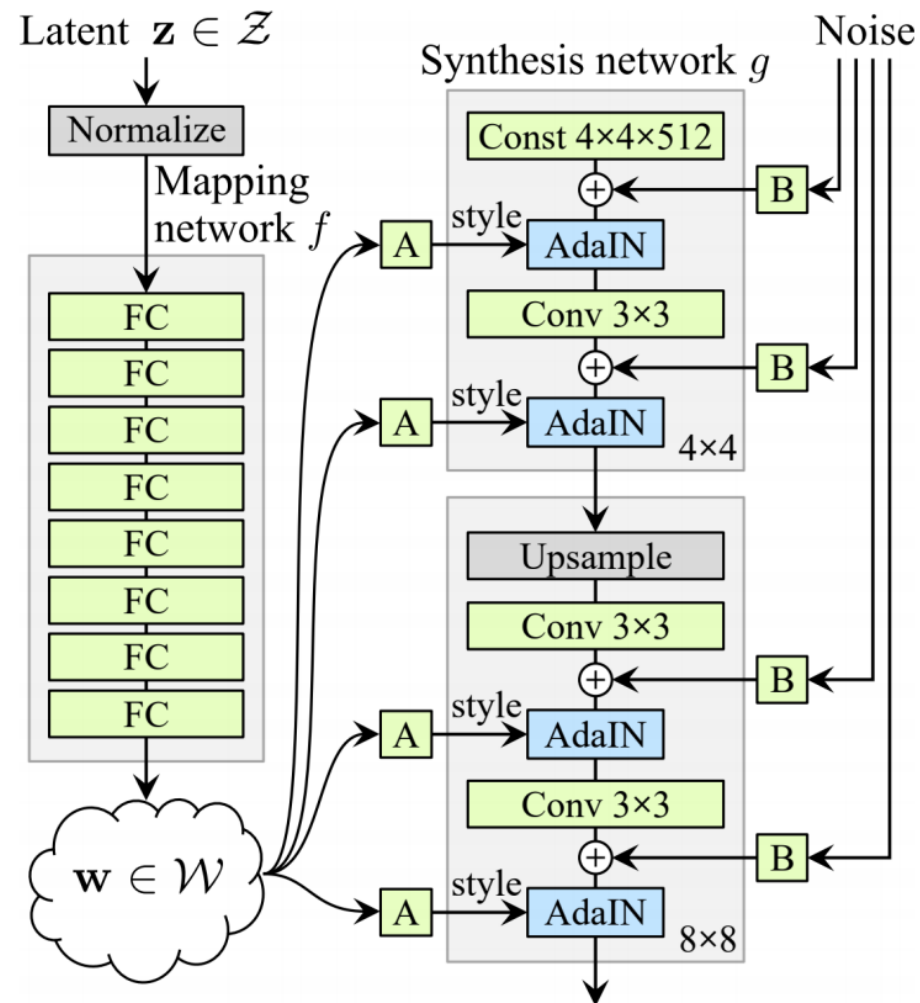


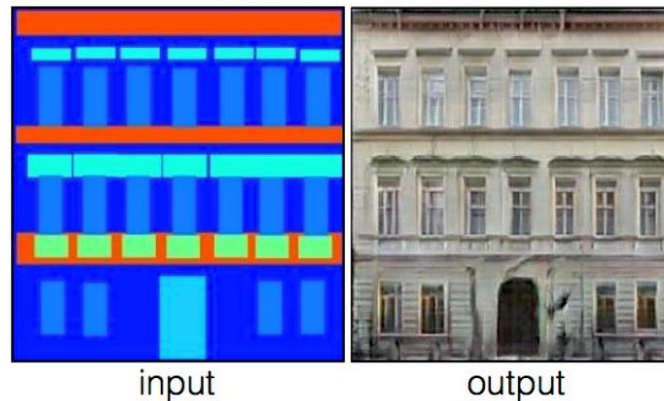
Figure 17-20. StyleGAN's generator architecture (part of figure 1 from the StyleGAN paper)<sup>19</sup>



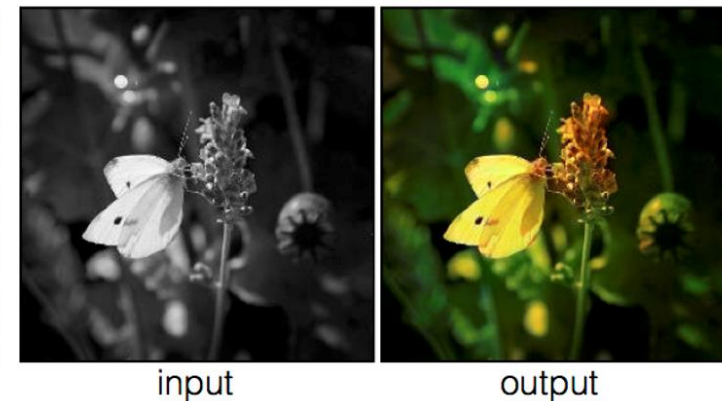
Labels to Street Scene



Labels to Facade



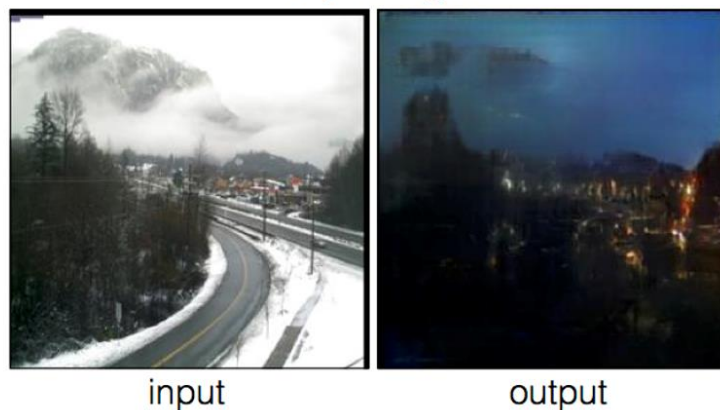
BW to Color



Aerial to Map



Day to Night



Edges to Photo



- Architecture: *DCGAN*-based architecture
- Training is conditioned on the images from the source domain.
- Conditional GANs provide an effective way to handle many complex domains without worrying about designing *structured loss functions* explicitly.

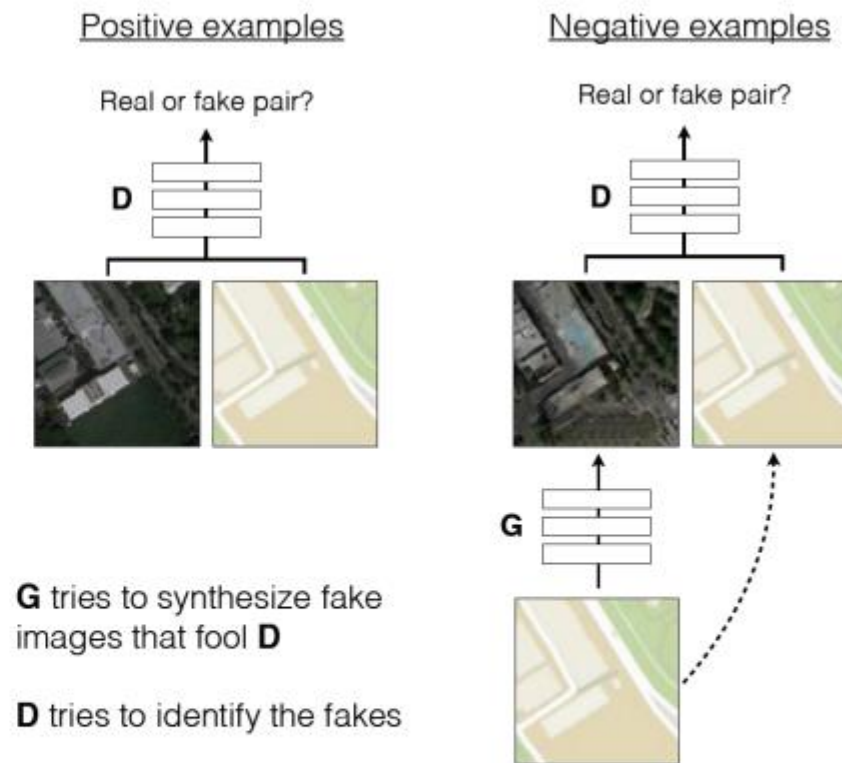
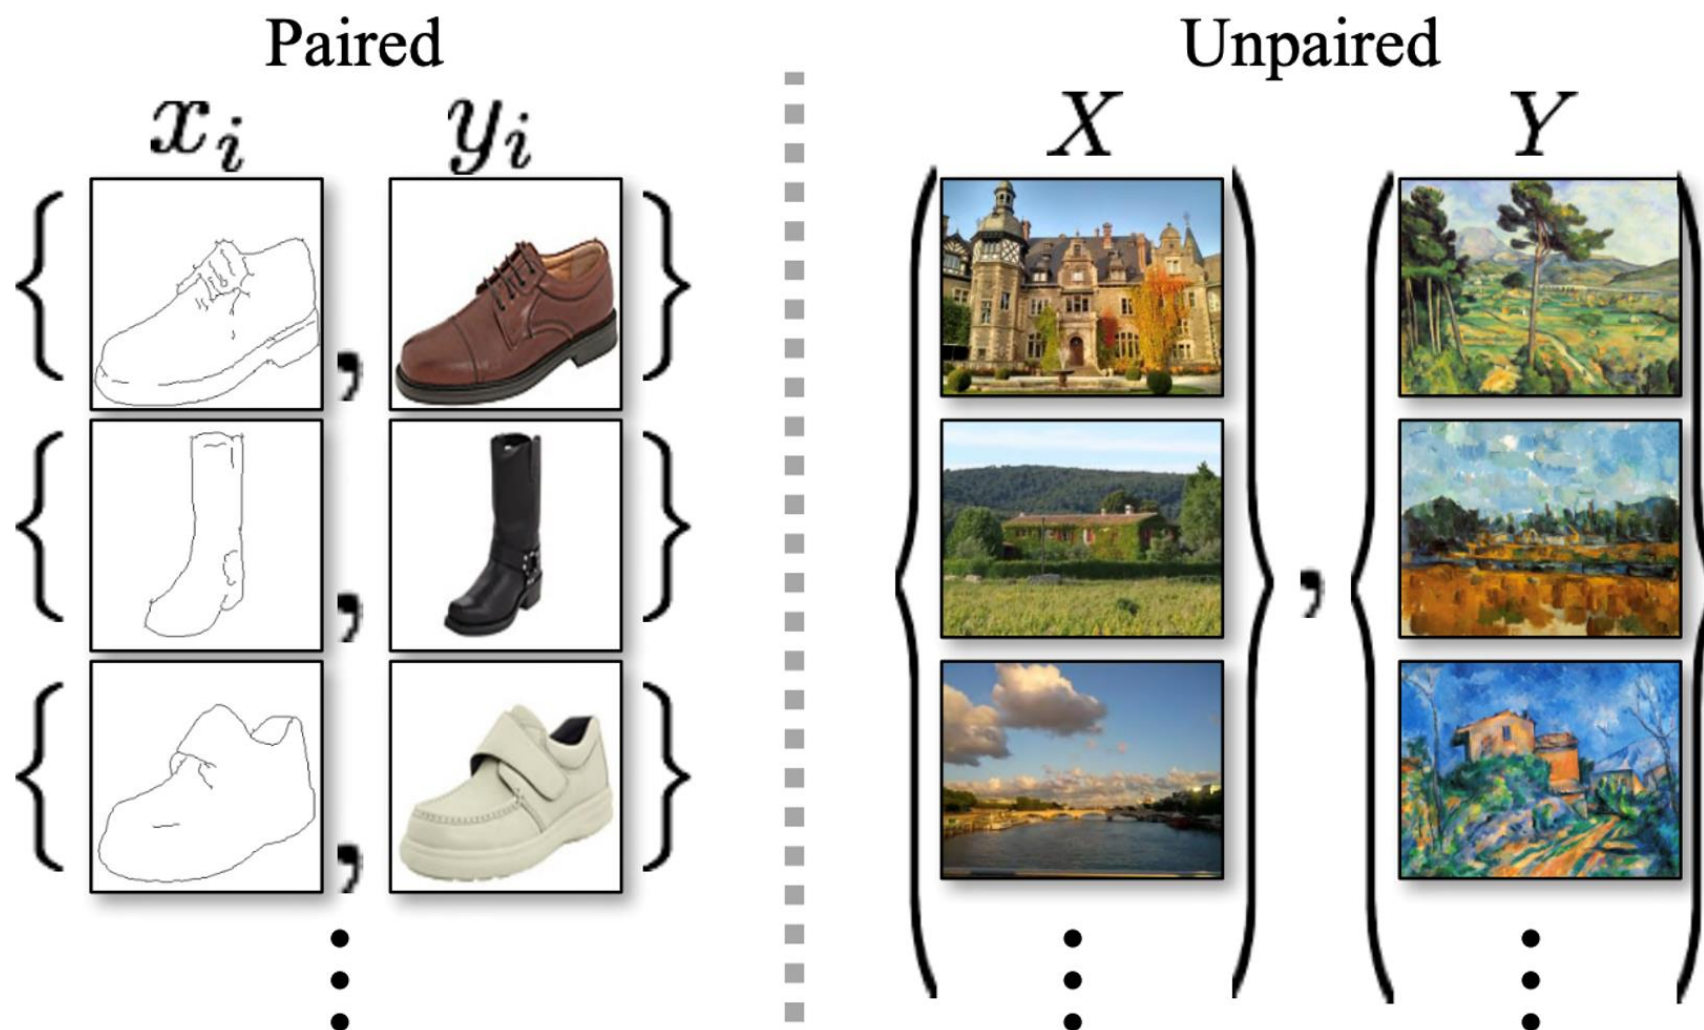


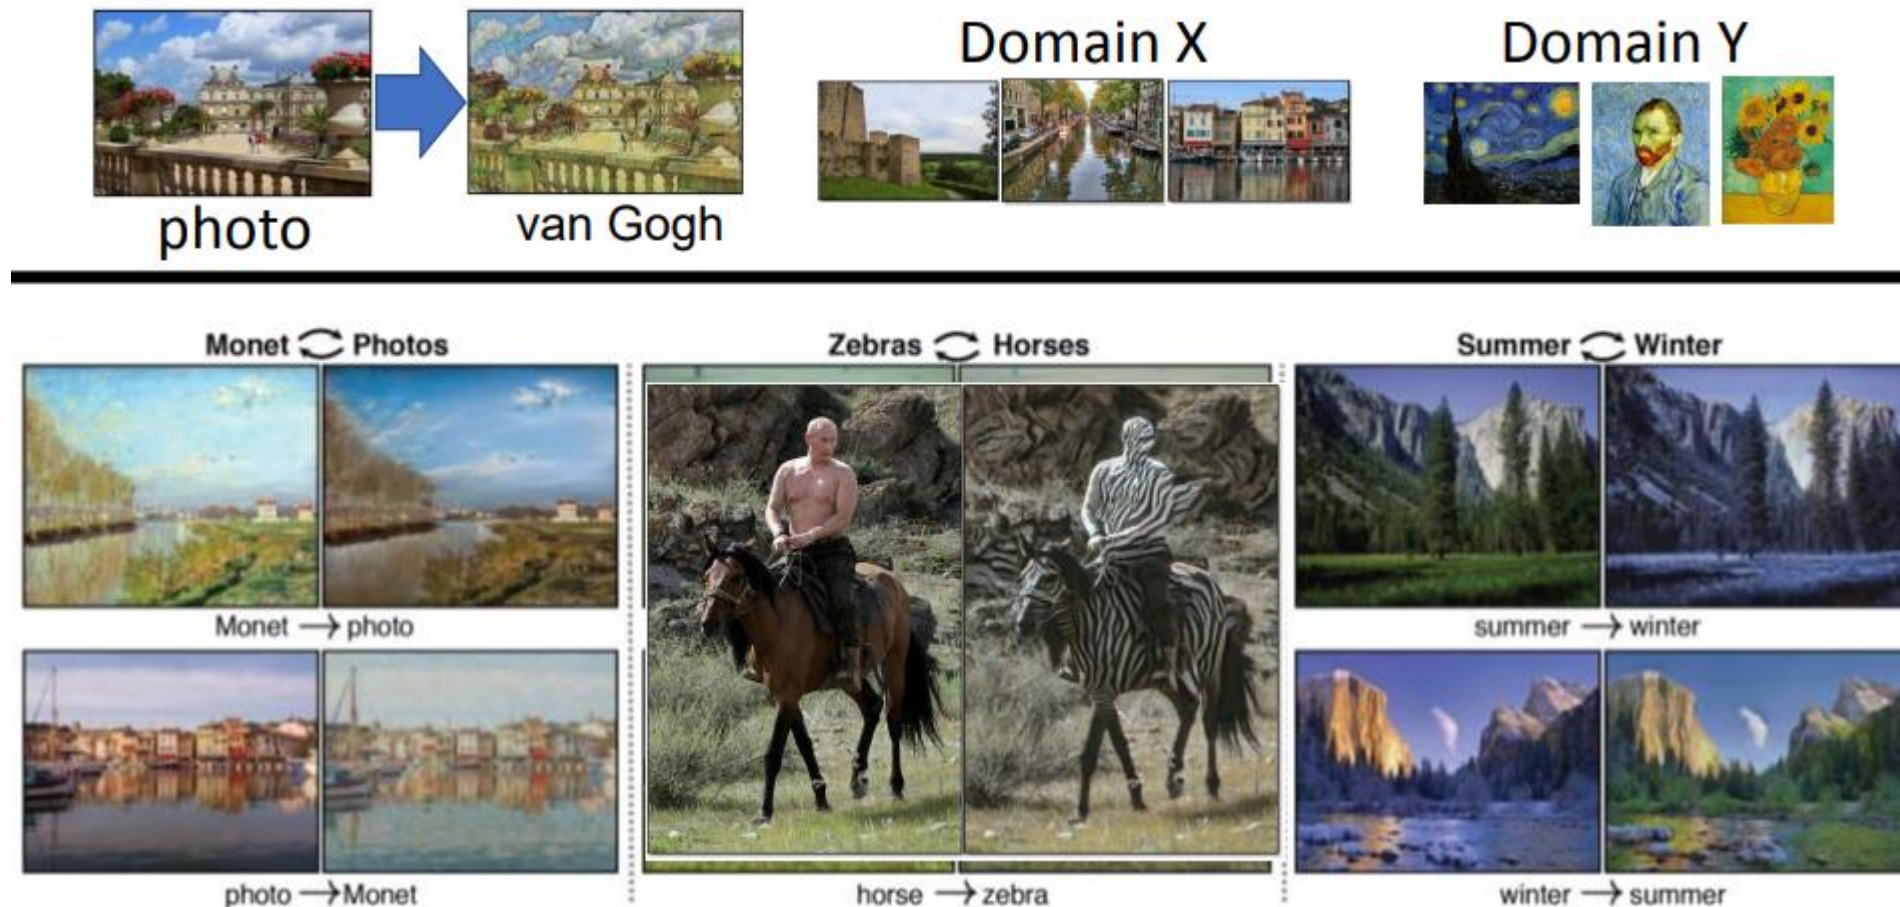
Figure 2 in the original paper.

Isola, P., Zhu, J. Y., Zhou, T., & Efros, A. A. "Image-to-image translation with conditional adversarial networks". arXiv preprint arXiv:1611.07004. (2016).





- Transform an object from one domain to another without paired data



- "Photo-realistic single image super-resolution using a generative adversarial network."
  - ✓ Ledig, Christian, et al.
  - ✓ arXiv preprint arXiv:1609.04802 (2016).





2022

Korea Institute of Science  
and Technology Information

TRUST  
KISTIL

