


# Advanced Topic in Research Data-centric Deep Learning

## Lec 14: Introduction to Deep Reinforcement Learning



## ❖ Background material

- ✓ Reinforcement Learning: An Introduction, Sutton & Barto 2018
  - <http://incompleteideas.net/book/the-book-2nd.html>
- ✓ Reinforcement Learning Lecture Series 2021
  - <https://www.deepmind.com/learning-resources/reinforcement-learning-lecture-series-2021>
- ✓ <https://huggingface.co/blog/deep-rl-intro>




# What is reinforcement learning?

# State Values and Policy Evaluation




- ❖ Reinforcement Learning can be thought of as learning from trial and error
  - ✓ An *agent*, that interacts with its *environment*, receives *rewards* that reflect its ability to accomplish some predefined goal.
- ❖ Reinforcement Learning can progressively move towards an agent that gives the maximum amount of reward and that solves the task at hand.
- ❖ RL consist of two distinct parts:
  - ✓ The *Prediction Problem*, in which the performance of the agent is evaluated.
  - ✓ The *Control Problem*, where the policy, used by the agent to select its actions, is modified to improve performance

- ❖ Once upon a time there was a Baby Robot who got lost in the mall
  - ✓ **reward**: a single numerical value that is used to measure how well the task at hand has been performed
  - ✓ **state**: each of these squares **environment** that we're working in.
    - When a state is independent of the prior states it is said to satisfy the *Markov Property*.





- ❖ Value: how good it is to be in a particular state
- ❖ Return: The expected total amount of reward
- ❖ Policy: The strategy used to select the next action
  - ✓ optimal : move in the direction of increasing value we are actually following best policy




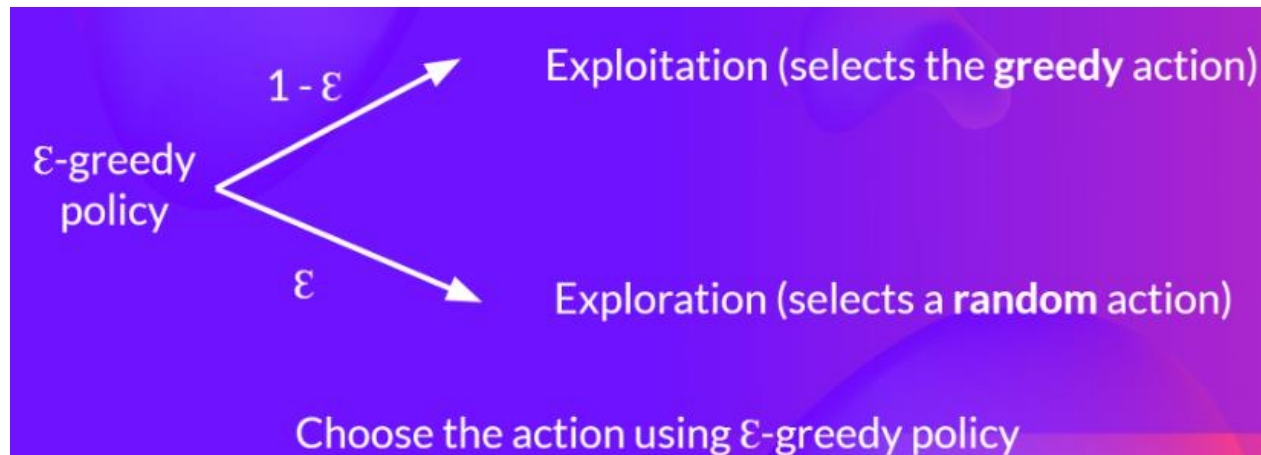
## ❖ The $\epsilon$ -greedy policy

✓ with probability  $1 - \epsilon$

- **we do exploitation** : our agent selects the action with the highest state-action pair value

✓ With probability  $\epsilon$ :

- **we do exploration** : trying random action and then gradually reduce it






- ❖  $r = \text{reward}$
- ❖  $a = \text{action}$
- ❖  $s = \text{state}, s' = \text{next state}$
- ❖ The rewards, states and actions are actually random variables:
  - ✓ there's a probability of getting a certain reward, taking a specific action/state
    - these probabilities are referred to using capital letters.
- ❖ the expected reward for a state-action pair:

$$r(s, a) = \mathbb{E}[R_t \mid S_{t-1} = s, A_{t-1} = a]$$

- ❖ Return ' $G_t$ ': the total amount of reward accumulated over an episode, starting at time ' $t$ '.
  - an *episode* refers to all the time steps that occur between entering and exiting a level.
- ✓ the return is just the sum of the future rewards

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$




- ❖ Policy: software agent uses to determine its actions which is commonly denoted by the symbol  $\pi$ 
  - ✓ So the value for state  $s$  under policy  $\pi$  is simply the expected return:

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t \mid S_t = s]$$

$$v_{\pi}(s) = r + v_{\pi}(s')$$

# Reinforcement Learning examples:



- (a) robotics,
- (b) Ms. Pac-Man,
- (c) Go player,
- (d) thermostat,
- (e) automatic trader

## ❖ DeepMind's AlphaGo Zero is really a scientific breakthrough

- ✓ Role of Deep Reinforcement Learning in achieving Strong AI
  - Mastering the game of Go without human knowledge , David Silver, et al. Nature(2017)




Figure 3 | Empirical evaluation of AlphaGo Zero. a

- ❖ **The agent :**
  - ✓ the program controlling a robot.
- ❖ **The environment :**
  - ✓ the real world
- ❖ **Observation : information from the env.**
  - ✓ partial description of the state of the world
- ❖ **State:**
  - ✓ complete description of the state of the world
- ❖ **Action :**
  - ✓ consist of sending signals to activate motors.
- ❖ **Reward: positive/negative**
  - ✓ it approaches the target destination,
  - ✓ it wastes time or goes in the wrong direction




$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{current value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{current value}} \right)}_{\text{temporal difference}}$$

new value (temporal difference target)



## ❖ Policy

- ✓ The algorithm a software agent uses to determine its actions
- ✓ A neural network taking observations as inputs and outputting the action to take

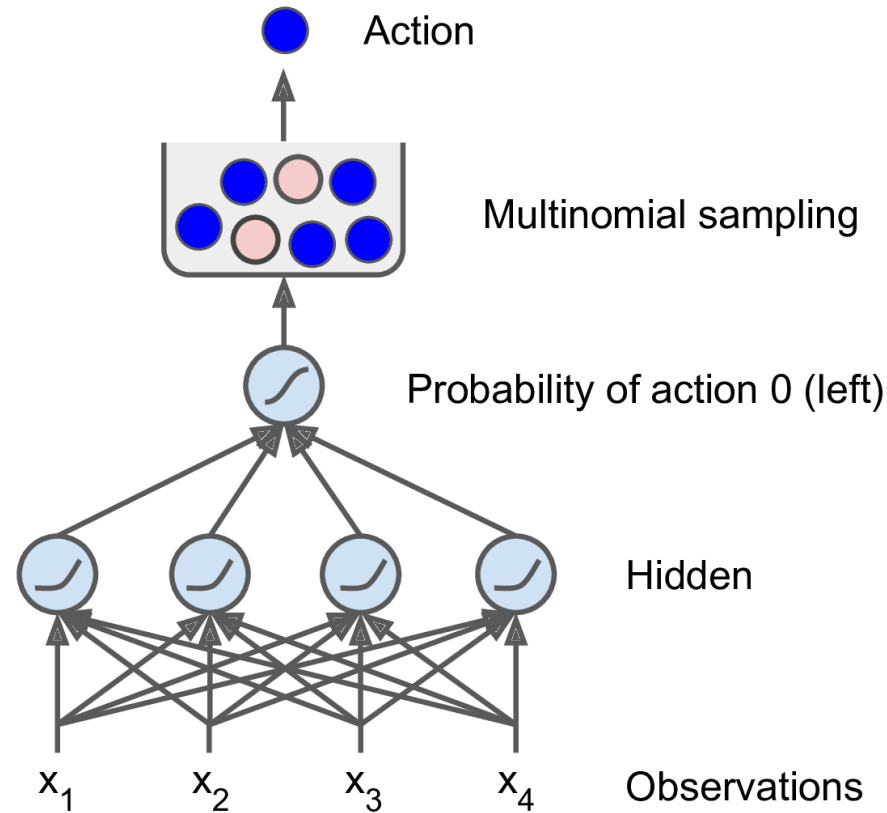


## ❖ Policy Search : Policy Gradient (PG)

- ✓ One approach is to use optimization technologies, by evaluating the gradients of the rewards with respect to policy parameters


## ❖ Two possible actions

- the probability  $p$  of action 0 (left) and the probability of action 1 (right) will be  $1-p$




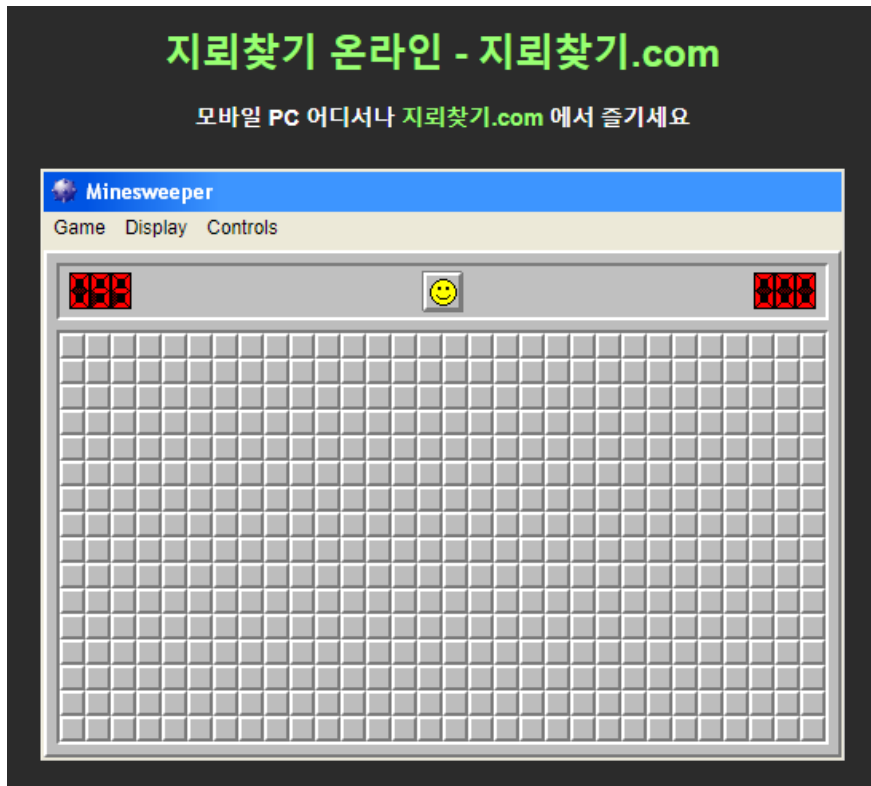
## ❖ Exploring and Exploitation

- ✓ picking a random actions
- ✓ *Exploitation*
  - *the process of taking benefits from things which we know about*
- ✓ *Exploration*
  - *to get knowledge about things which we didn't know*



## ❖ Exploring and Exploitation Strategy

✓ <https://zangsi.net/minesweeper/>



# Markov Decision Processes (MDP) and Bellman Equations

- ❖ Typically we can frame all RL tasks as MDPs
- ❖ The key in MDPs is the Markov Property
  - ✓ Essentially the future depends on the present and not the past
    - More specifically, the future is independent of the past given the present
- ❖ Types of Markov Models
  - ✓ Control over state transitions and completely observable states: MDPs
  - ✓ Control over state transitions and partially observable states: Partially Observable MDPs
  - ✓ No control over state transitions and completely observable states: Markov Chain
  - ✓ No control over state transitions and partially observable states: Hidden Markov Model

- ❖ Markov chains: In the early 20th century, the Andrey Markov studied stochastic processes with **no memory**
  - ✓ The probability for it to evolve from a state  $s$  to a state  $s'$  is fixed, and it depends only on the pair  $(s, s')$ , not on past states (this is why we say that the system has no memory)
- ❖ An example of a Markov chain with four states




Figure 18-7. Example of a Markov chain



## ❖ Markov Decision Processes were first described in the 1950s by Richard Bellman.

✓ They resemble Markov chains but with a twist:

- At each step, an agent can choose one of several possible actions, and the transition probabilities depend on the chosen action.
- Moreover, some state transitions return some reward (positive or negative),
- and the agent's goal is to find a policy that will maximize reward over time.

- ❖ Bellman found a way to estimate the optimal state value of any state  $s$ ,

$$\underbrace{V_{\pi}(s)}_{\text{Value of state } s} = \underbrace{\mathbf{E}_{\pi}}_{\text{Expected return}} \underbrace{[G_t | S_t = s]}_{\text{If the agent starts at state } s}$$

And uses the policy to choose its actions for all time steps

For each state,  
the state-value function outputs  
the expected return  
if the agent starts in that state  
and then follows the policy forever after.

## ❖ The action-value function

$$Q_{\pi}(s, a) = \mathbf{E}_{\pi}[G_t | S_t = s, A_t = a]$$

Value of state-action pair  $s, a$

Expected return


If the agent starts at state  $s$  and chooses action  $a$

And then uses the policy to choose its actions for all time steps

For each state and action, the action-value function outputs the expected return if the agent starts in that state and takes the action and then follows the policy forever after.

## ❖ For example:

- ✓ the MDP has three states (represented by circles) and up to three possible discrete actions at each step (represented by diamonds)
  - $R(s,a,s')$  is the reward that the agent gets when it goes from state  $s$  to state  $s'$ , given that the agent chose action  $a$ .



$$\bar{T}(s_2, a_1, s_0) = 0.8.$$

$$R(s_2, a_1, s_0) = +40.$$

Figure 18-8. Example of a Markov decision process


## ❖ The policy

- ✓ the agent's decision-making process
- ✓ The Policy  $\pi$  is the **brain of our Agent**
  - An agent **select the actions that maximize its expected cumulative reward**

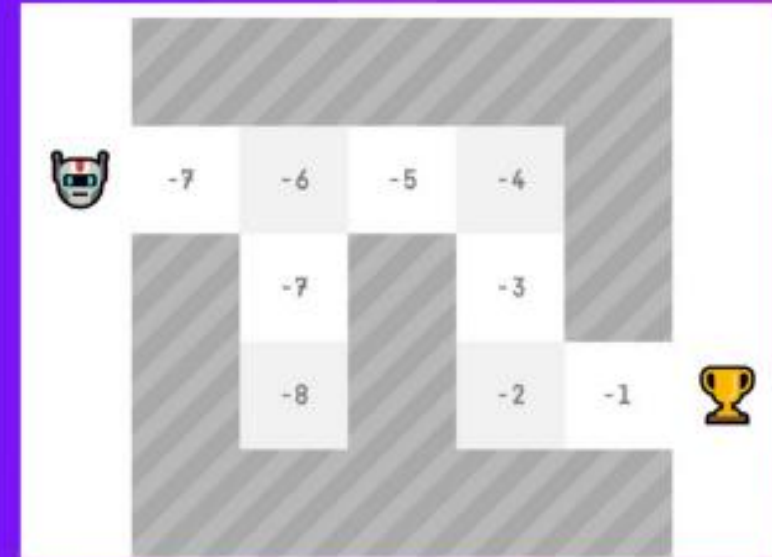


# There are two main types of RL methods

**Policy-Based methods:** train the agent to learn which **action to take**, given a state.



**Value-Based methods:** train the agent to learn which state is **more valuable** and take the action that leads to it.



- ❖ In policy-based: train directly the policy
  - ✓ Our policy is a Neural Network
  - ✓ No value function
- ❖ In value-based: don't train the policy
  - ✓ Our policy is a function defined by hand
  - ✓ Instead train a value-function that is a Neural Network
- ❖ Finding an optimal value function leads to having an optimal policy

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



## ❖ The Bellman Equation is one central to Markov Decision Processes

### ✓ The Bellman Equation

- “what is the maximum reward an agent can receive if they make the optimal action now and for all future decisions?”

The expected return (value) at the current state  $s$  is:

The expected reward for taking action  $a$  at state  $s$ ...

$$V(s) = \max_a (R(s, a) + \gamma V(s'))$$

The maximum value of any possible action  $a$  for:

...plus the discount factor (gamma) multiplied by the value of the next state

- ❖ Bellman found a way to estimate the optimal state value of any state  $s$

*Equation 18-1. Bellman Optimality Equation*

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \cdot V^*(s')] \quad \text{for all } s$$

$$T(s_2, a_1, s_0) = 0.8. \quad R(s_2, a_1, s_0) = +40. \quad \gamma \text{ is the discount factor.}$$

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \cdot V_k(s')] \quad \text{for all } s$$

- ❖ Knowing the optimal **state-values** can be useful, in particular to evaluate a policy, but it does not give us the optimal policy for the agent.
- ❖ **Q-Values (Quality Values):**
  - ✓ Bellman found a very similar algorithm to estimate the **optimal state-action values**
  - ✓  $Q^*(s, a)$ : the optimal Q-Value of the state-action pair (s, a):
  - ✓ Defining the optimal policy, noted  $\pi^*(s)$ :

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma \cdot \max_{a'} Q_k(s', a') \right] \quad \text{for all } (s/a)$$

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^*(s, a)$$

## ❖ Monte Carlo Approach:

- ✓ Monte Carlo uses an entire episode of experience before learning
- ✓ waits until the end of episode, then calculates return( $G_t$ ) and use it as a target for its value of policy

## ❖ Temporal Difference (TD)

- ✓ uses only a step to learn

Monte Carlo:  $V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$

TD Learning:  $V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$

# Temporal Difference (TD) Learning

## ❖ The TD Learning algorithm is very similar to the Value Iteration algorithm

✓ the TD Learning algorithm updates the estimates of the state values based on the transitions and rewards that are actually observed

✓ Goal:

- learn  $V_{\pi}(s)$  from episodes of experience under policy  $\pi$

$$V_{k+1}(s) \leftarrow V_k(s) + \alpha(R + \gamma V_k(s') - V_k(s))$$

*TD target*

*\alpha is the learning rate (e.g., 0.001)*

*TD error*

$$V(s) \leftarrow_{\alpha} r + \gamma \cdot V(s')$$

- ❖ Policy evaluation (the prediction problem):
  - ✓ for a given policy  $\pi$ , compute the state-value function  $V_\pi(s)$
- ❖ The simplest Temporal-Difference method TD(0):

$$V_{k+1}(s) \leftarrow V_k(s) + \alpha(\underbrace{R + \gamma V_k(s')}_{\text{TD target}} - V_k(s))$$


*TD target: an estimate of the return*



## ❖ SARSA is an on-policy algorithm

- ✓ while learning the optimal policy it uses the current estimate of the optimal policy to generate the behavior

## ❖ Estimate optimal policy $q_\pi$ for the current policy $\pi$



- ✓ After every transition from a nonterminal state,  $S_t$ , do this:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

If  $S_{t+1}$  is terminal, then define  $Q(S_{t+1}, A_{t+1}) = 0$

# Q-Learning

## ❖ Q-Learning is an off-policy value-based method

- ✓ that uses a Temporal Difference(TD) approach to train its action-value function
- ✓ Value-based method
  - finds the optimal policy indirectly by training a value or action-value function that will tell us **the value of each state or each state-action pair.**
- ✓ Uses a TD approach:
  - **updates its action-value function at each step instead of at the end of the episode.**
- ✓ Off-policy:
  - **using a different policy for acting and updating using epsilon-greedy policy again**

## ❖ Q-Learning is the algorithm to train an action-value function

- ✓ that determines the value of being at a particular state and taking a specific action at that state.
  - Given a state and action, Q function outputs a state-action value (also called Q-value)



<https://huggingface.co/blog/deep-rl-q-part2>

- ❖ In Q-learning the learned action-value function,  $Q$ , directly approximates the optimal action-value function, independent of the policy being followed.
  - the transition probabilities are unknown and the rewards are initially unknown
  - Q-Learning works by watching an agent play (e.g., randomly) and gradually improving its estimates of the Q-Values

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left( r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right)$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

- Current Q-table value we are updating
- Learning rate
- Reward
- Discount
- Estimated reward from our next action

## ❖ Off-policy: The Q-learning algorithm

- ✓ work by "*looking over someone's shoulder*."
- ✓ the algorithm attempts to learn about policy  $\pi$  from experience sampled from  $\mu$ .
- ✓ the policy being executed is completely random, while the policy being trained will always choose the actions with the highest Q-Values.
  - Q-Learning is capable of learning the optimal policy by just watching an agent act randomly
  - learning to play golf when your teacher is a drunk monkey

## ❖ On-policy: The Policy Gradients algorithm, and SARSA

- ✓ We can say that algorithms classified as **on-policy** are "*learning on the job*."
  - In other words, the algorithm attempts to learn about policy  $\pi$  from experience sampled from  $\pi$ .
  - it explores the world using the policy being trained.