

Deep Learning based Text Processing

Lec 09: Introduction to Recurrent Neural Network



hsyi@kisti.re.kr

Hongsuk Yi (이홍석)



❖ 성명 : 이홍석

❖ 소속

- ✓ KISTI 데이터기반문제해결연구단 책임연구원
- ✓ 과학기술연합대학원(UST-KISTI) 응용AI전공 교수

❖ 연구내용

- ✓ 딥러닝기반 도심지 교통혼잡 예측 (2018~2021)
- ✓ 지능형 인프라 기술 연구 (2018~2019)
- ✓ 도심지 교통예측을 위한 트랜스포머 모델 연구(2022~2025)

❖ 컴퓨팅 기술

- ✓ GPU 기반 가속컴퓨팅 기술 : CUDA, OpenCL
- ✓ 슈퍼컴퓨팅 병렬처리 기술 : MPI, OpenMP, ManyCore Computing
- ✓ Deep Learning 기술 : Tensorflow, PyTorch, Keras, Theano



Course Materials : github.com/hongsuk.yi

Information : sites.google.com/view/deepai/home

publication :<https://www.researchgate.net/profile/Hongsuk-Yi/publications>



The screenshot shows a personal website for Hongsuk Yi. The header features a dark blue background with a geometric pattern. On the left is a sidebar with a black background and white text, listing 'DeepAI', 'Home', 'Research', 'Lecture', 'Publications', and 'Contacts Us'. The main content area has a light gray background. At the top right is a photo of Hongsuk Yi speaking at a podium. To his right is his name in Korean and English, and below it is the text 'Korea Institute of Science and Technology Information'. A dark blue horizontal bar below the photo contains the heading 'Work Experiences' and a list of three positions: 'Principal Researcher, Korea Institute of Science and Technology Information (KISTI)/Supercomputing Center(1999.12~present)', 'Adjunct Professor, University of Science and Technology (UST) (2020.2~present)', and 'Adjunct Professor, Korea University (2009.8~2013.12)'. Another dark blue horizontal bar below that contains the heading 'Research Interests' and a list of three topics: '텐서플로우 2.0 딥러닝 알고리즘 성능 최적화', '라지스케일 GPU 분산 병렬 컴퓨팅', and 'CCTV 영상에서 차량/사람/사물을 탐색, 추적, 재인식 기술'.

Reviewing the last class:

LSTM

- ❖ **Basic knowledge of NLP**

- ✓ Natural language processing concepts and its applications

- ❖ **Word Embedding Theory**

- ✓ NPLM, Word2Vec, FastTex

- ❖ **Word Embedding - Practice 1**

- ✓ Word embedding based on Word2Vec

- ❖ **Word Embedding - Practice 2**

- ✓ Word embedding and its Applications

- ❖ **The concept of machine learning and deep learning**
 - ✓ practice of libraries for deep learning (Numpy, Pandas)
- ❖ **Understanding regression problems with FCN**
 - ✓ Regression problem practice (housing price prediction)
- ❖ **Understanding classification problems with FCN**
 - ✓ Classification problem practice (diabetes onset prediction)
- ❖ **Text Analysis Practice with FCN**
 - ✓ Text classification with movie review

❖ Introduction to Recurrent Neural Network

- ✓ Simple RNN, BPTT, Memory Cell
- ✓ Code: Implementing an RNN with Keras

❖ Introduction to Long-Short Term Memory

- ✓ Cell state, LSTM, and GRU, and Applications
- ✓ A Visual Guide to Recurrent Layers in Keras
- ✓ Code: A simple LSTM layers

❖ Text generation with RNN

- ✓ Tokenizer, Character-Level Language model
- ✓ Code: Alice's Adventures in Wonderland

❖ Sequence to Sequence Learning model with RNN

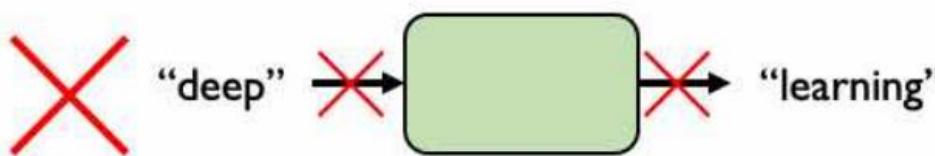
- ✓ Introduction to Seq2Seq and Attention model
- ✓ Code: Character-Level Neural Machine Translation

“This morning I took the dog for a walk.”

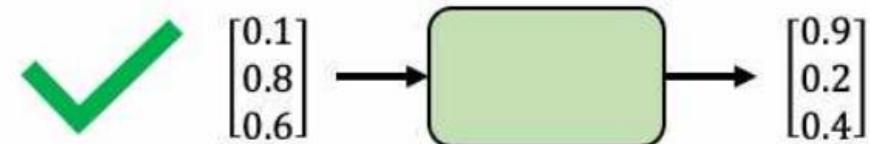
given these words

*predict the
next word*

Representing Language to a Neural Network



Neural networks cannot interpret words



Neural networks require numerical inputs

given these two words

“This morning I took the dog **for a** **walk.**” *predict the next word*

One hot feature encoding : tells us what each word is

for a [1 0 0 0 0 0 1 0 0 0] —→ prediction

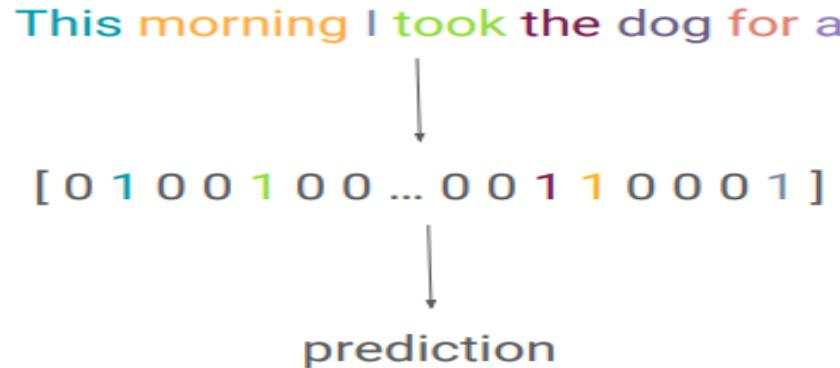
Problem : Can't Model Long-Term Dependencies

“France is where I grew up, but I now live in Boston. I speak fluent ____.”

We need information from **the distant past** to accurately predict the current word.

Trial: Use Entire Sequence as Set of Counts

- ❖ Bag-of-words model: a text is represented as the bag of its words.



Problem : Counts Don't Preserve Order



“The food was good, not bad at all.”

vs

“The food was bad, not good at all.”



Trial : Use a Really Big Fixed Window

"This morning I took the dog for a walk." given these 7 words, predict the next word

[1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 ...]

morning I took the dog ...

Problem : No Parameter Sharing

this morning
[1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 ...]

each of these inputs has a separate parameter

this morning
[0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 1 0 0 ...]

- ❖ Design criteria to model sequences, we need:

1. to deal with variable-length sequences
2. to maintain sequence order
3. to keep track of long-term dependencies
4. to share parameters across the sequence



RNN meet these sequence modeling design critera

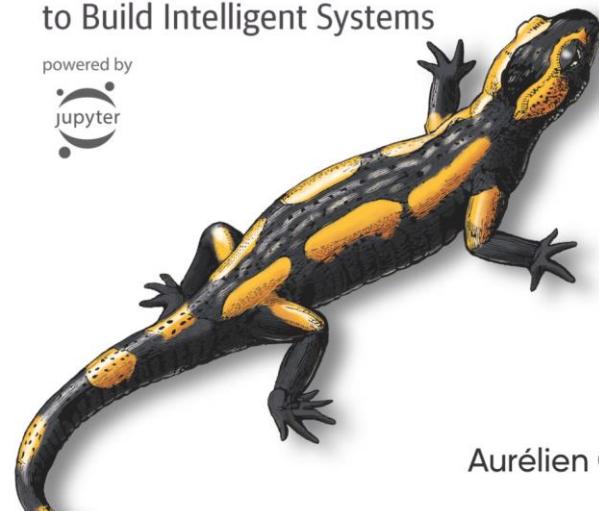
Introduction to Recurrent Neural Network

O'REILLY®

Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow

Concepts, Tools, and Techniques
to Build Intelligent Systems

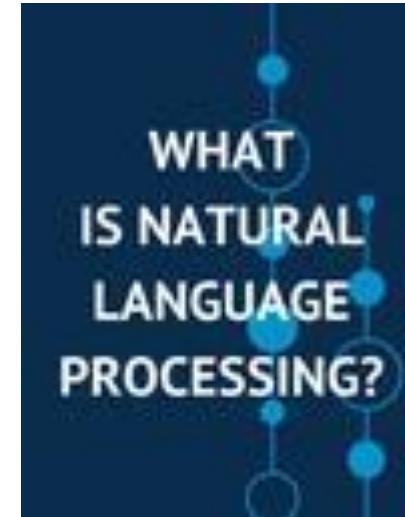
powered by



Aurélien Géron

2nd Edition
Updated for
TensorFlow 2

딥 러닝을 이용한 자연어 처리 입문
<https://wikidocs.net/book/2155>

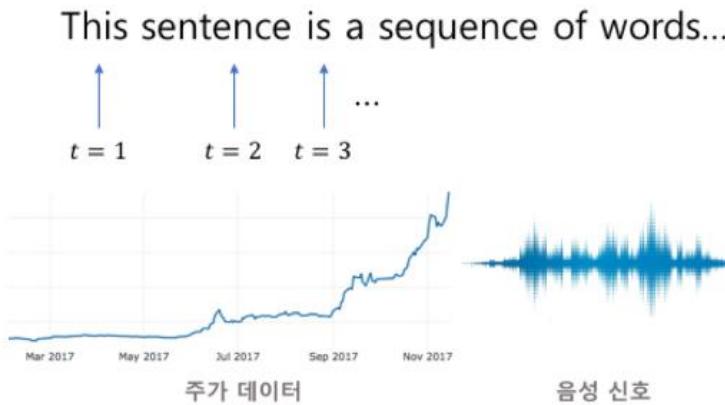


What is a time series problem?

❖ Time series deals with data over time:

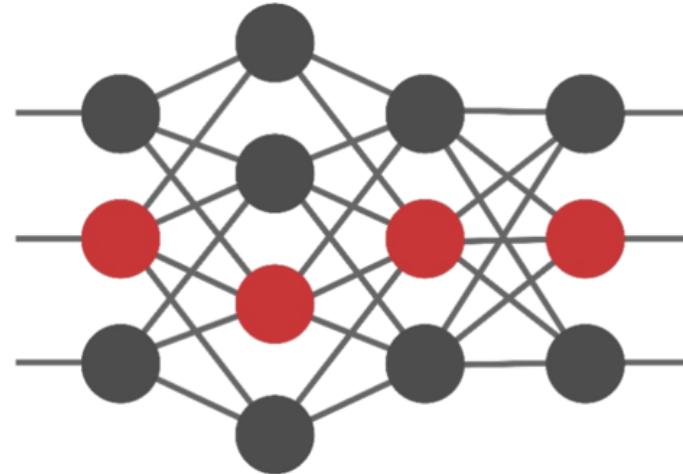
✓ Predicting

- weather (short term forecasts), climate (long term forecasts)
- sales accross different geographies for different products in the catalog



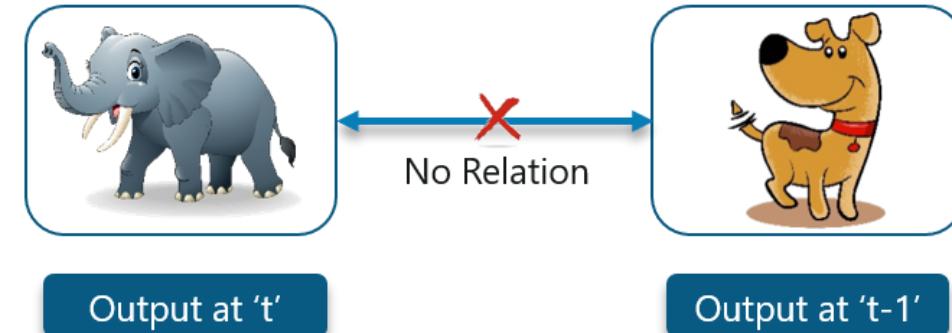
Why Not Feedforward Networks

In the existing FeedForWared Networks,



1) The input is a dog image, and a well-learned neural network outputs a dog label.

(3) However, elephants and dogs are not related to each other and are independent.

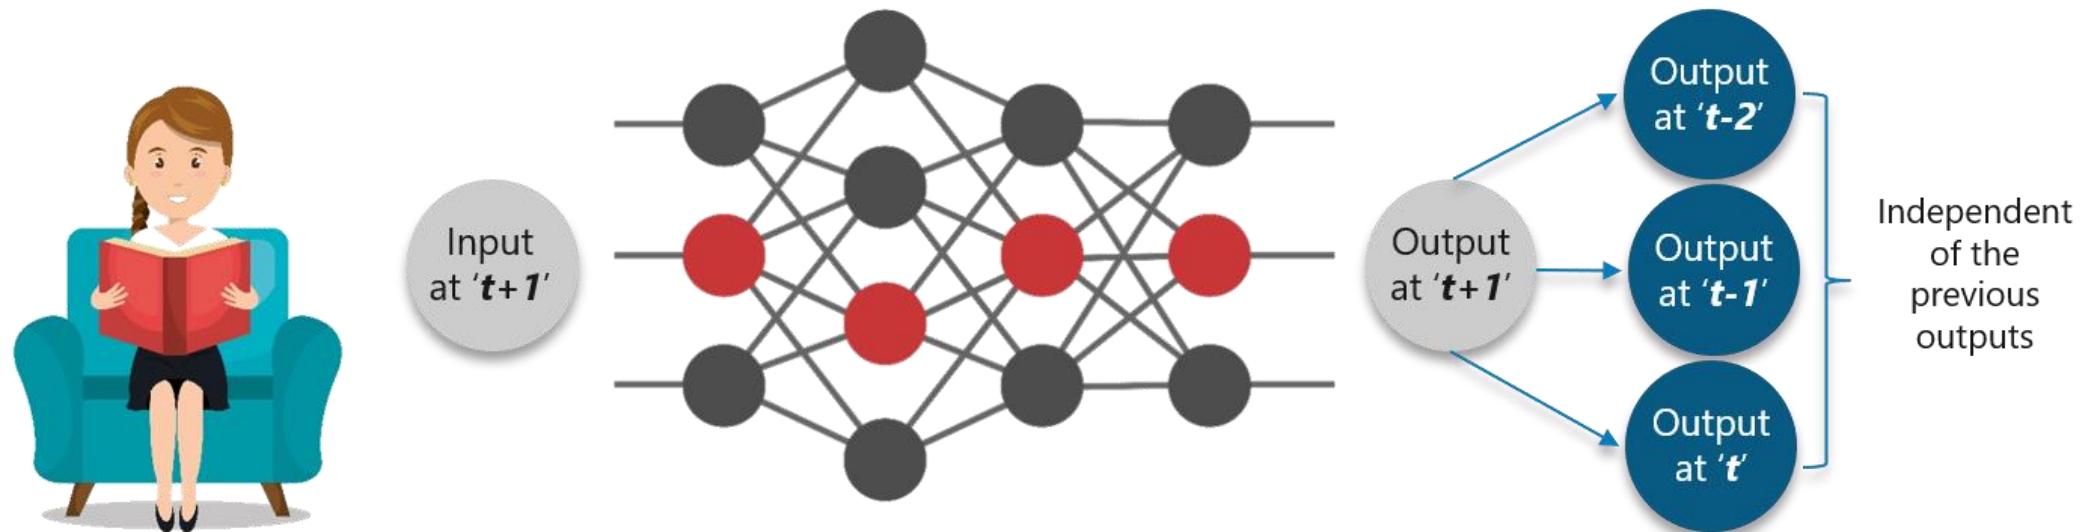


(2) The next elephant image is input to output an elephant (label).

Why Not Feedforward Networks

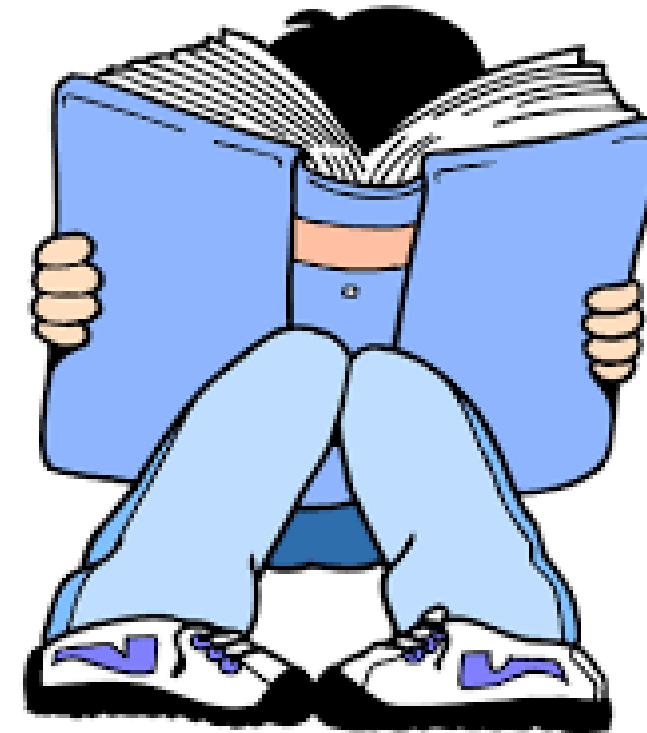
- FN cannot predict the relationship between the following words with previous words.

If you are reading a book, you should be well aware of the previous page.



❖ What is Memory Size?

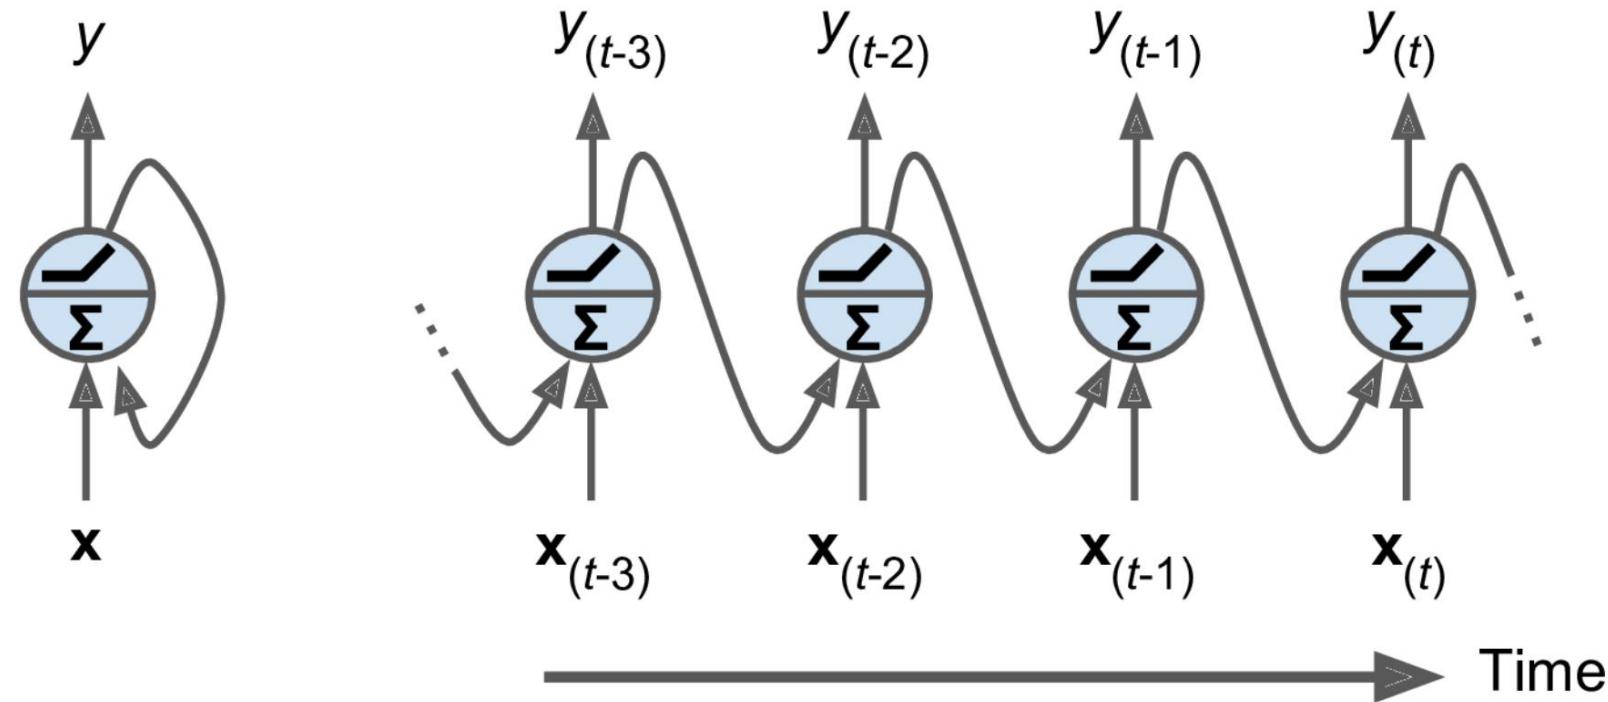
- ✓ 1 page
- ✓ 10 pages
- ✓ 100 pages
- ✓ 200 pages



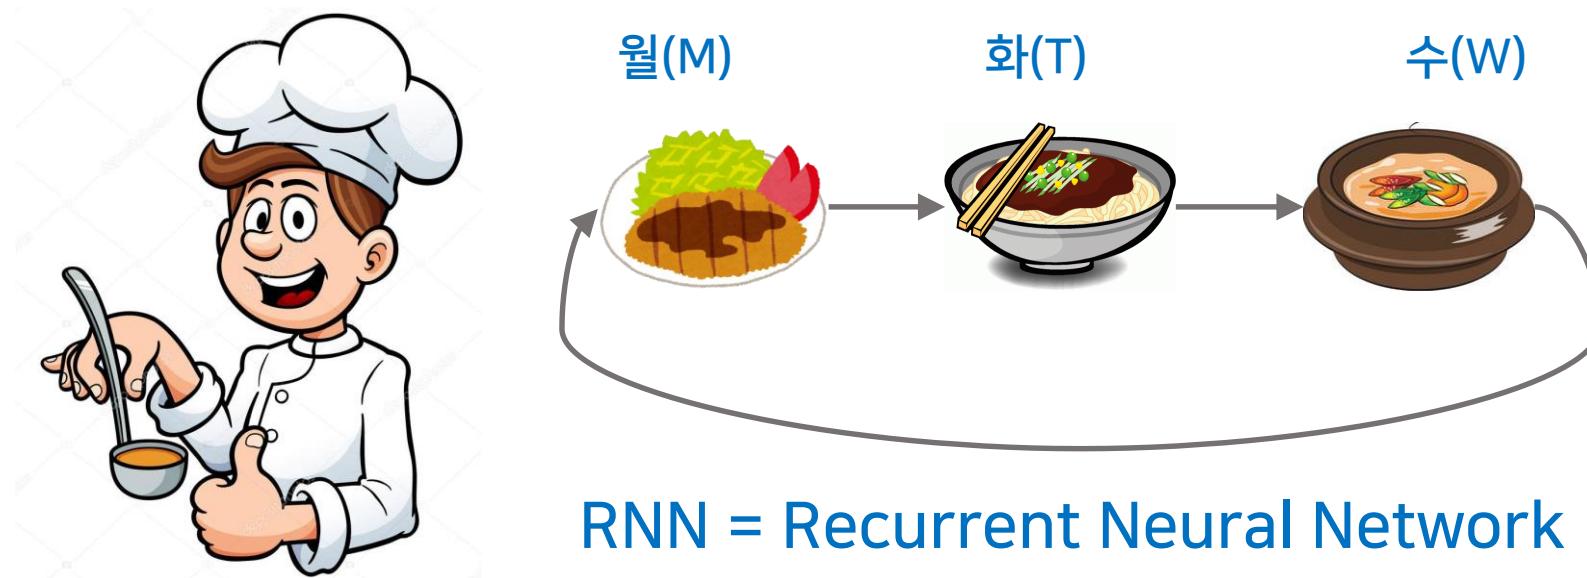
Recurrent Neurons and Memory Cells

❖ Memory Cells ~ a form of memory

- ✓ A RNN looks like a feedforward NN, except it also has connections pointing backward.



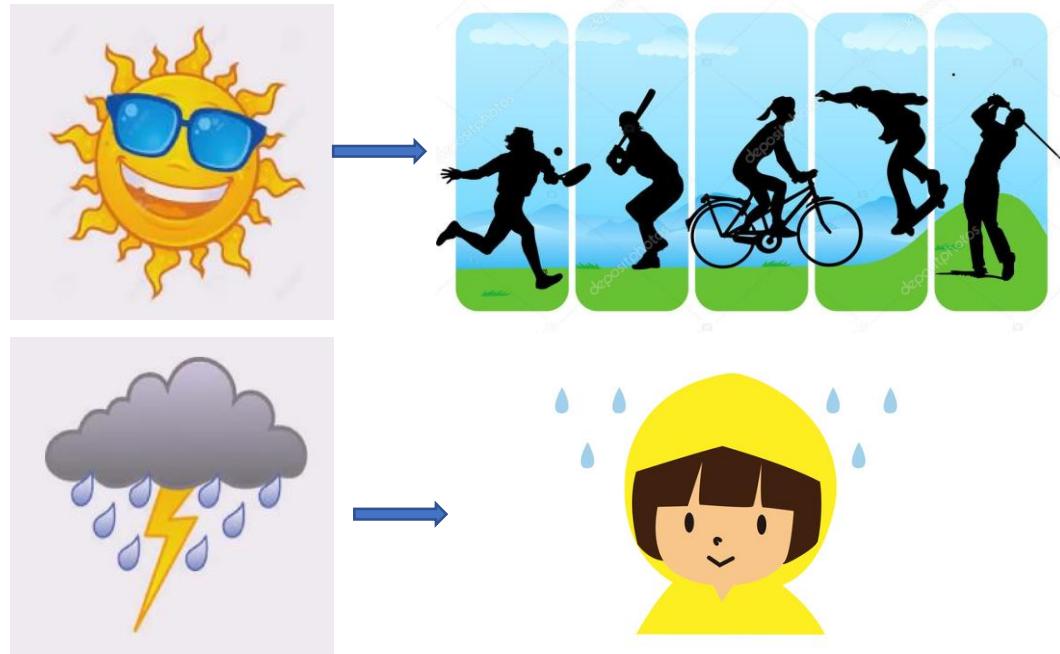
- ❖ There is a pattern in dinner.
 - ✓ Only three types of meals are available.



A simple of RNN (2/3)

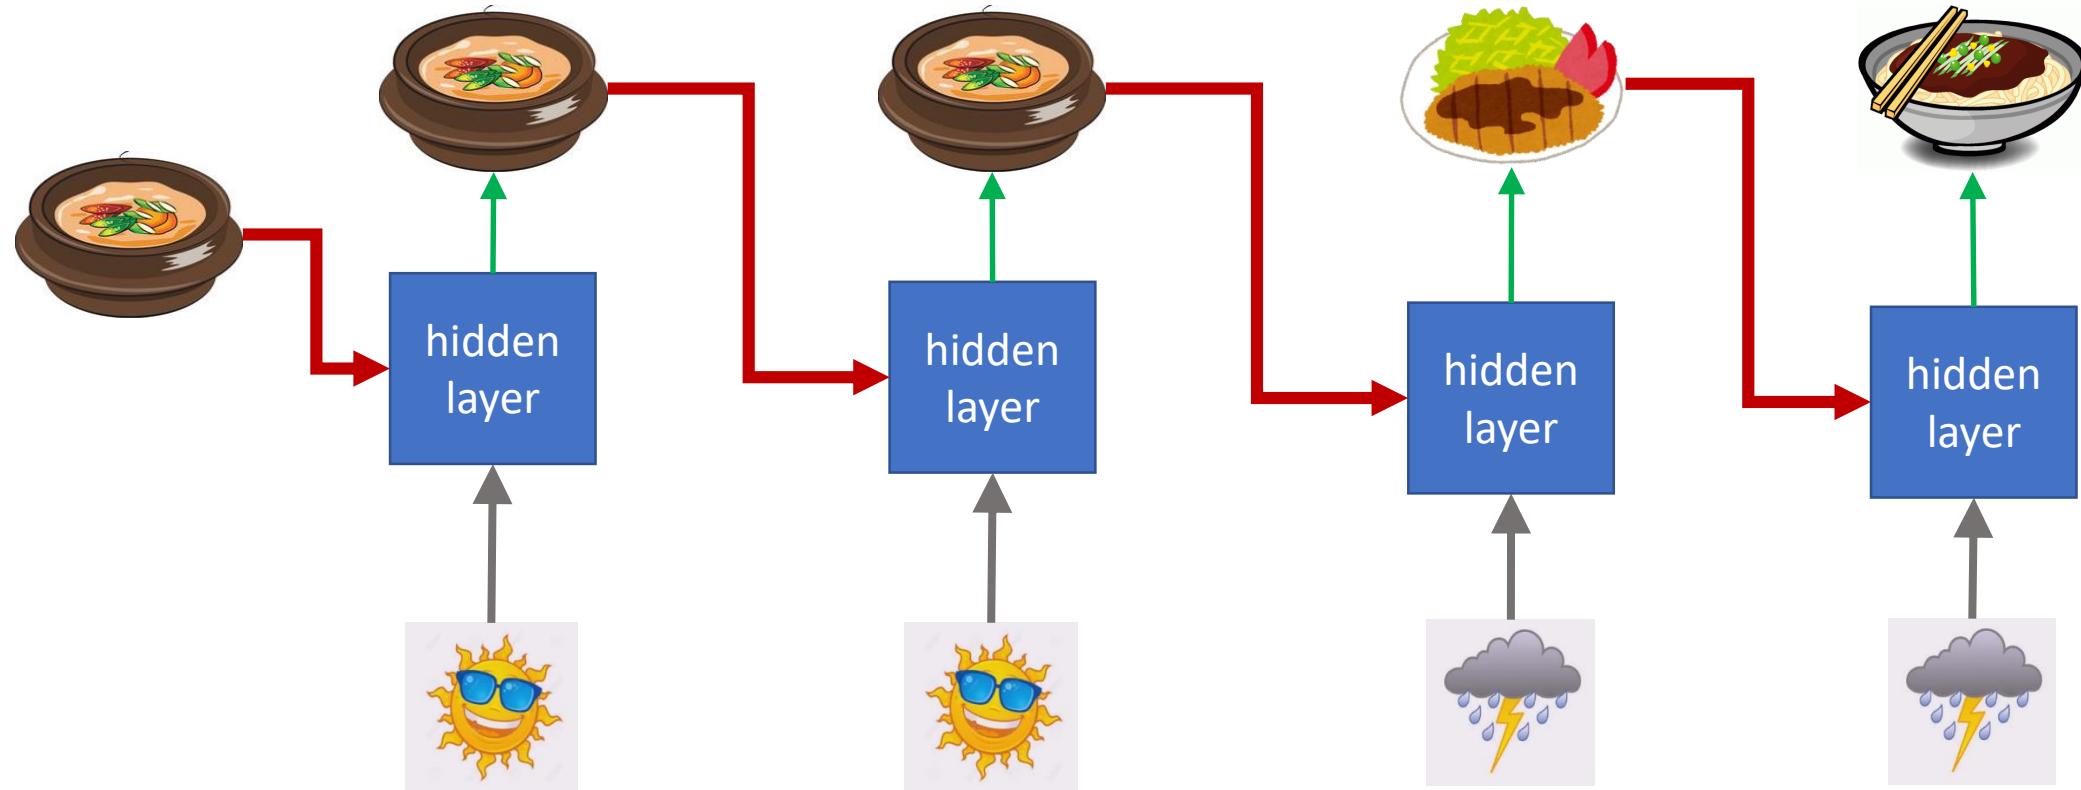
❖ The menu varies depending on the weather

- ✓ If the weather is nice, I eat food the day before.
- ✓ When it rains, I eat food according to the pattern



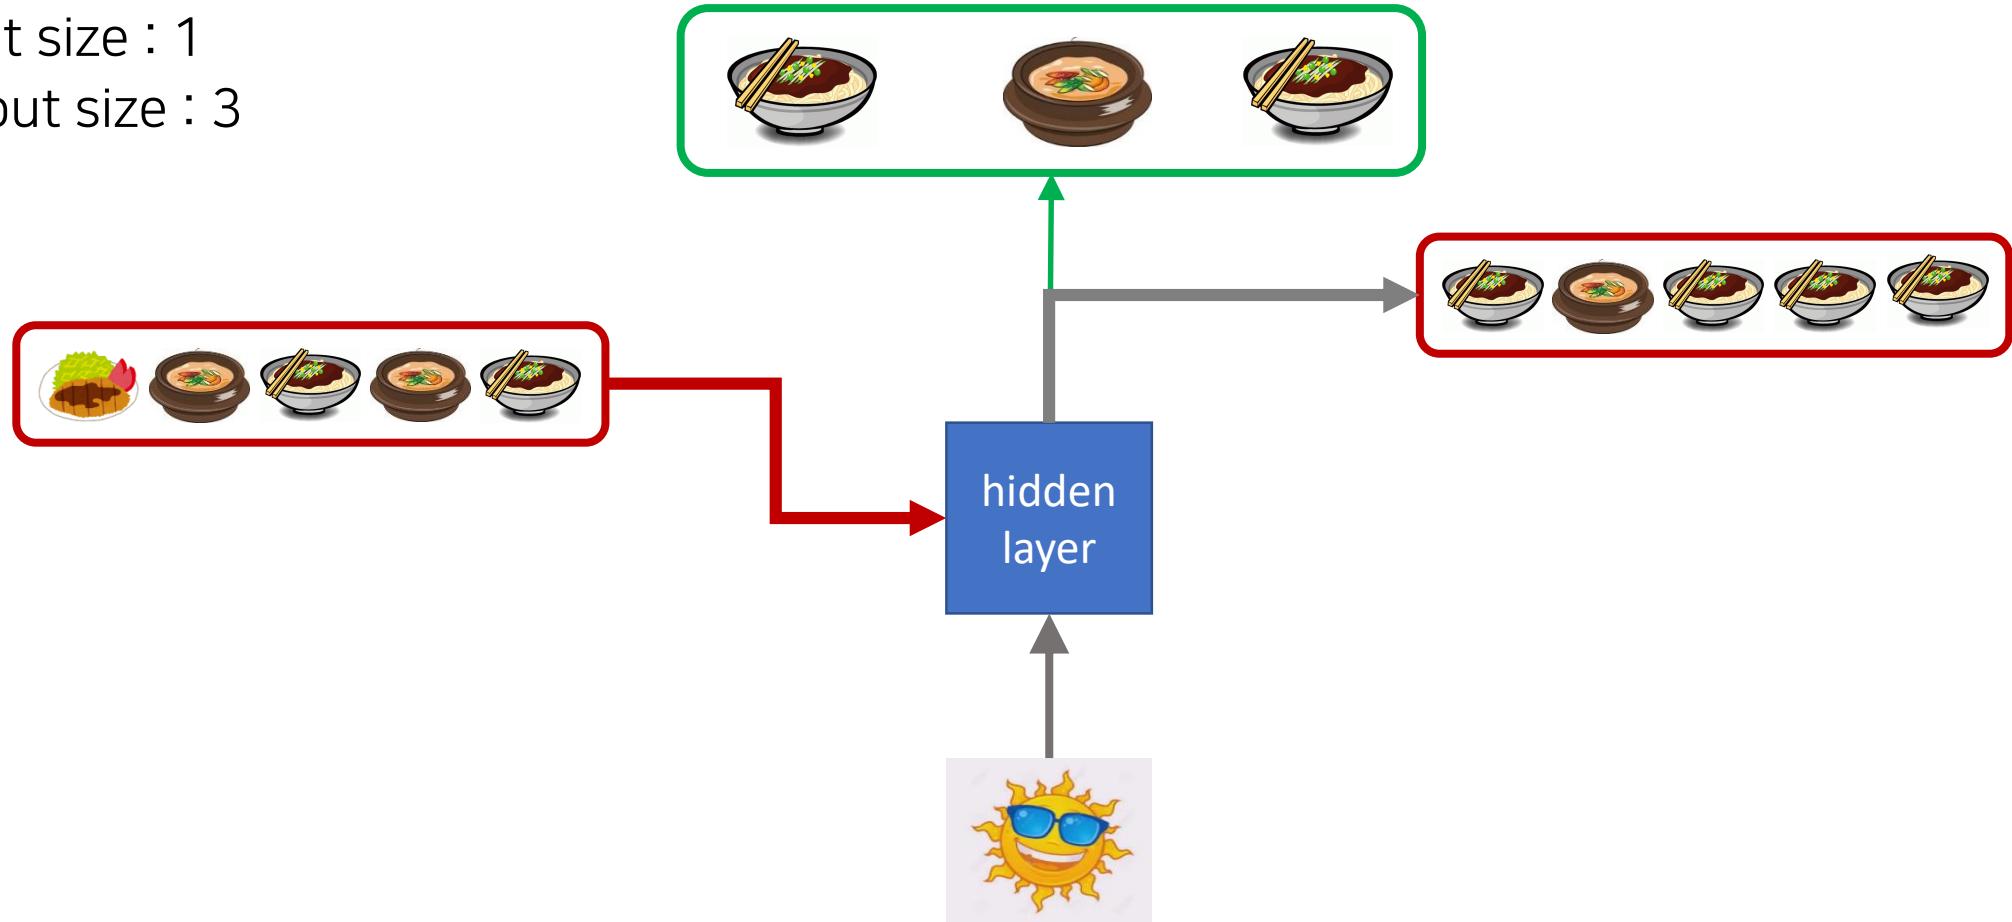
A simple of RNN (3/3)

- ❖ You can predict food depending on the weather



❖ A example of Length of memory

- ✓ memory size: 5
- ✓ input size : 1
- ✓ output size : 3



Neurons in RNN

Output of a recurrent layer for a single instance

$$\mathbf{y}(t) = \phi(\mathbf{W}_x^\top \mathbf{x}(t) + \mathbf{W}_y^\top \mathbf{y}(t-1) + \mathbf{b})$$

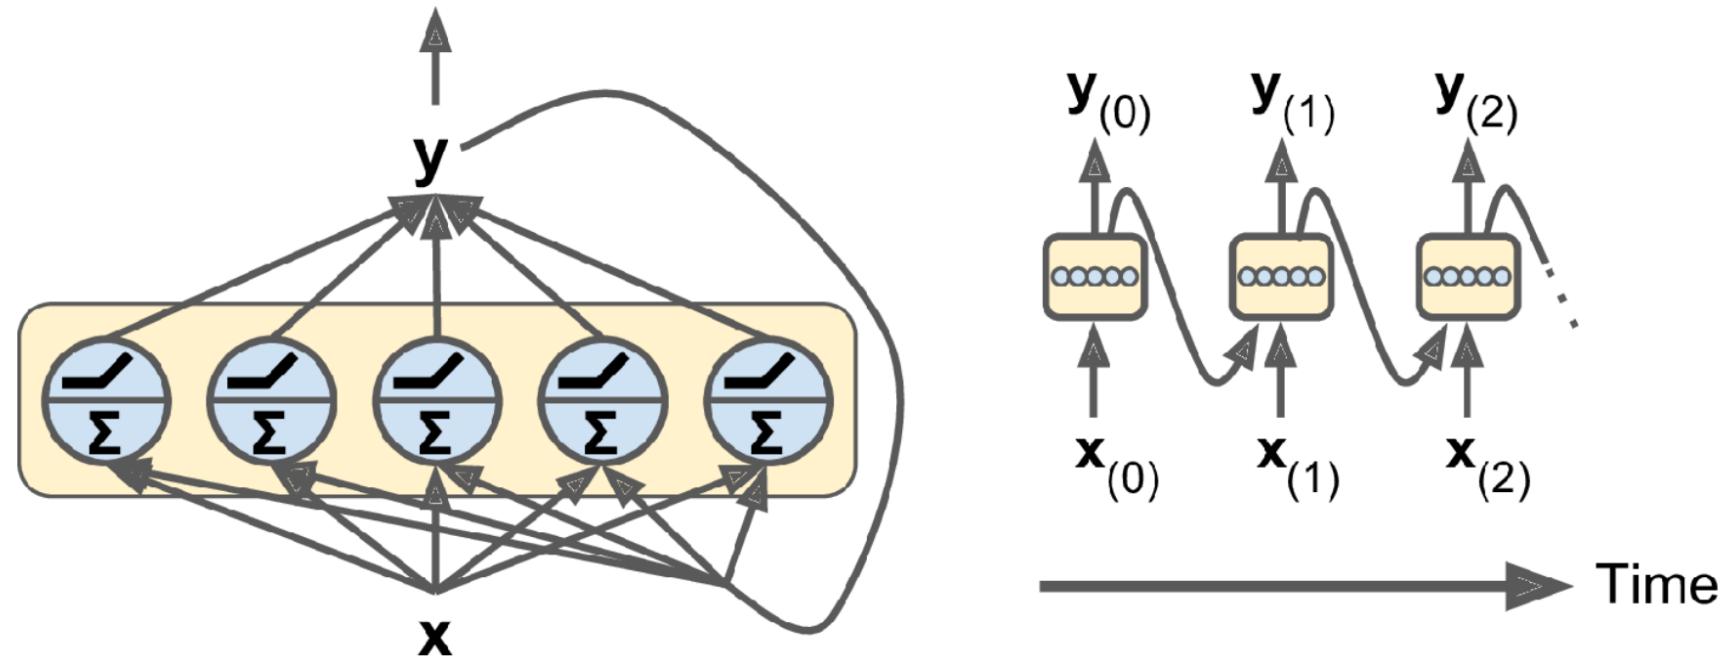
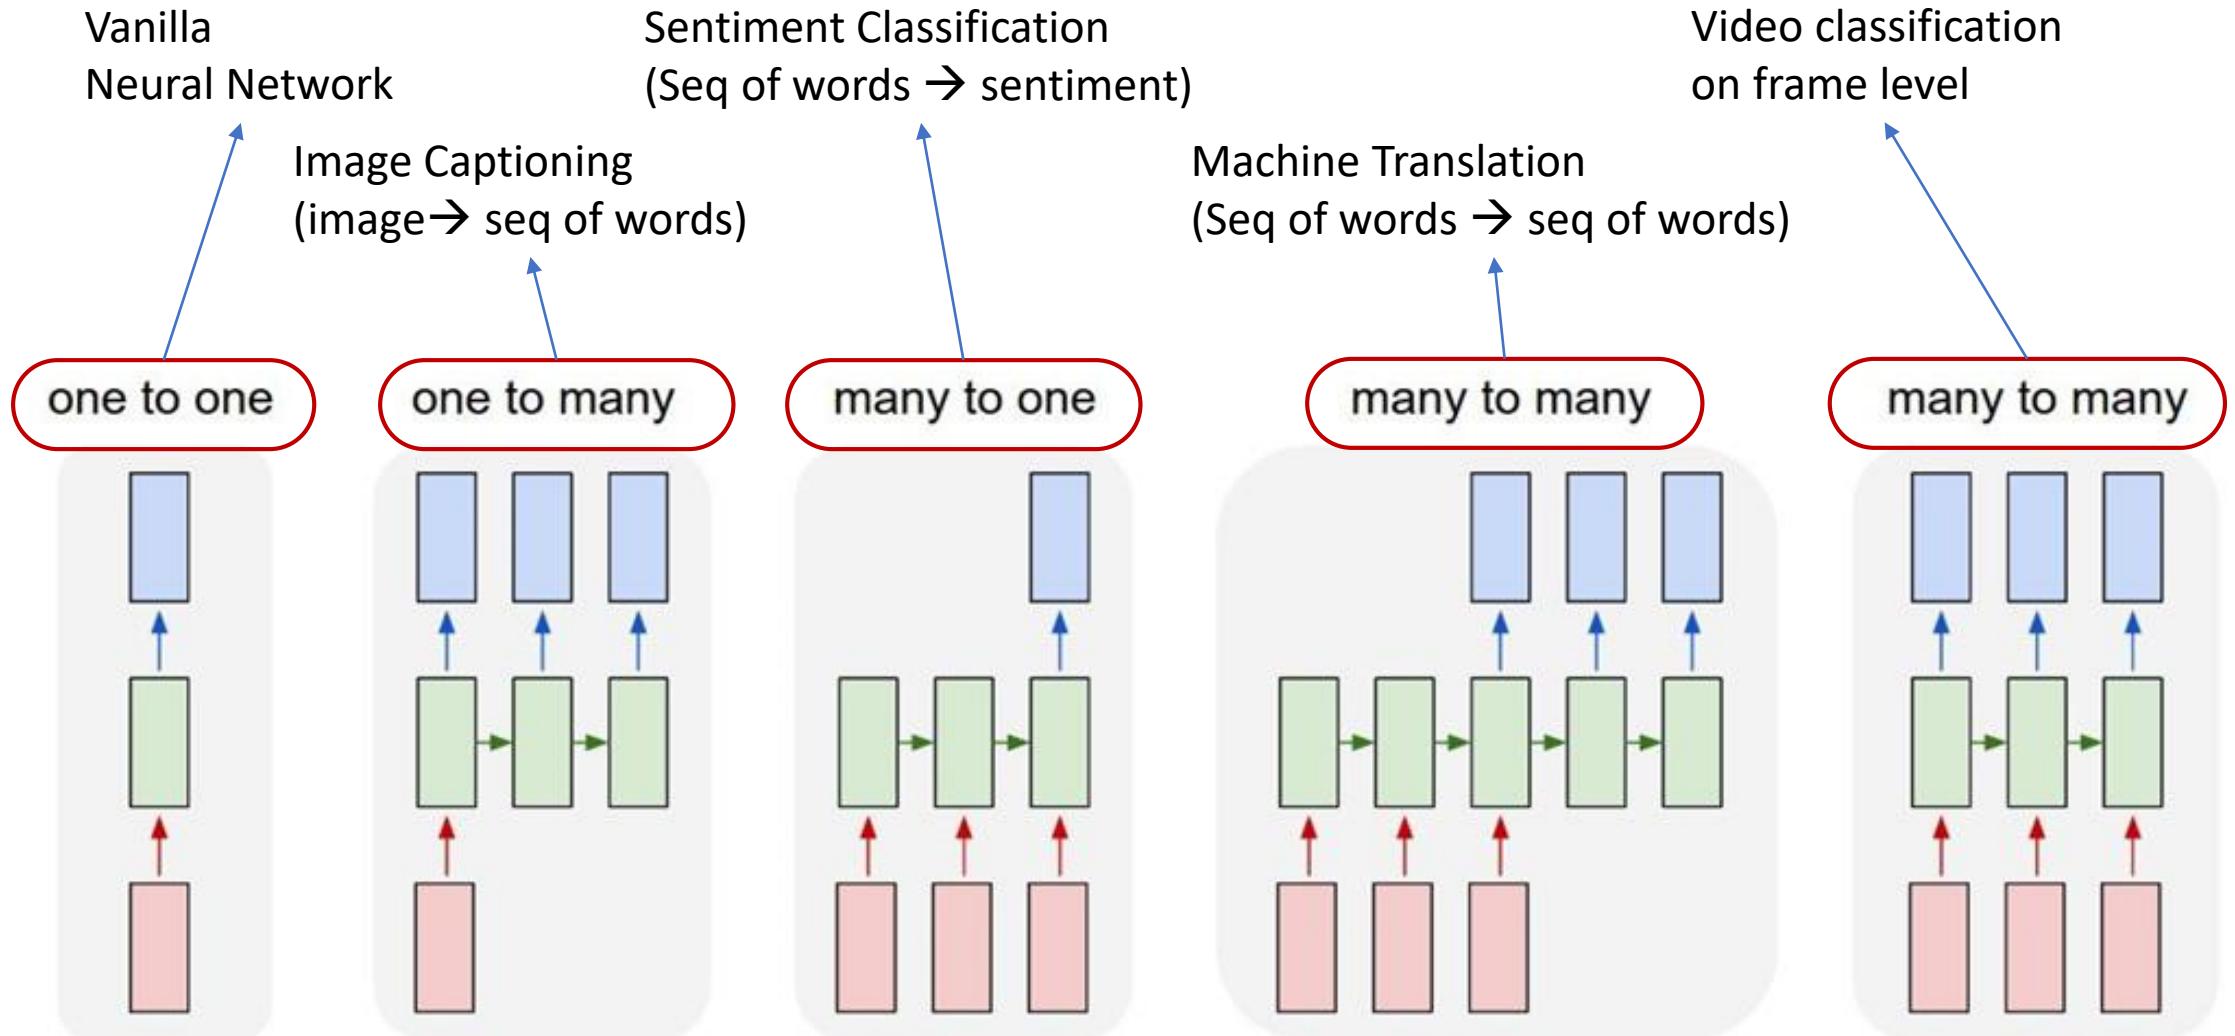


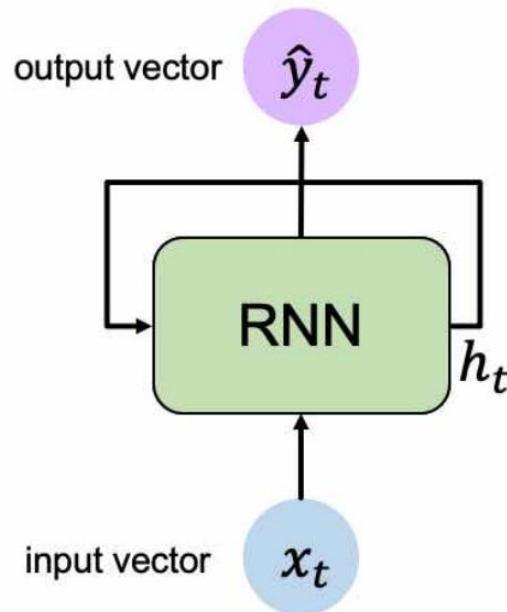
Figure 15-2. A layer of recurrent neurons (left) unrolled through time (right)

Recurrent Neural Networks: Process Sequences



How to train RNN?

RNN State Update and Output



Output Vector

$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

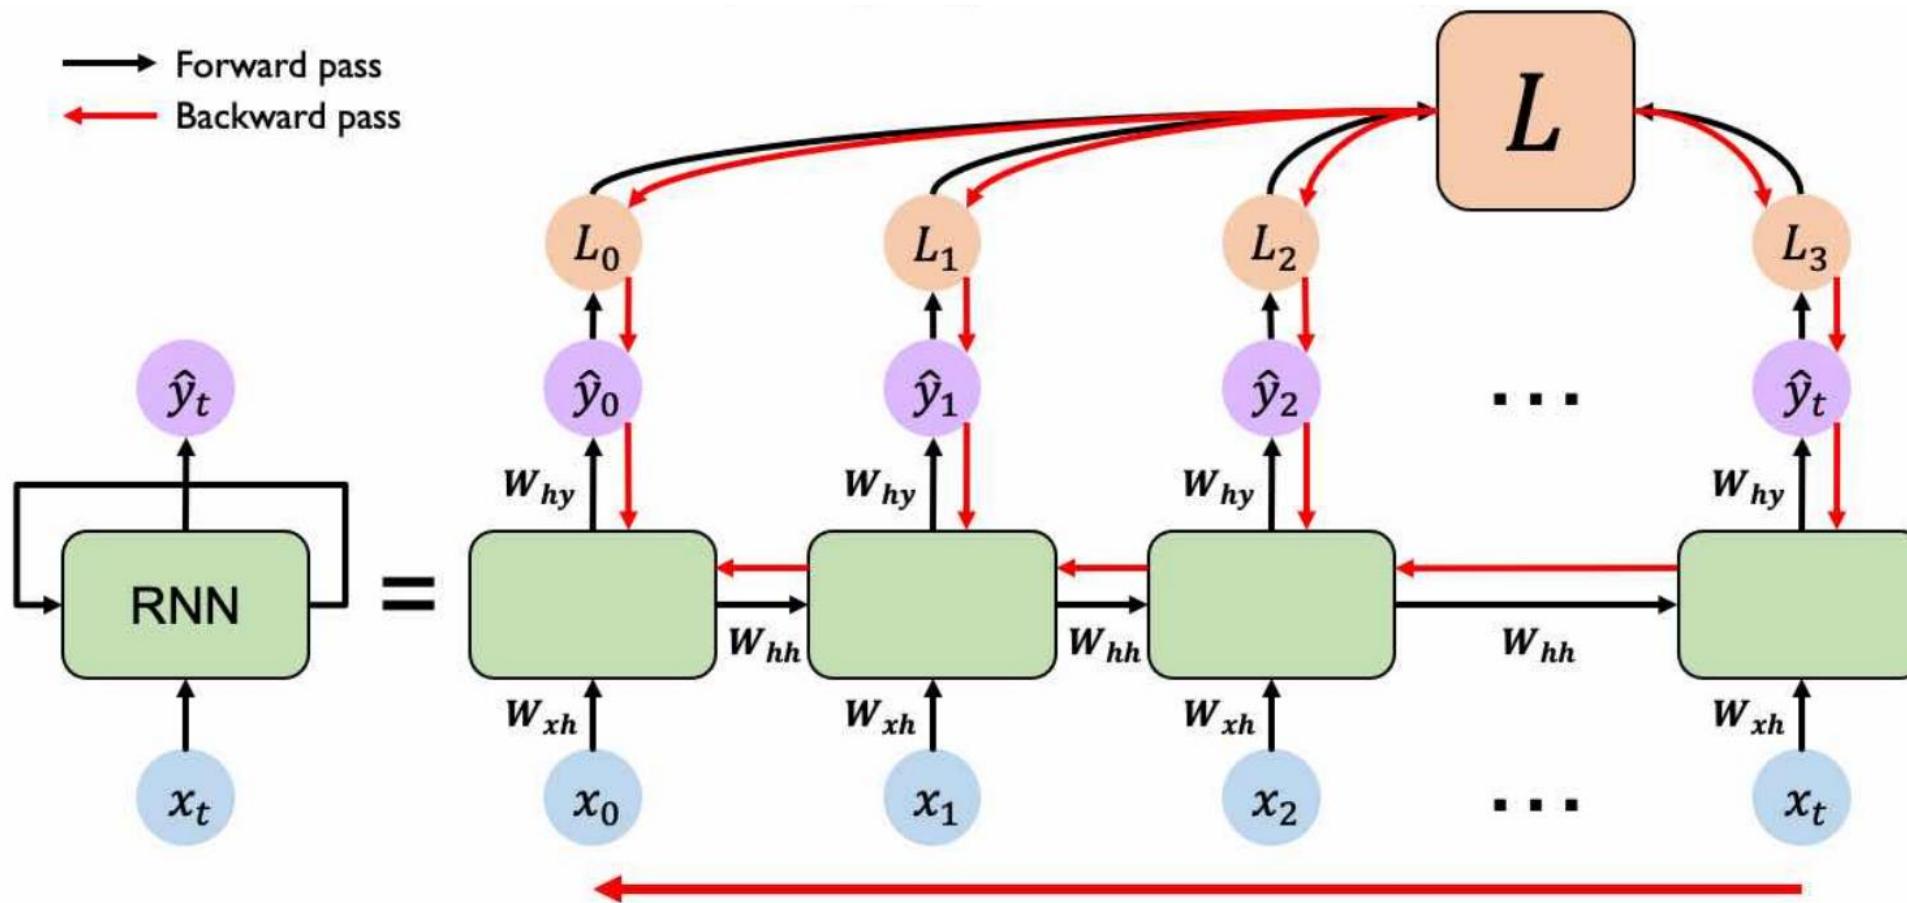
Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

Input Vector

$$x_t$$

BPTT: Backpropagation Through Time



Standard RNN Gradient Flow

$$\frac{\partial L^{(t)}}{\partial \mathbf{W}_{hh}} = \frac{\partial L^{(t)}}{\partial y^{(t)}} \cdot \frac{\partial y^{(t)}}{\partial \mathbf{h}^{(t)}} \cdot \left(\sum_{k=1}^t \frac{\partial \mathbf{h}^{(t)}}{\partial \mathbf{h}^{(k)}} \cdot \frac{\partial \mathbf{h}^{(k)}}{\partial \mathbf{W}_{hh}} \right)$$

computed as a multiplication of adjacent time steps:

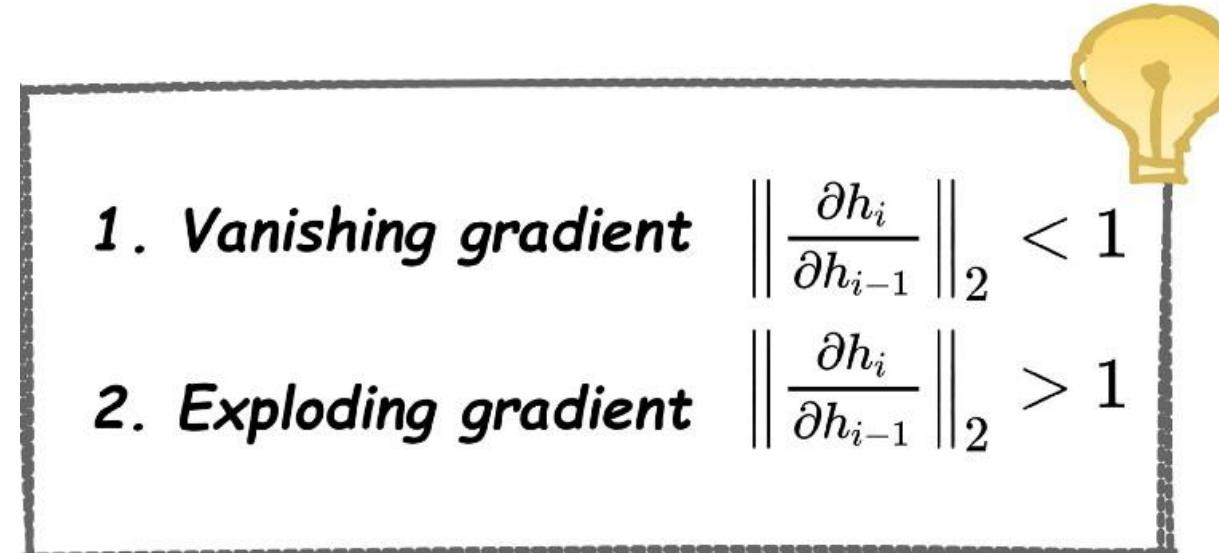
$$\frac{\partial \mathbf{h}^{(t)}}{\partial \mathbf{h}^{(k)}} = \prod_{i=k+1}^t \frac{\partial \mathbf{h}^{(i)}}{\partial \mathbf{h}^{(i-1)}}$$

❖ VGP(Vanishing Gradient Problem)

- ✓ when the gradient becomes too small, the parameter updates become insignificant. This makes the learning of long data sequences difficult.

❖ EGP(Exploding Gradient Problem)

- ✓ If the slope tends to grow exponentially instead of decaying



❖ Truncated backpropagation through time (TBPTT)

- ✓ simply limits the number of time steps the signal can backpropagate after each forward pass.
 - E.g., even if the sequence has 100 elements/steps, we may only backpropagate through 20 or so

❖ Long short-term memory (LSTM)

- ✓ uses a memory cell for modeling long-range dependencies and avoid vanishing gradient problems

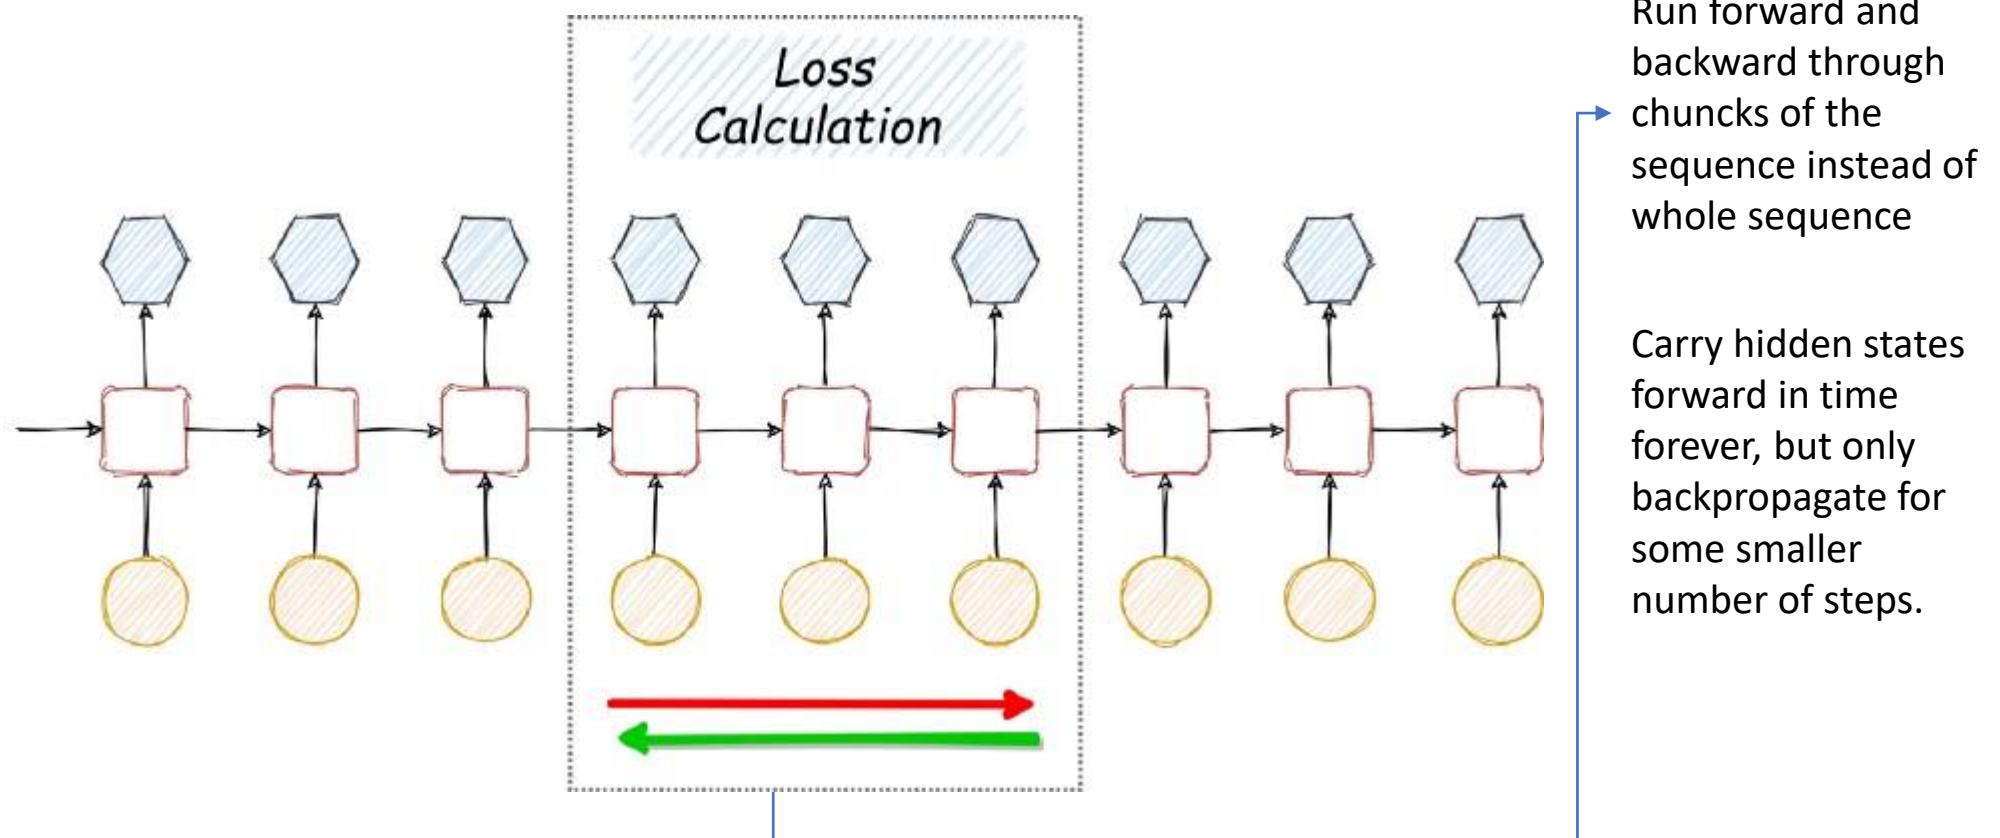
❖ Gradient Clipping

- ✓ set a max value for gradients if they grow to large
 - solves only exploding gradient problem)

(1) TBPTT: Truncated Backpropagation Through Time

❖ TBPTT

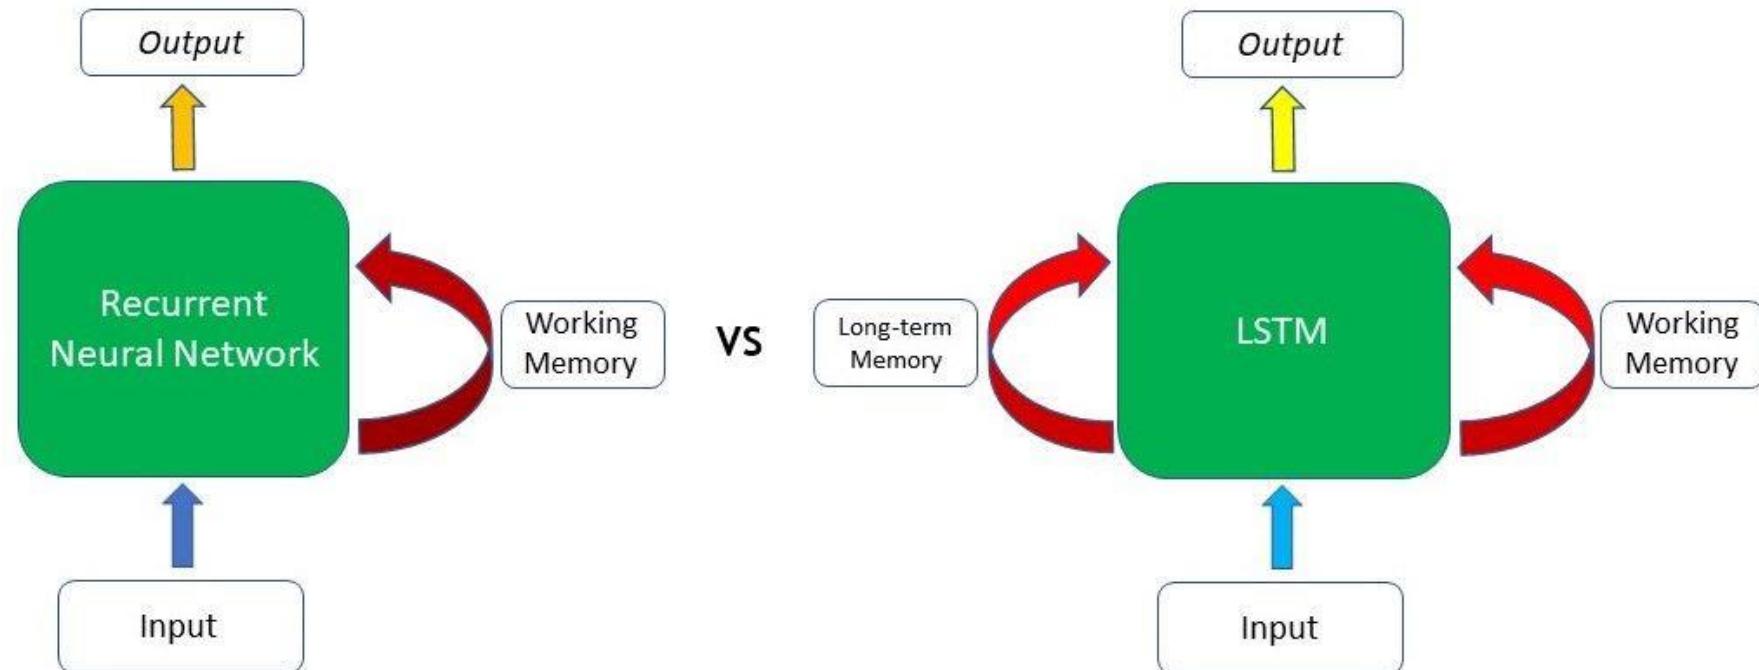
- ✓ Truncated BPTT trick tries to overcome the VGP by considering a moving window through the training process



(2) Long Short-Term Memory networks (LSTMs)

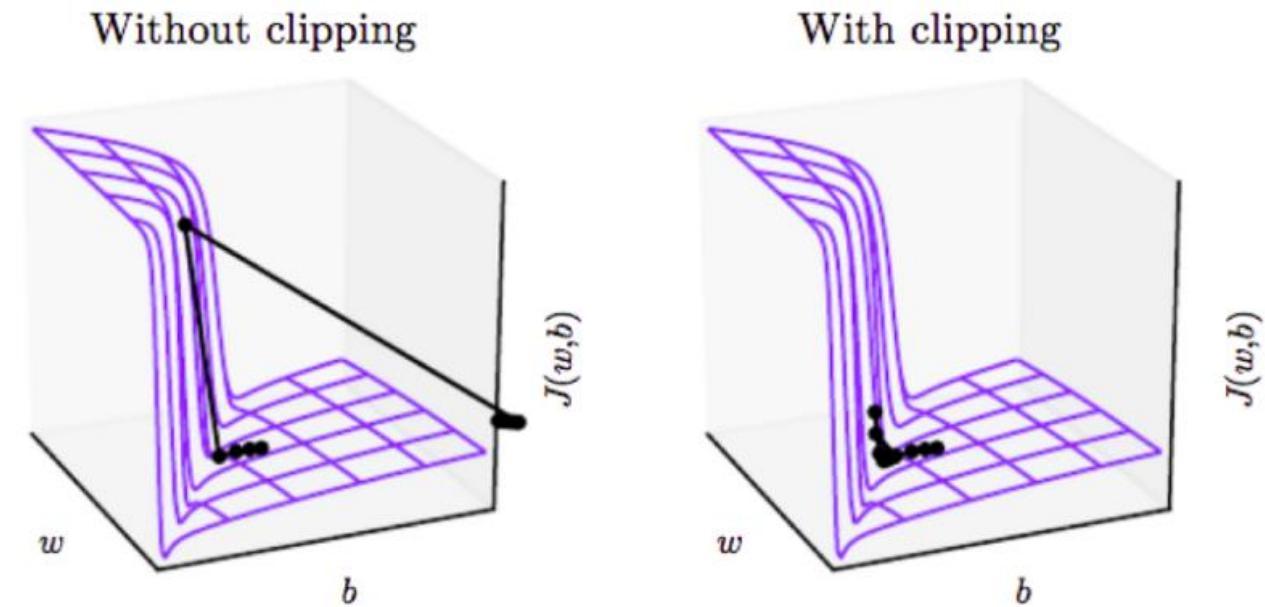
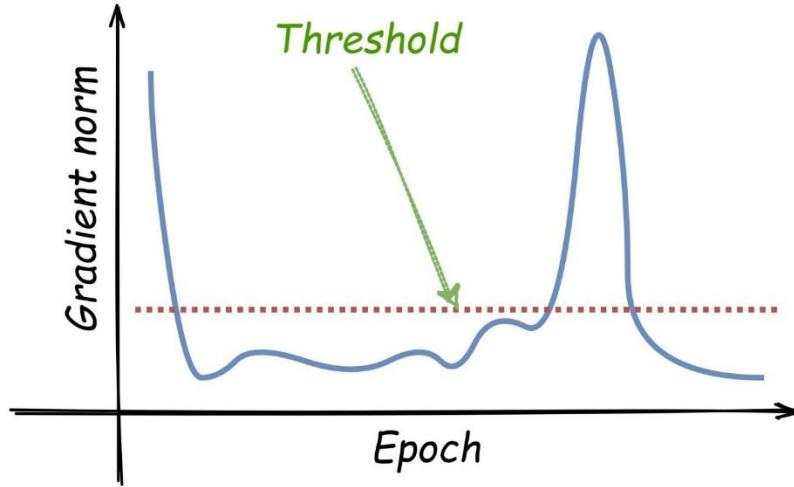
❖ Allows learning of long-term dependencies

- ✓ A RNN addresses the vanishing/exploding gradient problem
- ✓ Capable of learning long-term dependencies by remembering information



(3) Gradient Clipping

```
if  $\|\hat{g}\| \geq threshold$  then  
     $\hat{g} \leftarrow \frac{threshold}{\|\hat{g}\|} \hat{g}$   
end if
```

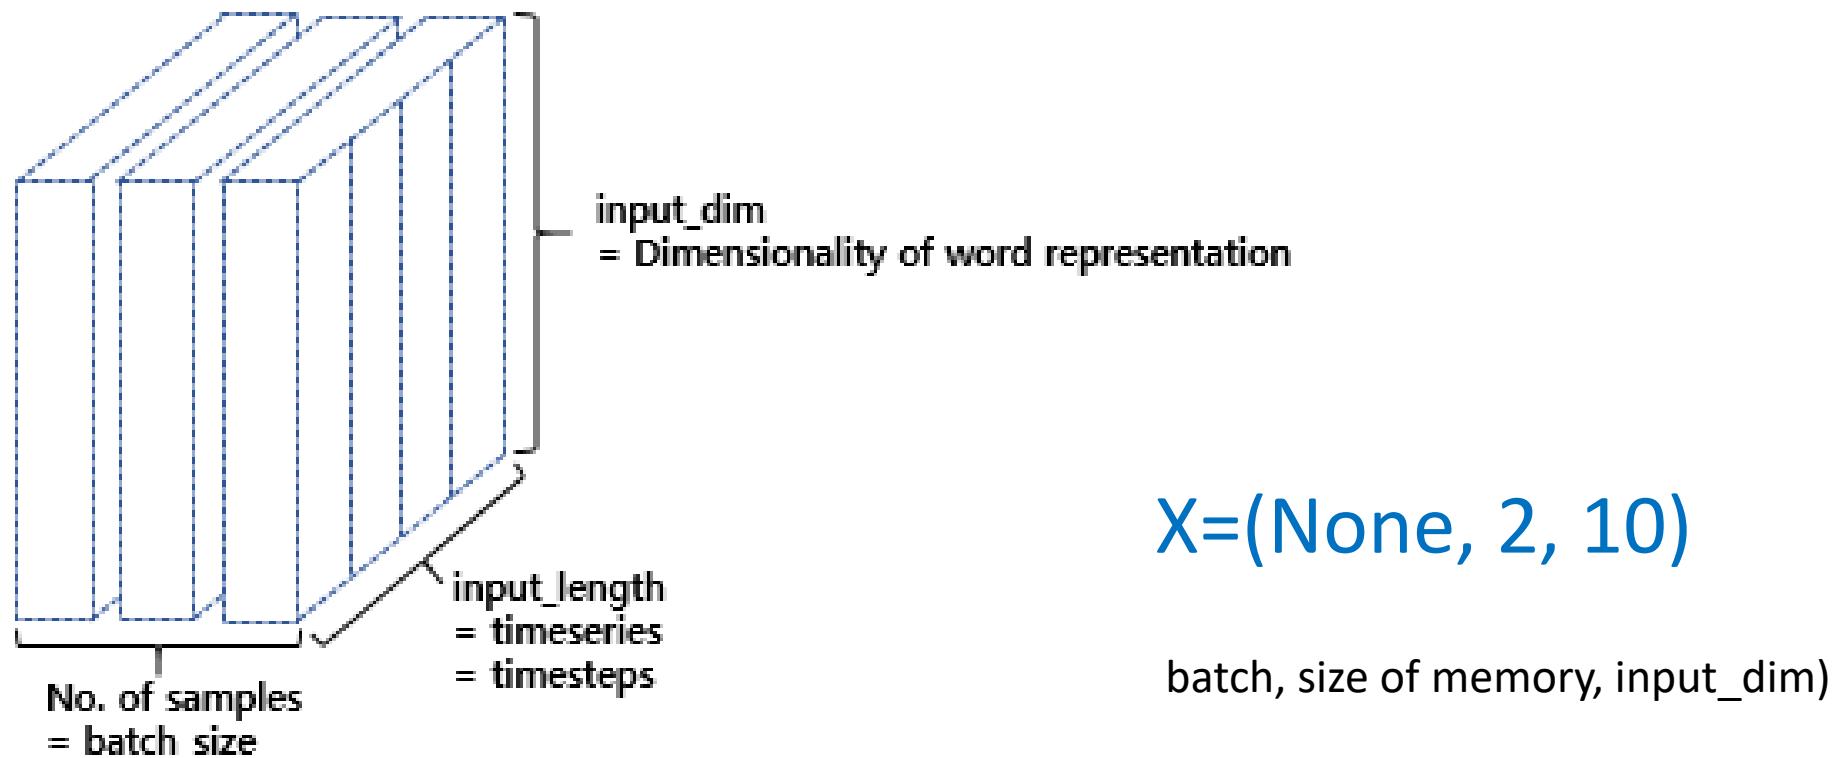


source: <https://eehoeskrap.tistory.com/582>

Implementing an RNN with Keras

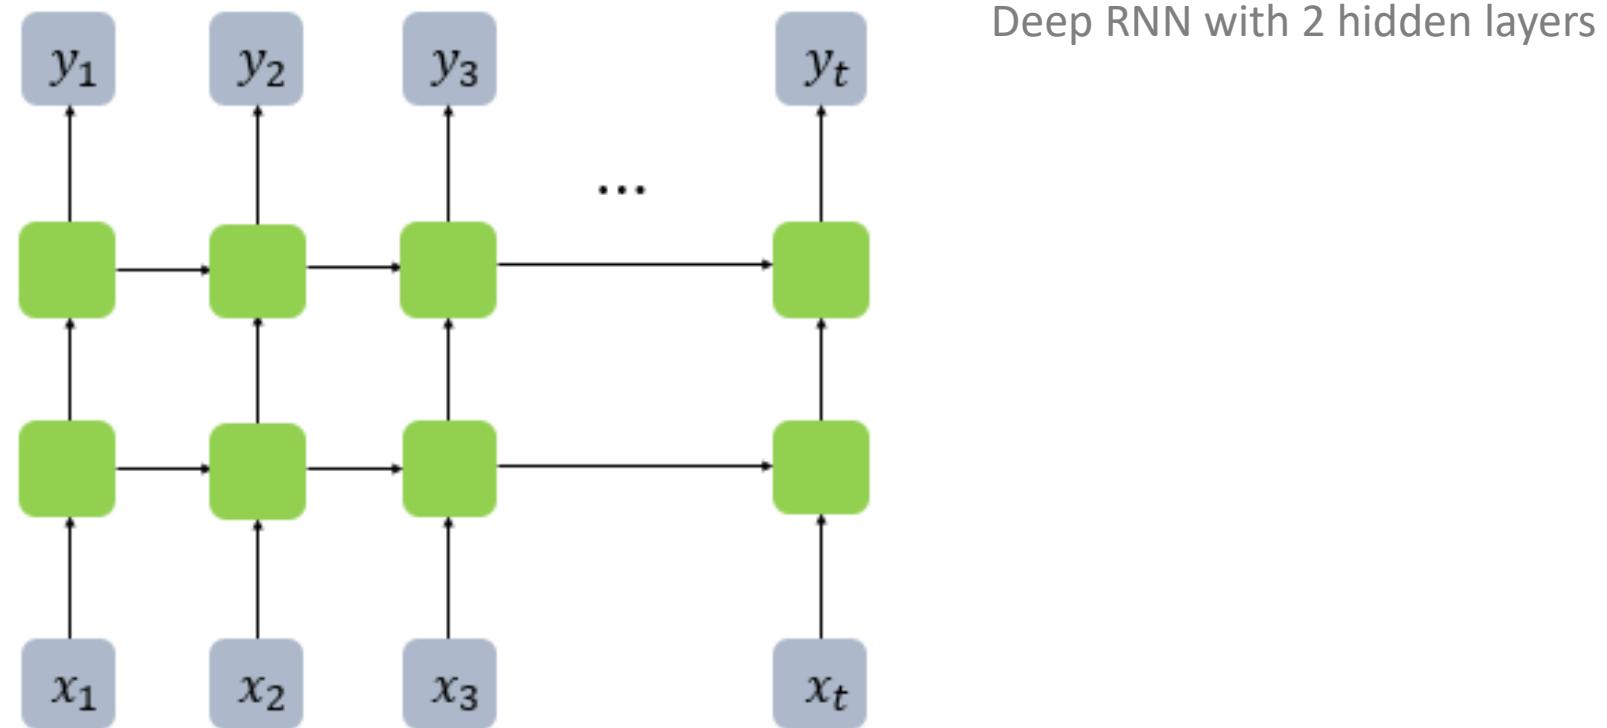
RNN Input : X Dimension

- ❖ The RNN layer receives a 3D tensor of size as input
 - ✓ batch_size, timesteps, and input_dim



A example: Deep Recurrent Neural Network

```
model.add(SimpleRNN(hidden_units, input_length=10,  
input_dim=5, return_sequences=True))
```



SimpleRNN without Batch (1/4)

output_dim : (batch_size, hidden_units)
`model.add(SimpleRNN(3, input_shape=(2,10)))`

Model: "sequential"

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 3)	42

=====
Total params: 42
Trainable params: 42
Non-trainable params: 0

X=(None, 2, 10)

SimpleRNN with batch size (2/4)

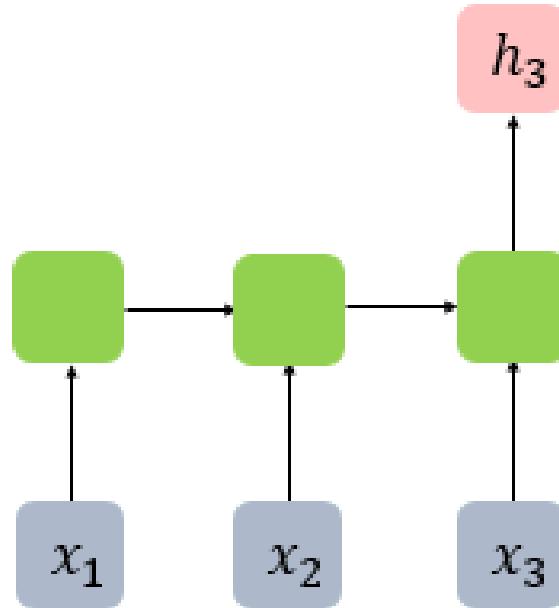
with Batch size = 8

```
model.add(SimpleRNN(3, batch_input_shape=(8,2,10)))
```

Layer (type)	Output Shape	Param #
simple_rnn_1 (SimpleRNN)	(8, 3)	42

- ❖ Setting the return_sequences of the RNN layer

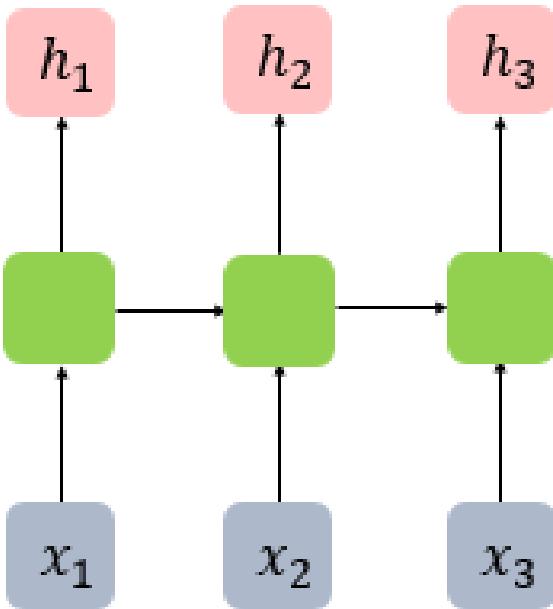
`return_sequences=False`



다음층으로 마지막 은닉 상태만 전달

many-to-one task

`return_sequences=True`



다음층으로 모든 은닉 상태 전달

many-to-many task

```
model.add(SimpleRNN(3, batch_input_shape=(8,2,10), return_sequences=True))
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
simple_rnn_2 (SimpleRNN)	(8, 2, 3)	42

[HW] Let's Code! Simple RNN

Input Data and Train Data Shape

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.keras.layers import SimpleRNN, GRU, LSTM, Bidirectional
4 train_X = [[0.1, 4.2, 1.5, 1.1, 2.8], [1.0, 3.1, 2.5, 0.7, 1.1], [0.3, 2.1, 1.5, 2.1, 0.1], [2.2, 1.4, 0.5, 0.9, 1.1]]
5 print(np.shape(train_X))
```

```
1 train_X=np.reshape(train_X, (-1,4,5))
2 print(train_X.shape)
```

(4, 5)
(1, 4, 5)

SimpleRNN and Hidden states

```
1 rnn = SimpleRNN(3)
2 # rnn = SimpleRNN(3, return_sequences=False, return_state=False)와 동일.
3 hidden_state = rnn(train_X)
4
5 print('hidden state : {}, shape: {}'.format(hidden_state, hidden_state.shape))
```

hidden state : [[0.9206875 -0.9963834 -0.99492633]], shape: (1, 3)

```
1 rnn = SimpleRNN(3, return_sequences=True)
2 hidden_states = rnn(train_X)
3
4 print('hidden states : {}, shape: {}'.format(hidden_states, hidden_states.shape))
```

hidden states : [[[[-0.91097903 -0.99876505 -0.95570076]
[0.9831649 -0.52204293 -0.94400156]
[-0.7801549 0.03199867 -0.44760188]
[0.8878632 -0.20588721 0.7986243]]], shape: (1, 4, 3)

SimpleRNN and Return Sequence

```
1 rnn = SimpleRNN(3, return_sequences=True, return_state=True)
2 hidden_states, last_state = rnn(train_X)
3
4 print('hidden states : {}, shape: {}'.format(hidden_states, hidden_states.shape))
5 print('last hidden state : {}, shape: {}'.format(last_state, last_state.shape))
```

```
hidden states : [[[0.99741   0.5095164  0.9999336 ]
 [0.97162277 0.32083237 0.99995947]
 [0.95783615 0.86253214 0.99940664]
 [0.99748254 0.7658603  0.9743646 ]]], shape: (1, 4, 3)
last hidden state : [[0.99748254 0.7658603  0.9743646 ]], shape: (1, 3)
```

Simple RNN and Return_Sequence

```
1 rnn = SimpleRNN(3, return_sequences=False, return_state=True)
2 hidden_state, last_state = rnn(train_X)
3
4 print('hidden state : {}, shape: {}'.format(hidden_state, hidden_state.shape))
5 print('last hidden state : {}, shape: {}'.format(last_state, last_state.shape))
```

```
hidden state : [[0.9096651 0.9990103 0.7855833]], shape: (1, 3)
last hidden state : [[0.9096651 0.9990103 0.7855833]], shape: (1, 3)
```

2022

Korea Institute of Science
and Technology Information

TRUST
KISTI

