

## Lec 12: Sequence to Sequence Model with RNN



hsyi@kisti.re.kr

Hongsuk Yi (이홍석)



- ❖ Introduction to Recurrent Neural Network
  - ✓ Simple RNN, BPTT, Memory Cell
  - ✓ Code: Implementing an RNN with Keras
- ❖ Introduction to Long-Short Term Memory
  - ✓ Cell state, LSTM, and GRU, and Applications
  - ✓ A Visual Guide to Recurrent Layers in Keras
  - ✓ Code: A simple LSTM layers
- ❖ Text generation with RNN
  - ✓ Tokenizer, Character-Level Language model
  - ✓ Code: Alice's Adventures in Wonderland
- ❖ **Sequence to Sequence Learning model with RNN**
  - ✓ Introduction to Seq2Seq and Attention model
  - ✓ Code: Character-Level Neural Machine Translation

# Reviewing the last class:

## Character-level language model

## “Modeling word probabilities is really difficult”

	Supervised learning	Sequence modelling
Data	$\{x, y\}_i$	$\{x\}_i$
Model	$y \approx f_{\theta}(x)$	$p(x) \approx f_{\theta}(x)$
Loss	$\mathcal{L}(\theta) = \sum_{i=1}^N l(f_{\theta}(x_i), y_i)$	$\mathcal{L}(\theta) = \sum_{i=1}^N \log p(f_{\theta}(x_i))$
Optimisation	$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta)$	$\theta^* = \arg \max_{\theta} \mathcal{L}(\theta)$

## Simplest model:

Assume independence of words

$$p(\mathbf{x}) = \prod_{t=1}^T p(x_t)$$

$$p(\text{"modeling"}) \times p(\text{"word"}) \times p(\text{"probabilities"}) \times p(\text{"is"}) \times p(\text{"really"}) \times p(\text{"difficult"})$$

Word	$p(x_t)$
the	0.049
be	0.028
...	...
really	0.0005
...	...

## More realistic model:

Assume conditional dependence of words

$$p(x_T) = p(x_T | x_1, \dots, x_{T-1})$$

Modeling word probabilities is really

?

Context

Target

$p(x|\text{context})$

difficult

0.01

hard

0.009

fun

0.005

...

...

easy

0.00001



## The chain rule

Computing the joint  $p(\mathbf{x})$  from conditionals

$$p(\mathbf{x}) = \prod_{t=1}^T p(x_t | x_1, \dots, x_{t-1})$$

## Modeling

Modeling **word**

Modeling word **probabilities**

Modeling word probabilities **is**

Modeling word probabilities is **really**

Modeling word probabilities is really **difficult**

$$p(x_1)$$

$$p(x_2 | x_1)$$

$$p(x_3 | x_2, x_1)$$

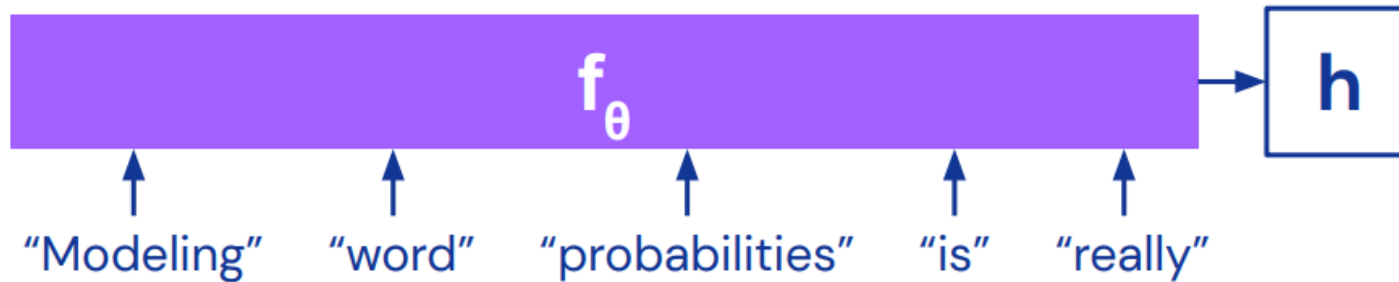
$$p(x_4 | x_3, x_2, x_1)$$

$$p(x_5 | x_4, x_3, x_2, x_1)$$

$$p(x_6 | x_5, x_4, x_3, x_2, x_1)$$

## Learning to model word probabilities

✓ Vectorising the context



$f_{\theta}$  summarises the context in  $h$  such that:

$$p(x_t | x_1, \dots, x_{t-1}) \approx p(x_t | h)$$

Desirable properties for  $f_{\theta}$ :

- Order matters
- Variable length
- Learnable (differentiable)

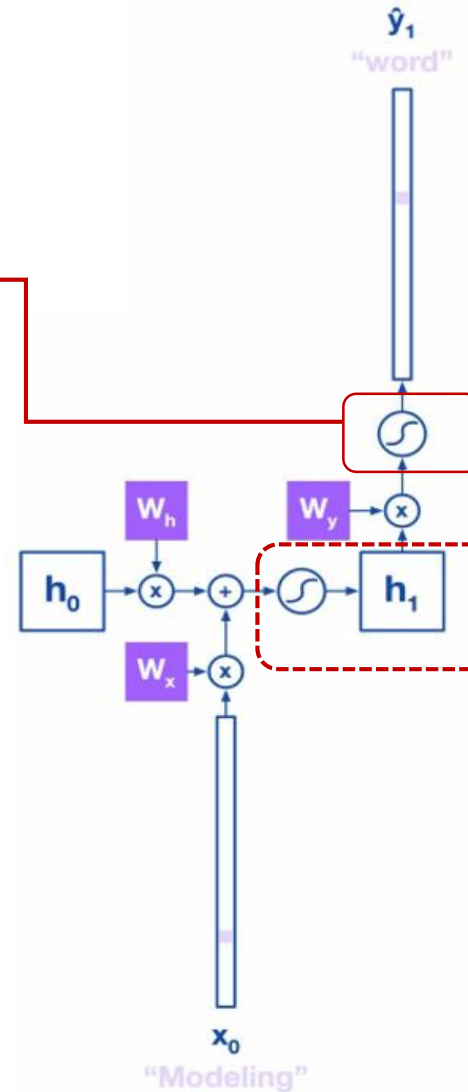


RNNs predict the target  $\mathbf{y}$  (the word) from the state  $\mathbf{h}$ .

Persistent state variable  $\mathbf{h}$  stores information from the context observed so far

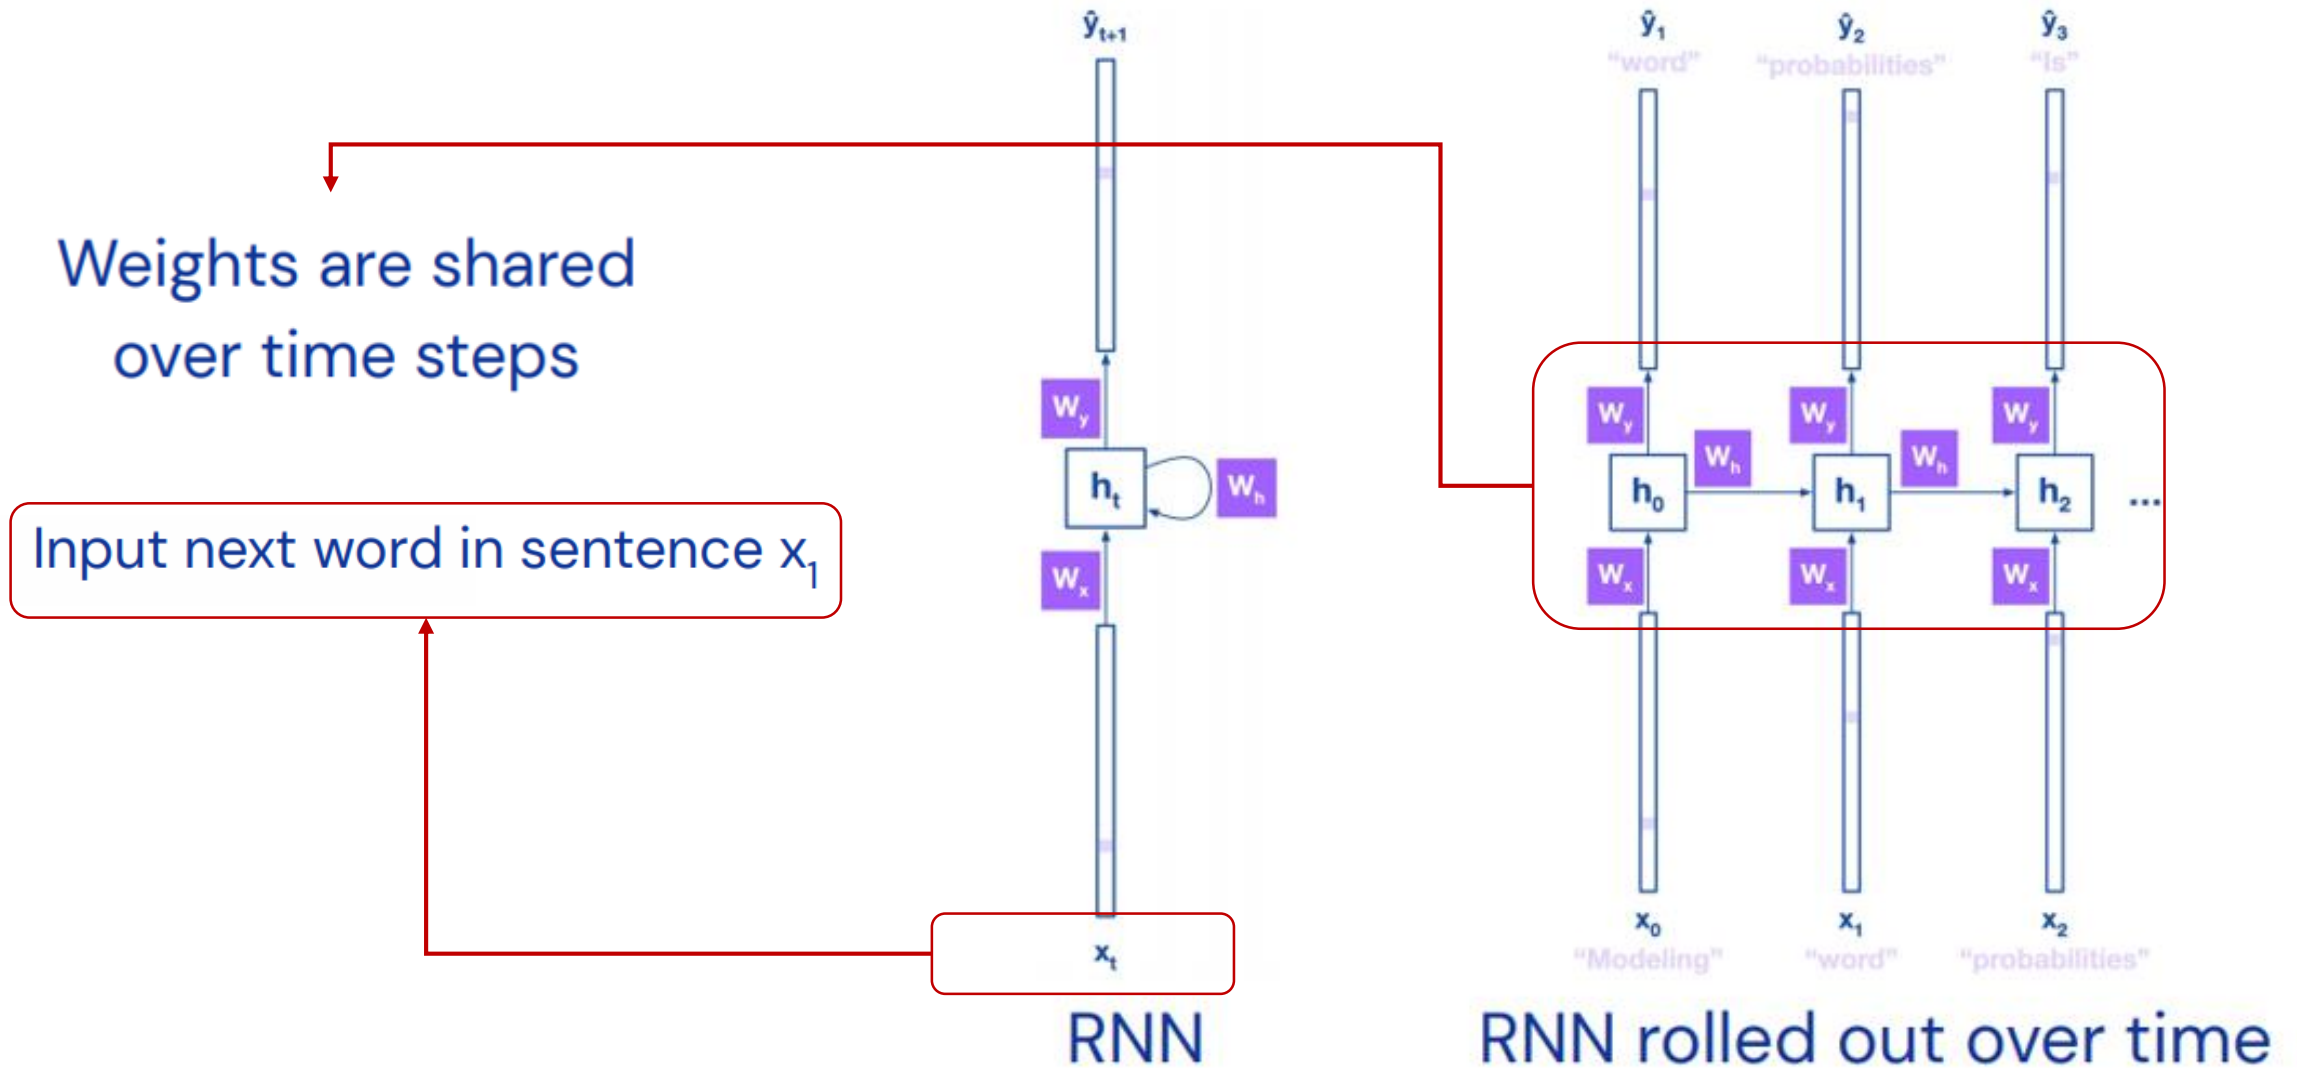
$$p(\mathbf{y}_{t+1}) = \text{softmax}(\mathbf{W}_y \mathbf{h}_t)$$

Softmax ensures we obtain a distribution over all possible words.



$$\mathbf{h}_t = \tanh(\mathbf{W}_h \mathbf{h}_{t-1} + \mathbf{W}_x \mathbf{x}_t)$$

# Recurrent Neural Networks (RNNs)



Next word prediction is essentially a classification task where the number of classes is the size of the vocabulary.


As such we use the cross-entropy loss:

For one word:  $\mathcal{L}_{\theta}(\mathbf{y}, \hat{\mathbf{y}})_t = -\mathbf{y}_t \log \hat{\mathbf{y}}_t$

For the sentence:


$$\mathcal{L}_{\theta}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{t=1}^T \mathbf{y}_t \log \hat{\mathbf{y}}_t$$

With parameters  $\theta = \{\mathbf{W}_y, \mathbf{W}_x, \mathbf{W}_h\}$




$$\mathbf{h}_t = \tanh(\mathbf{W}_h \mathbf{h}_{t-1} + \mathbf{W}_x \mathbf{x}_t)$$
$$p(\mathbf{x}_{t+1}) = \text{softmax}(\mathbf{W}_y \mathbf{h}_t)$$
$$\mathcal{L}_\theta(\mathbf{y}, \hat{\mathbf{y}})_t = -\mathbf{y}_t \log \hat{\mathbf{y}}_t$$

$$\frac{\partial \mathbf{L}}{\partial W} = \sum_{i=0}^T \frac{\partial \mathcal{L}_i}{\partial W} \propto \sum_{i=0}^T \left( \prod_{i=k+1}^y \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W}$$

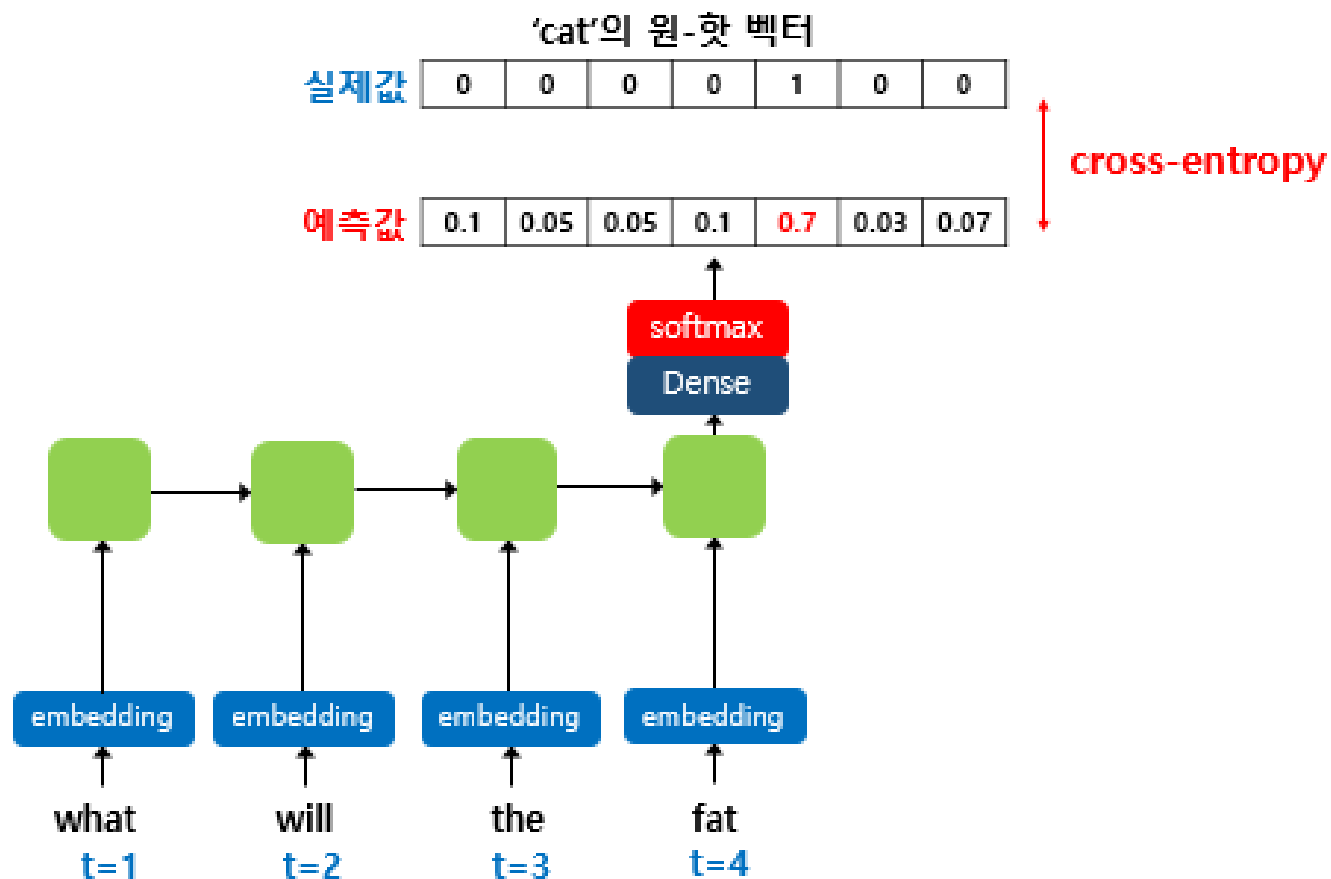


- ❖ A model that predicts the next word from a word sequence


예문 : 'what will the fat cat sit on'







- ❖ The model does not use the predicted value at point  $t$  as input at point  $t+1$ , but uses the label at point  $t+1$  as input at point  $t+1$





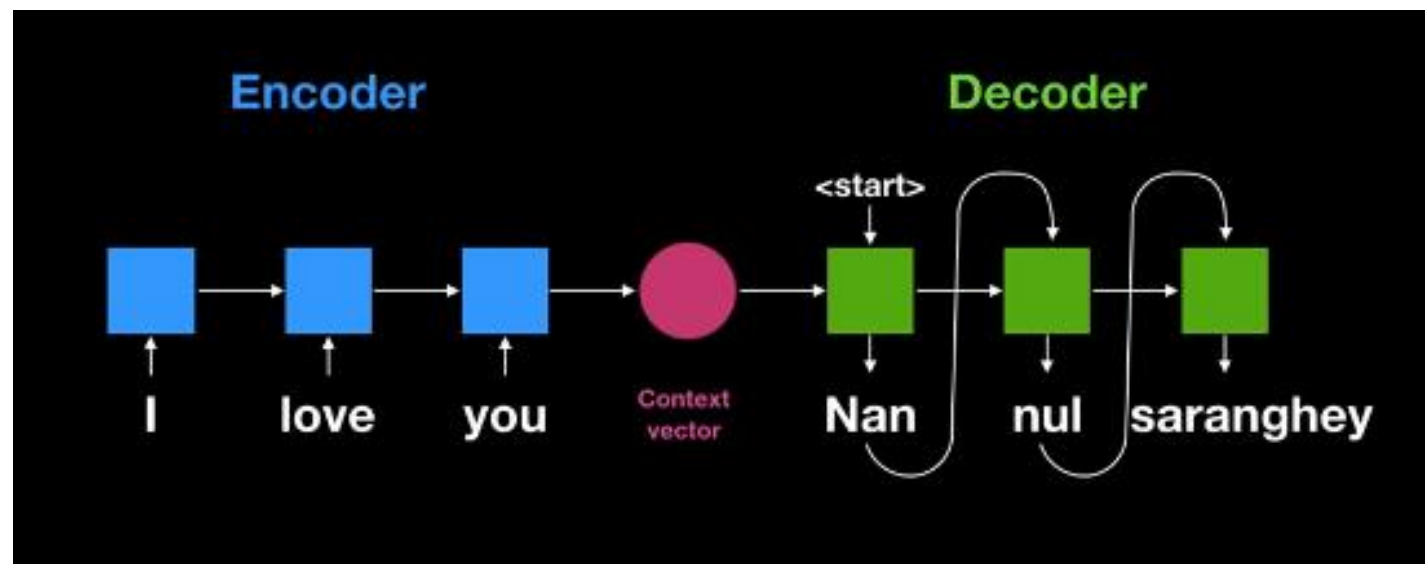


# Sequence to Sequence Neural Machine Translation



- ❖ The seq2seq model compresses the input sequence into one fixed-size vector representation, called the context vector, through which the decoder produces the output sequence.

context vector: the final RNN cell states "I love you".



(source) <https://www.kaggle.com/code/jeongwonkim10516/attention-mechanism-for-nlp-beginners/notebook>

## ❖ Sequence-to-sequence learning (Seq2Seq)


- ✓ Seq2Seq is about training models to convert sequences from one domain (e.g. sentences in English) to sequences in another domain (e.g. the same sentences translated to French).

```
"the cat sat on the mat" -> [Seq2Seq model] -> "le chat etait assis sur le tapis"
```

- ✓ This can be used for machine translation or for free-form question answering (generating a natural language answer given a natural language question)
- ✓ in general, it is applicable any time you need to generate text

## ❖ machine translation

- ✓ input sequences and output sequences have different lengths



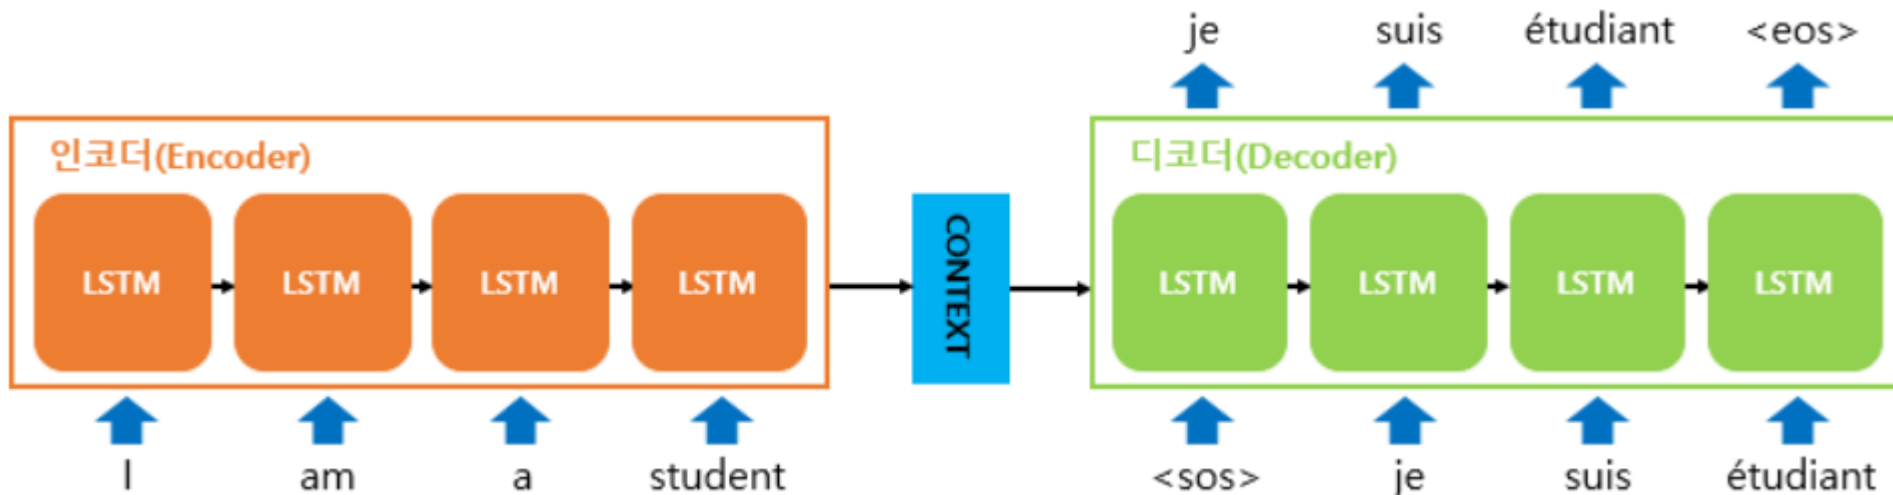


## ❖ A Seq2Seq model usually consists of:

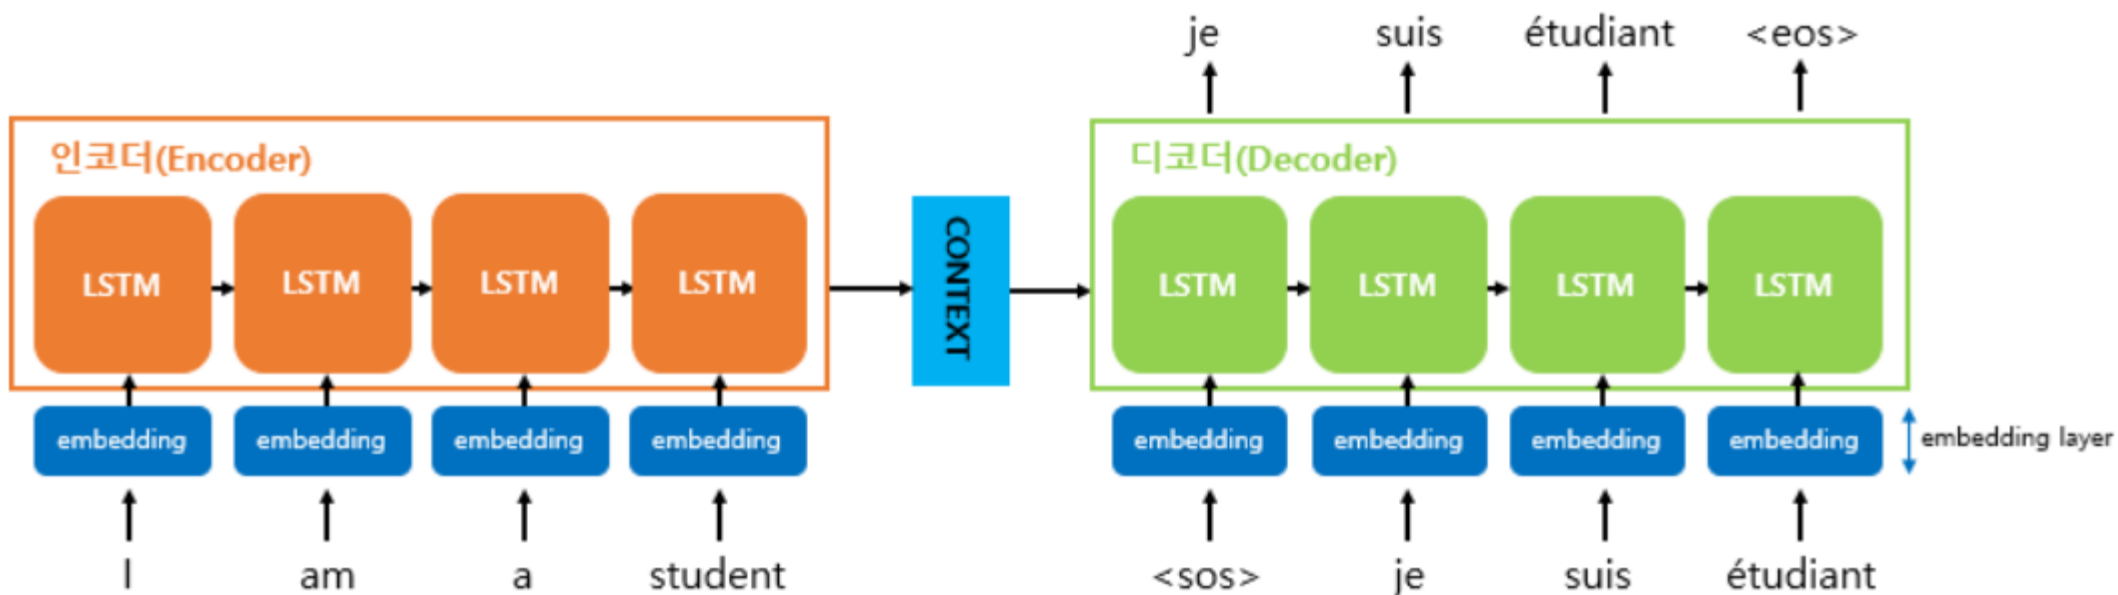
- ✓ an Encoder
  - The **encoder** processes all the inputs by transforming them into a single vector, called **context** (usually with a length of 256, 512, or 1024).
- ✓ a Decoder
  - The context contains all the information that the encoder was able to detect from the input.
- ✓ a Context (*vector*)
  - Finally, the vector is sent to the **decoder** which formulates the output sequence.

## ❖ context vector

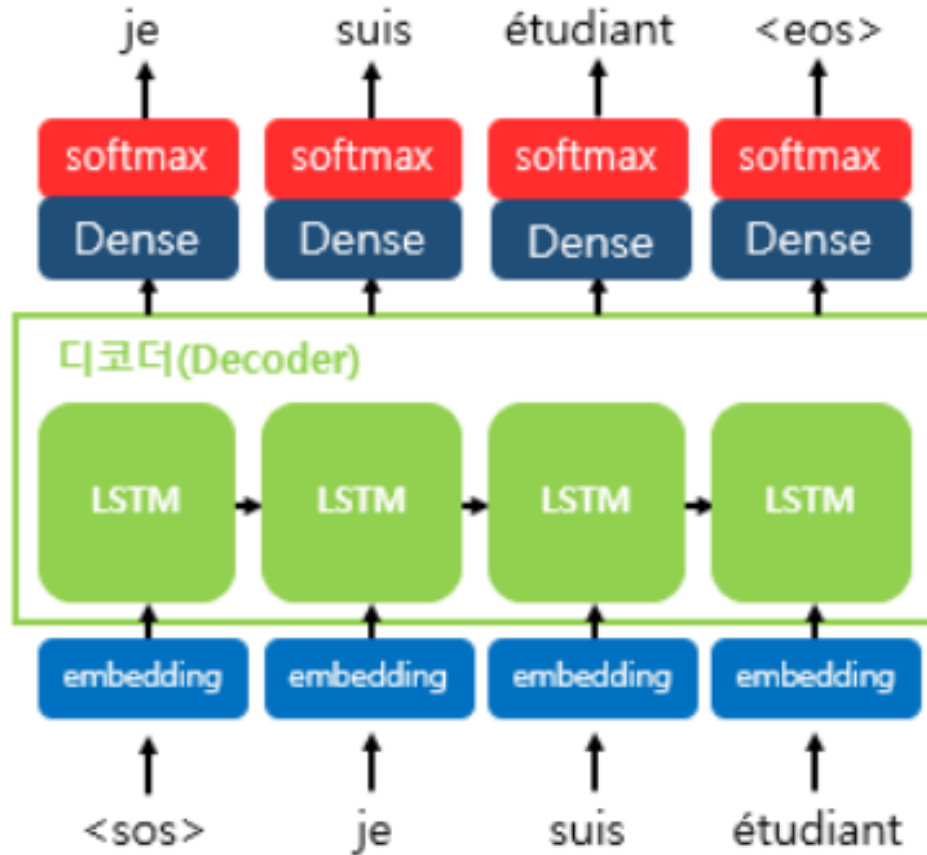
- ✓ The context vector is the first hidden state of the decoder RNN cell



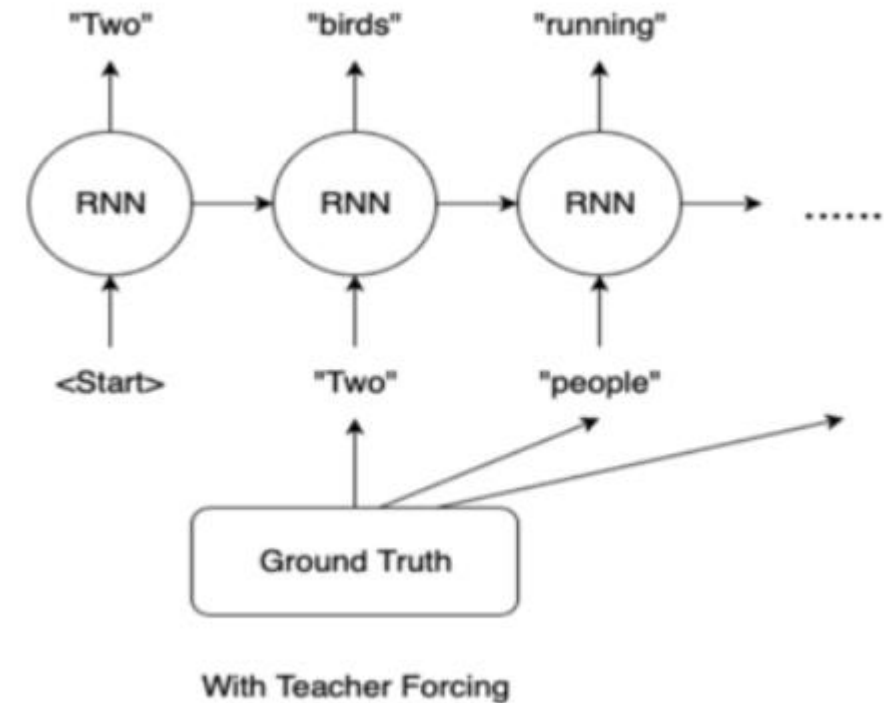
- ✓ Decoder is essentially an RNNLM (RNN Language Model)
  - Many-to-Many



## ❖ Softmax for next prediction word



## Teacher Forcing Learning

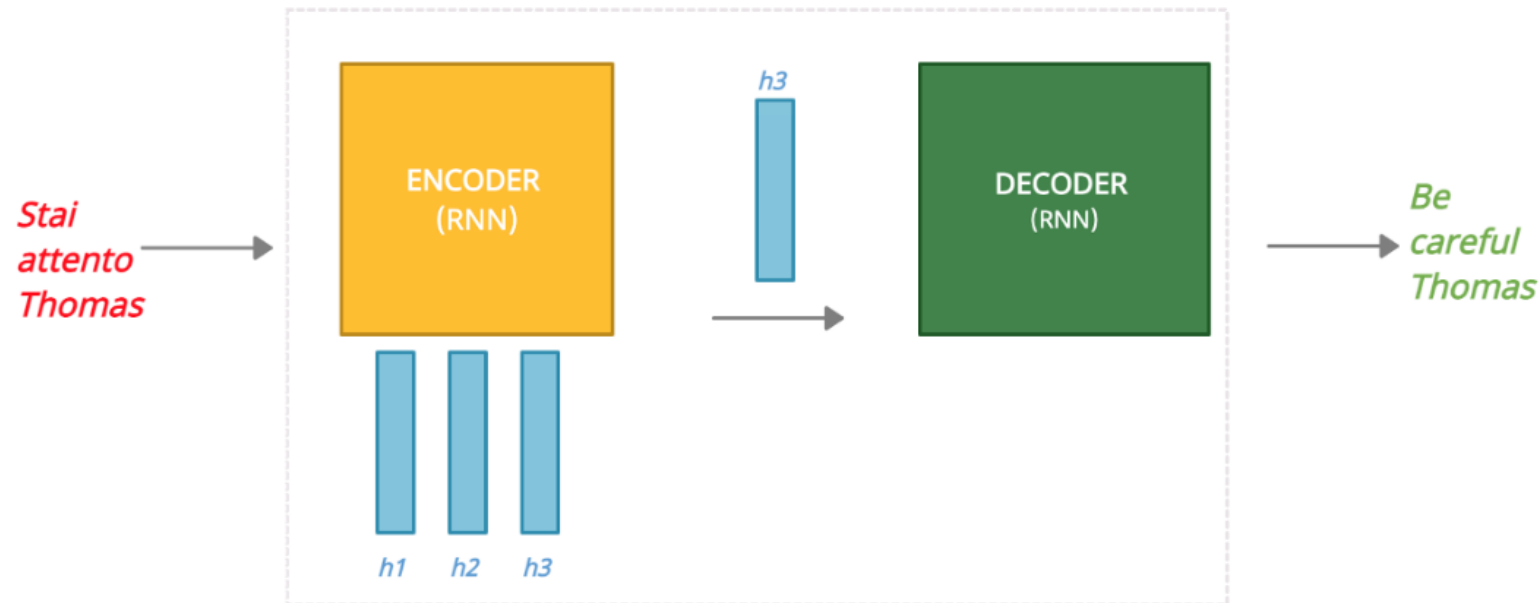


## ❖ Main problem with seq2seq models

- ✓ compress all the information into one fixed-size vector results in information loss.

## ❖ This is the problem that attention solves!

- ✓ The last **hidden state ( $h_3$ )** becomes the content that is sent to the decoder
- ✓ the encoder is “forced” to send only **a single vector**, regardless of the length of our input



## ❖ Attention mechanism proposed by Bahdanau et al. (2015)

- ✓ We compute a "summary" (weighted average) of the states which correspond to some notion of "importance"




image borrowed from [Richard Johansson](#) (Chalmers Technical University and University of Gothenburg)



## ❖ A general formulation

- ✓ for the attention weights, we apply the softmax
- ✓ the "summary" is computed as a weighted sum

$$\alpha_j = \frac{\exp e_j}{\sum_{j=1}^n \exp e_j}$$

“importance score”

attention weights

$$s = \sum_{i=1}^n \alpha_i h_i$$

“summary”

each RNN state  $h_i$

image borrowed from [Richard Johansson](#) (Chalmers Technical University and University of Gothenburg)

# **[HW4] Let's Code!**

## **Character-Level Neural Machine Translation**

## ❖ Seq2Seq

✓ <https://blog.keras.io/a-ten-minute-introduction-to-sequence-to-sequence-learning-in-keras.html>

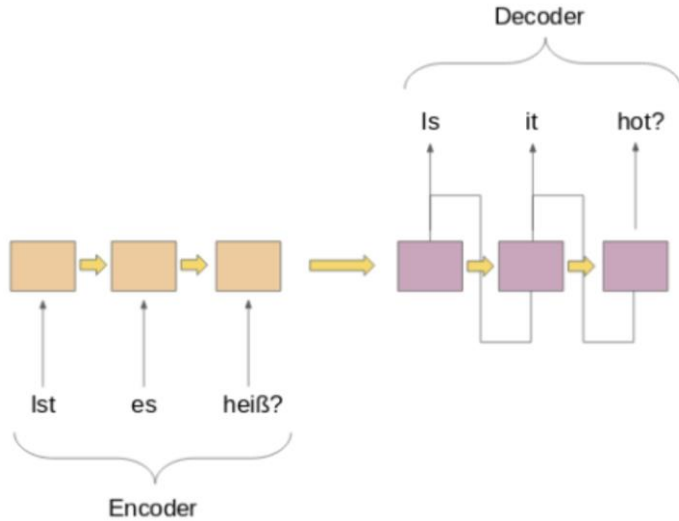
❖ Seq2seq may basically have different lengths of the input sequence and the output sequence

## ❖ Corpus with two or more languages in parallel

✓ <http://www.manythings.org/anki>

- fra-eng.zip
- kor-eng.zip

```
In [2]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Embedding, Bidirectional,
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import load_model
from tensorflow.keras import optimizers
import matplotlib.pyplot as plt
```



```
In [5]: data = read_text("kor.txt")
kor_eng = to_lines(data)
kor_eng = np.array(kor_eng)
print(len(kor_eng))
kor_eng.shape
```

3729

Out [5]: (3729, 3)

## (b) Text to Sequence Conversion

```
kor_eng
```

```
인 것 같아.],
```

```
'CC-BY 2.0 (France) Attribution: tatoeba.org #953635 (CK) & #8384140 (Eunhee)],
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3729 entries, 0 to 3728
```

```
Data columns (total 2 columns):
```

```
#   Column  Non-Null Count  Dtype
```

```
---  ---
```

```
0   eng     3729 non-null    int64
```


```
1   kor     3729 non-null    int64
```

```
dtypes: int64(2)
```

```
memory usage: 58.4 KB
```

```
df.plot.hist(bins = 300,figsize=(10, 6))
```

```
plt.show()
```





- ✓ The maximum length of the Korean sentences is 15 and that of the English is 20.

```
In [15]: # function to build a tokenizer  
def tokenization(lines):  
    tokenizer = Tokenizer()  
    tokenizer.fit_on_texts(lines)  
    return tokenizer
```

```
In [16]: # prepare english tokenizer  
eng_tokenizer = tokenization(kor_eng[:, 0])  
eng_vocab_size = len(eng_tokenizer.word_index) + 1  
  
eng_length = 20  
  
print('English Vocabulary Size: %d' % eng_vocab_size)
```

English Vocabulary Size: 2561

In [18]:

```
# encode and pad sequences
def encode_sequences(tokenizer, length, lines):
    # integer encode sequences
    seq = tokenizer.texts_to_sequences(lines)
    # pad sequences with 0 values
    seq = pad_sequences(seq, maxlen=length, padding='post')
    return seq
```

We will now split the data into train and test set for model training and evaluation, respectively.

In [20]:

```
from sklearn.model_selection import train_test_split
train, test = train_test_split(kor_eng, test_size=0.2, random_state = 12)
```

```
# build NMT model
```

```
def build_model(in_vocab, out_vocab, in_timesteps, out_timesteps, units):
    model = Sequential()
    model.add(Embedding(in_vocab, units, input_length=in_timesteps, mask_zero=True))
    model.add(LSTM(units))
    model.add(RepeatVector(out_timesteps))
    model.add(LSTM(units, return_sequences=True))
    model.add(Dense(out_vocab, activation='softmax'))
    return model
```

We are using RMSprop optimizer in this model as it is usually a good choice for recurrent neural networks.

```
model = build_model(kor_vocab_size, eng_vocab_size, kor_length, eng_length, 64)
rms = optimizers.RMSprop(learning_rate=0.001)
model.compile(optimizer=rms, loss='sparse_categorical_crossentropy', metrics=['acc'])
```

loss='sparse\_categorical\_crossentropy': because it allows us to use the target sequence as it is instead of one hot encoded format.

- One hot encoding the target sequences with such a huge vocabulary might consume our system's entire memory.
- We will train it for 30 epochs and with a batch size of 64.

```
history = model.fit(trainX, trainY.reshape(trainY.shape[0], trainY.shape[1], 1),  
                    epochs=30, batch_size=64,  
                    validation_split = 0.2, verbose=1)
```

```
38/38 [=====] - 1s 22ms/step - loss: 1.5617 - acc:  
1.8888 - val_acc: 0.7428
```


Let's compare the training loss and the validation loss.

```
In [28]: history_dict = history.history  
history_dict.keys()
```


```
Out [28]: dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
```

```
In [30]: plt.plot(history.history['loss'], 'bo')  
plt.plot(history.history['val_loss'], 'r')  
plt.legend(['train', 'validation'])  
plt.show()
```

Loss



Accuracy



2022

Korea Institute of Science  
and Technology Information

TRUST  
KISTIL

